

COMPUTER VISION

LECTURE 3 11.9.2019

Laura Ruotsalainen, Associate Professor
Department of Computer Science




IMAGE PROCESSING, TODAY'S LECTURE

- Point processing, neighbourhood processing
 - Convolution
 - Box filter, Gaussian filter, other filters
 - Fourier transformation
-
- Szeliski chapter 3
 - However, Forsyth & Ponce provides much more comprehensive and logical explanation which is found from the slides








IMAGE PROCESSING



- Image = interaction of 3D scene objects, lighting, camera optics and sensors
- First step in computer vision is processing the image to be suitable for the computations needed => image processing
 - Exposure correction and color balancing
 - Noise reduction
 - Increasing sharpness
 - Rotating the image
- Point operators
- Neighborhood operators (linear filtering, non-linear filtering)
- Global operators


WHAT IS AN IMAGE?




color image patch

How many bits are
the intensity
values?


red






green



blue




colorized for visualization



actual intensity values per channel

Each channel
is a 2D array
of numbers.




WHAT IS AN IMAGE?

$$f(x)$$




grayscale image

What is the range
of the image
function f ?



$$\text{domain } \mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

A (grayscale)
image is a 2D
function.




TYPES OF IMAGE TRANSFORMATIONS



Filtering,
changes
pixel values



Warping,
changes
pixel
locations






POINT VS NEIGHBORHOOD OPERATION




- Brightness and contrast adjustments
- Color transformations

Point Operation




point processing

Neighborhood Operation



“filtering”



POINT PROCESSES

- Image processing operator (h) takes an input image (f) and produces an output image (g)


$$g(i, j) = h(f(i, j)).$$

- Common point processes are multiplication and addition with a constant

$$g(\mathbf{x}) = af(\mathbf{x}) + b.$$



- a is called a gain parameter controlling the contrast and b is called a bias parameter controlling brightness
- Both can be spatially varying for e.g. cool effects



EXAMPLES OF POINT PROCESSING



x

invert



$$255 - x$$



darken

$$x - 128$$

lighten



$$x + 128$$



lower contrast

$$\frac{x}{2}$$

raise contrast



$$x \times 2$$



non-linear lower contrast

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

non-linear raise contrast



$$\left(\frac{x}{255}\right)^2 \times 255$$




MANY OTHER TYPES OF POINT PROCESSING



camera output



image after stylistic tonemapping



NEIGHBORHOOD OPERATIONS



Szeliski

- Filtering
 - Local tone adjustment
 - Adding soft blur
 - Sharpening details
 - Accentuate edges
 - Remove noise
- Linear filtering = weighted combinations of pixels
- Non-linear filtering = morphological filters, distance transforms



(a)



(b)




(c)



(d)

Figure 3.11 Some neighborhood operations: (a) original image; (b) blurred; (c) sharpened; (d) smoothed with edge-preserving filter




LINEAR FILTERING AND CONVOLUTION



- *Shift invariant filtering*
 - value of the output depends on the pattern in an image neighborhood and not in its location
 - linear means that the output for the sum of two images equals the sum of the outputs obtained for the images separately
- *Kernel (mask, filter)* = pattern of weights used for linear filtering (h)
- *Convolution* = process of applying the linear filter

$$G_{ij} = \sum_{u,v} F_{i-u,j-v} H_{u,v}$$

where f is the original image and g is the resulting image



DEFINING CONVOLUTION 1




- F is the image, H is the kernel and G is the output image
- Kernel size is $2k+1 \times 2k+1$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

- This is a dot product between a certain local neighborhood and the kernel for each pixel
- Cross-correlation

$$G = H \otimes F$$




DEFINING CONVOLUTION 2



- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

- This is called a (discrete) **convolution** operation: $G = H * F$
- $H = \text{impulse response function}$, when kernel function H is convolved with impulse signal $\delta(i, j) \Rightarrow H * \delta = H$ whereas correlation produces the reflected signal



MORE PROPERTIES OF CONVOLUTION


Convolution is linear

Convolution is shift-invariant

Convolution is commutative ($w * f = f * w$)

Convolution is *associative* ($v * (w * f) = (v * w) * f$)

Every linear shift-invariant operation is a convolution



EXAMPLE: THE BOX FILTER




- also known as the 2D rect filter
- also known as the square mean filter

$$\text{kernel } h[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- replaces pixel with local average
- has smoothing (blurring) effect







MEAN FILTERING

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

H



0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0




F



	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
10	20	30	30	30	30	20	10		
10	10	10	0	0	0	0	0		



G




MEAN FILTERING/MOVING AVERAGE

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$




MEAN FILTERING/MOVING AVERAGE

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10							




MEAN FILTERING/MOVING AVERAGE

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20						




MEAN FILTERING/MOVING AVERAGE

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30					




MEAN FILTERING/MOVING AVERAGE

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30				



MEAN FILTERING/MOVING AVERAGE



- Kernels are usually shift- invariant = behave similarly everywhere in the image
- However, shift-variant kernels may be used e.g. for blurring an image due to variable depth-dependent defocus


$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

What about the edges?



COMPUTATIONAL REQUIREMENTS



- Convolution requires K^2 multiply-add operations per pixel, K is the size (width w or height h) of the kernel \Rightarrow Every entry takes $O(k^2)$ operations, Total time complexity: $O(whK^2)$
- The process is much faster if first a one-dimensional horizontal convolution followed by vertical is performed \Rightarrow separable kernels required
- Horizontal kernel (h), vertical (v) \Rightarrow Kernel $K = vh^T$, time complexity of separable version : $O(whk)$
- Kernel is separable if only the first singular value σ_0 is non-zero when Singular Value Decomposition (SVD) is taken of the Kernel K .



$$K = \sum_i \sigma_i \mathbf{u}_i \mathbf{v}_i^T$$


w, h = Image size

- Vertical kernel $v = \sqrt{\sigma_0} \mathbf{u}_0$, horizontal kernel $h = \sqrt{\sigma_0} \mathbf{v}_0^T$



SMOOTHING WITH BOX FILTER, RESULT






GAUSSIAN FILTER




- Averaging gives as much emphasize for the distant neighbors as for the close ones => blurring
- Kernel values sampled from the 2D Gaussian function


$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$



Small std, weights of distant pixels very small => little effect

Large std, consensus of pixels => noise removed effectively but blurring exists





GAUSSIAN FILTER


$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$


Ignore factor in front, instead,
normalize filter to sum to 1

0.003	0.013	0.022	0.013	0.003
0.013	0.060	0.098	0.060	0.013
0.022	0.098	0.162	0.098	0.022
0.013	0.060	0.098	0.060	0.013
0.003	0.013	0.022	0.013	0.003


5x5, $\sigma=1$

GAUSSIAN FILTERING EXAMPLE






GAUSSIAN VS BOX FILTERING



original



7x7 Gaussian



7x7 box



OTHER TYPES OF FILTERS

- Band-pass filtering (e.g. Sobel filters), sophisticated filters made by
 - First smoothing the image with a Gaussian filter
 - Then taking the first or second derivatives of the image (Laplacian operator)
- Non-linear filters for improved performance e.g. when the noise is not Gaussian distributed
 - e.g. Median filtering => median value from each pixels neighbourhood



Shot noise




Gaussian filtered




Median filtered





Bilaterally filtered




DERIVATIVE OF GAUSSIAN FILTER




x-direction




y-direction






FOURIER TRANSFORMS 1

- So far the signal $f(x,y)$ has been considered as being a weighted sum of a large on infinite number of small box functions
- We need to consider two problems:
 - we don't know what is lost when representing the continuous signal with a discrete one => **sampling**
 - we don't know how to shrink an image, can't just take every k th pixel
- Problems are related to fast changes in the image => something important might be missed
- These can be studied by the change of basis => **Fourier transform**
 - basis will be a set of sinusoids, signal infinite weighted sum of an infinite number of sinusoids



FOURIER TRANSFORMS 2




The Fourier transform as a response of a filter $h(x)$ to an input sinusoid $s(x)$ yielding an output sinusoid $o(x)$

$$H(k) = \frac{1}{N} \sum_{x=0}^{N-1} h(x) e^{-j \frac{2\pi kx}{N}},$$

Discrete Fourier Transform (DFT)


- Fourier transform
 - used to understand the difference between continuous and discrete images => what has been lost
 - is simply a tabulation of the magnitude and phase response at each frequency



FOURIER TRANSFORMS 3

- A video showing how Fourier Transforms are used for images

<https://www.youtube.com/watch?v=gwaYwRwY6PU>



FOURIER TRANSFORMS 4




- A good tutorial is found from <http://homepages.inf.ed.ac.uk/rbf/HIPR2/fourier.htm>


Sonnet for Lena

O dear Lena, your beauty is so vast
It is hard sometimes to describe it fast.
I thought the entire world I would impress
If only your portrait I could compress.
Alas! First when I tried to use VQ
I found that your cheeks belong to only you.
Your silky hair contains a thousand lines
Hard to match with sums of discrete cosines.
And for your lips, sensual and tactual
Thirteen Crays found not the proper fractal.
And while those setbacks are all quite severe
I might have fixed them with hacks here or there
But when filters took sparkle from your eyes
I said, 'Damn all this. I'll just digitize.'

Thomas Colton



Logarithm of DFT's magnitude



Thresholded magnitude of the Fourier image

For a square image of size $N \times N$, the two-dimensional DFT is given by:



$$F(k, l) = \sum_{i=0}^{N-1} \sum_{j=0}^{N-1} f(i, j) e^{-i2\pi(\frac{ki}{N} + \frac{lj}{N})}$$



FOURIER TRANSFORMS 5



- Rotate the image 45 degrees







IMAGE SCALING

This image is too big to fit on the screen. How can we generate a half-sized version?



Source: S. Seitz





IMAGE SUB-SAMPLING



1/2




1/4 (2x zoom)




1/8 (4x zoom)

Throw away every other row and column to create a 1/2 size image: called *image sub-sampling*




ALIASING



- Occurs when sampling rate is not high enough to capture the amount of detail in your image
- Can give the wrong signal/image—an *alias*
- To do sampling right, need to understand the structure of the signal/image
- To avoid aliasing:
 - Sampling rate $\geq 2 * \text{max frequency in the image}$ = \geq two samples per cycle
 - This minimum sampling rate is called the **Nyquist rate**
 - k in DFT should be in the range $k \in [-N/2, N/2]$
- When downsampling by a factor of two the original image has frequencies that are too high



SUBSAMPLING WITH GAUSSIAN PRE-FILTERING



Gaussian $1/2$



G $1/4$




G $1/8$

- Solution: filter the image, *then* subsample



GAUSSIAN PRE- FILTERING

- Solution: filter the image, *then* subsample



F_0



F_1



F_2



blur



subsample




blur




subsample

...



F_0^*


H





F_1^*

H





*Gaussian
pyramid*



F_0



F_1



F_2



blur



subsample




blur




subsample

...



F_0^*



H




F_1^*

H









GAUSSIAN PYRAMIDS [BURT AND ADELSON, 1983]

Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N=2^k$)



- A precursor to *wavelet transform*
- Gaussian Pyramids have all sorts of applications in computer vision



UPSAMPLING



This image is too small for this screen:

How can we make it 10 times as big?

Simplest approach:

repeat each row

and column 10 times

(“Nearest neighbor interpolation”)







IMAGE INTERPOLATION



$d = 1$ in this example

Digital images are formed as follows:

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale





IMAGE INTERPOLATION



$d = 1$ in this example

Digital images are formed as follows:

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale





IMAGE INTERPOLATION



$d = 1$ in this example

- What if we don't know f ?

- Guess an approximation: \tilde{f}
- Can be done in a principled way: filtering
- Convert F to a continuous function:

$$f_F(x) = F\left(\frac{x}{d}\right) \text{ when } \frac{x}{d} \text{ is an integer, } 0 \text{ otherwise}$$

- Reconstruct by convolution with a *reconstruction filter*, h

$$\tilde{f} = h * f_F$$





Image interpolation






Image interpolation

Original image:  x 10



Nearest-neighbor interpolation



Bilinear interpolation



Bicubic interpolation



60° 10 1.2 N, 24° 57 18 E