

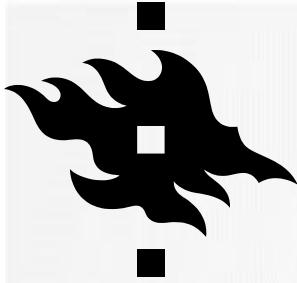
# TODAY'S LECTURE



1. Feature detectors
  1. Harris feature detector
2. Feature descriptors
  1. SIFT
3. Feature matching
4. Edge detection
5. Line detection

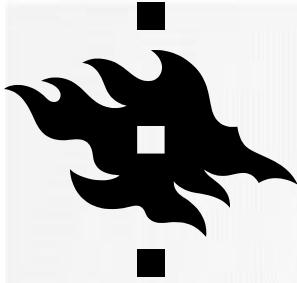
Features are used for:

Image alignment (e.g., mosaics)  
3D reconstruction  
Motion tracking (e.g. for AR)  
Object recognition  
Image retrieval  
Navigation  
... other



# USEFUL FEATURES

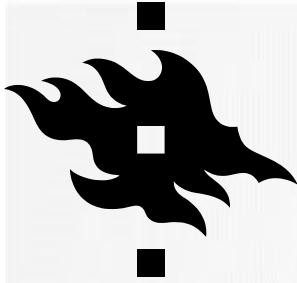




# APPROACH

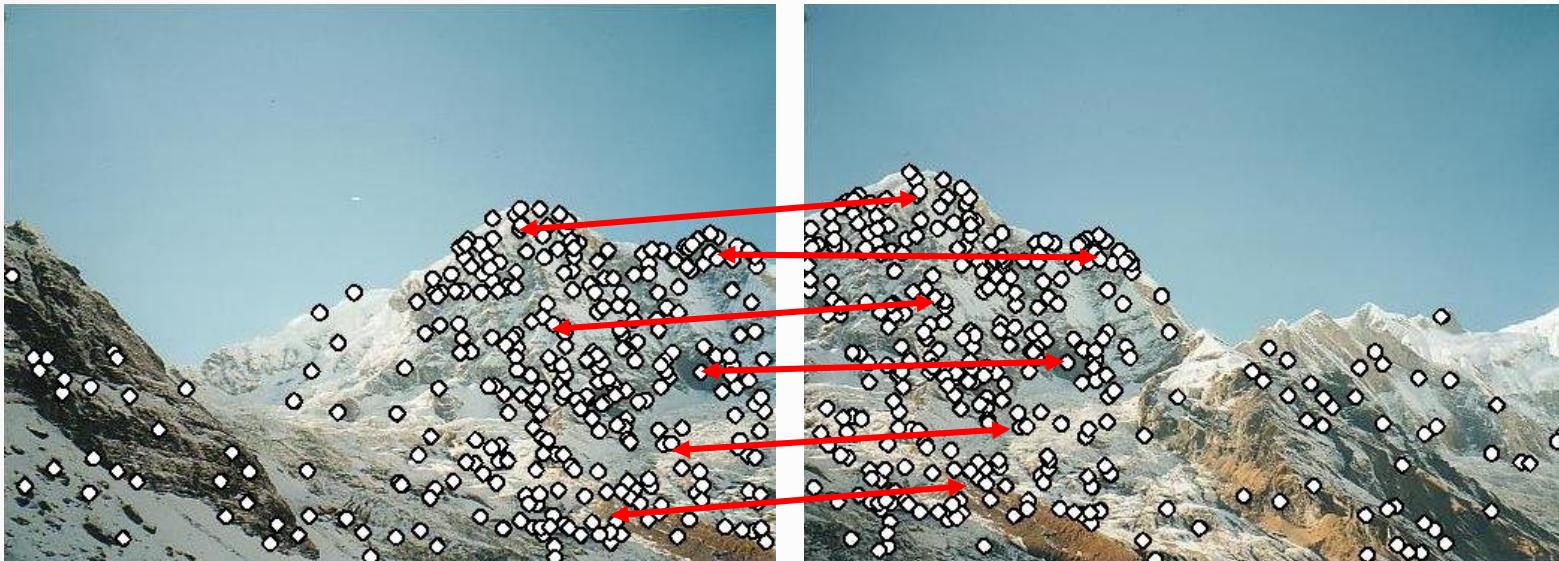


- 1. Feature detection:** find it
- 2. Feature descriptor:** represent it
- 3. Feature matching:** match it / **Feature tracking:** track it, when motion

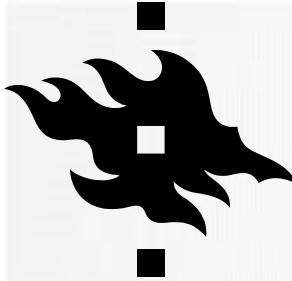


# EXTRACTING FEATURES

1. Feature detection
2. Feature descriptor
3. Feature matching



Slide credit: Kris Kitani



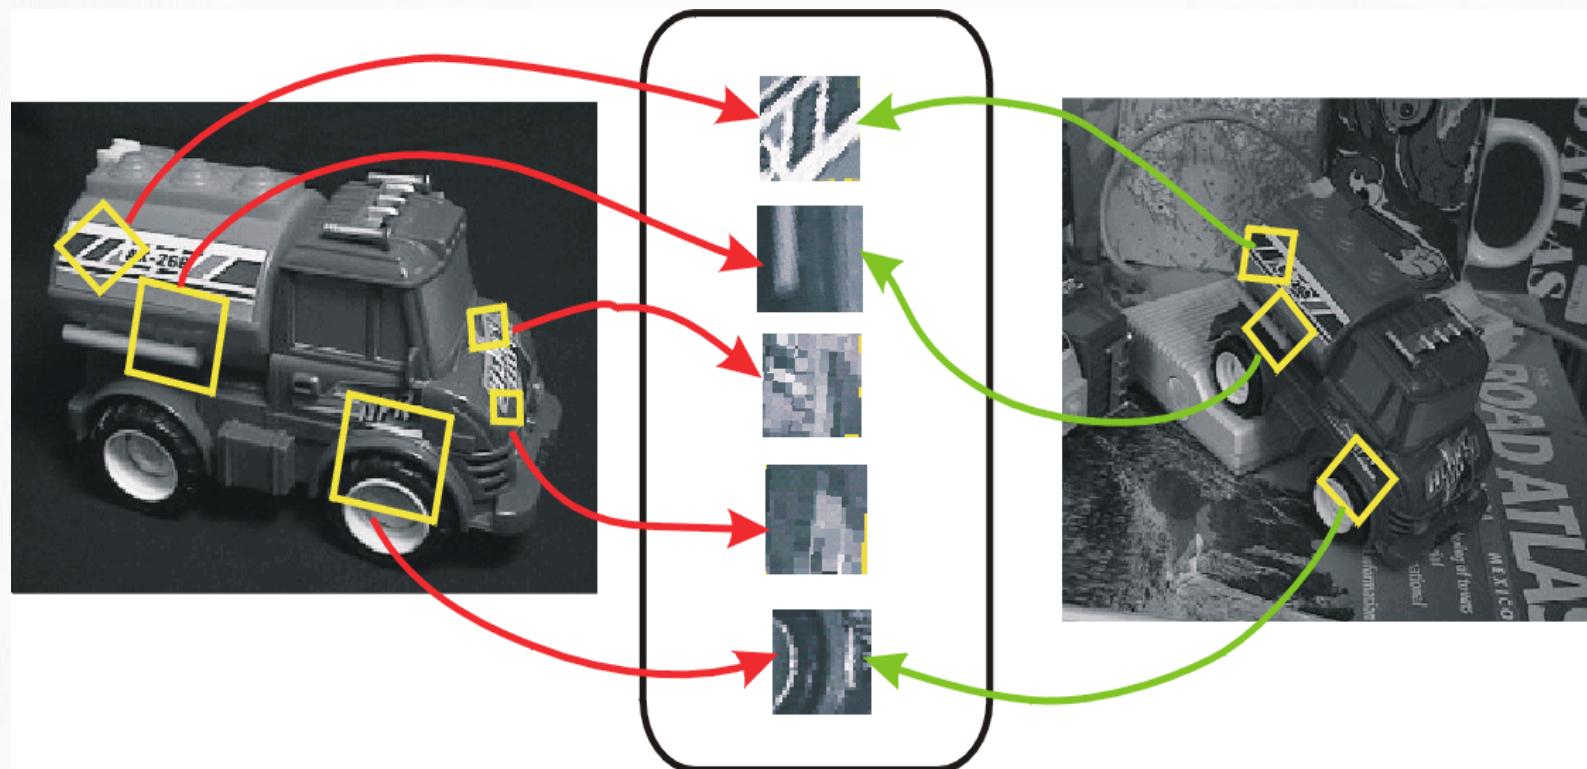
# INVARIANT LOCAL FEATURES

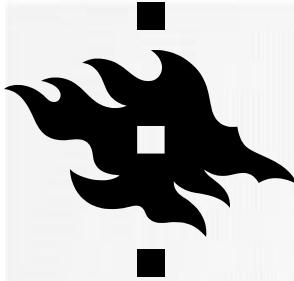


Find features that are invariant to transformations

geometric invariance: translation, rotation, scale

photometric invariance: brightness, exposure, ...





# ADVANTAGES OF LOCAL FEATURES

Locality

features are local, so robust to occlusion and clutter

Quantity

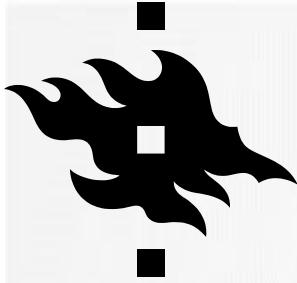
hundreds or thousands in a single image, in a favorable scene, please see my example later

Distinctiveness:

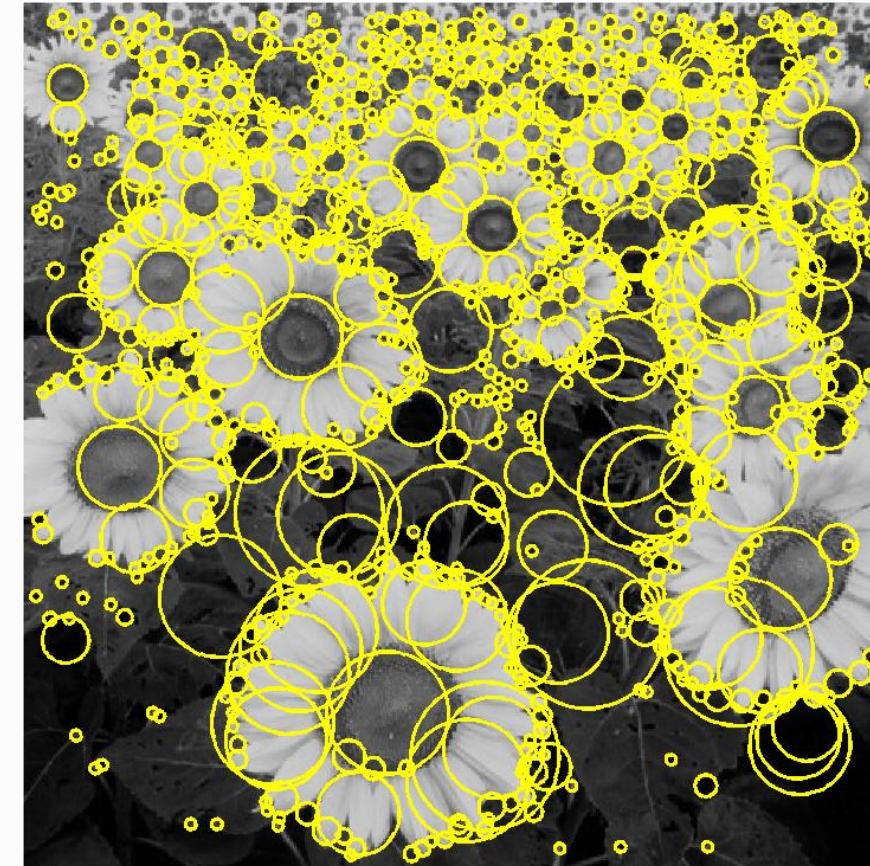
can differentiate a large database of objects

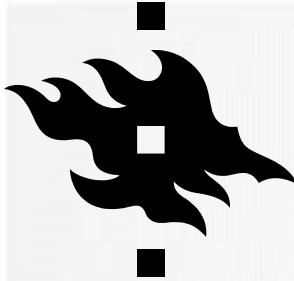
Efficiency

real-time performance achievable



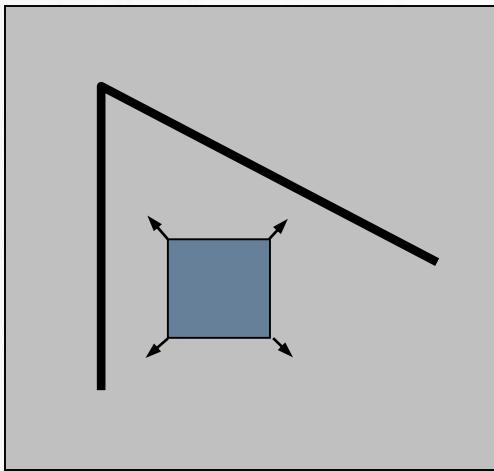
# FEATURE EXTRACTION: CORNERS AND BLOBS



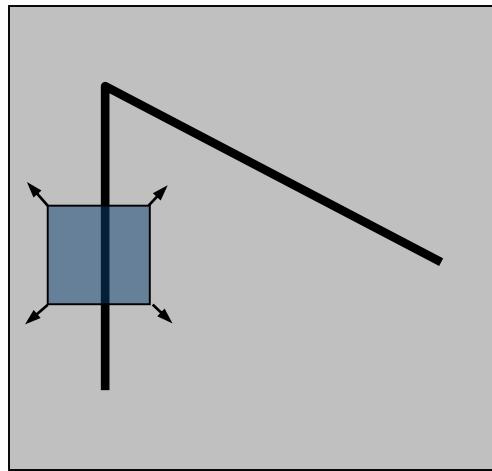


# LOCAL MEASURES OF UNIQUENESS

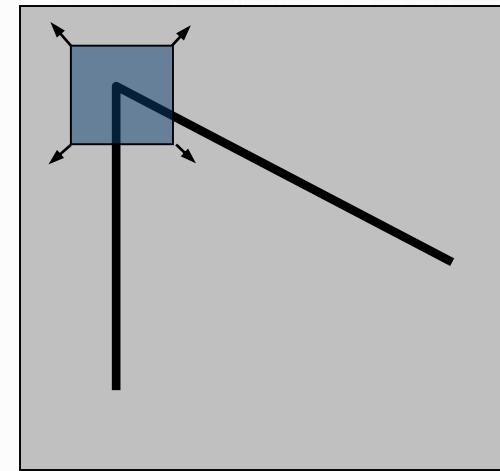
- How does the window change when you shift it?
- Shifting the window in any direction causes a big change



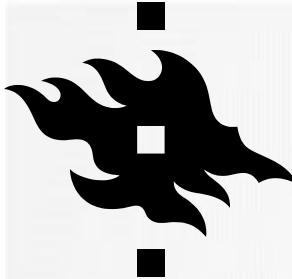
“flat” region:  
no change in all  
directions



“edge”:  
no change along  
the edge direction



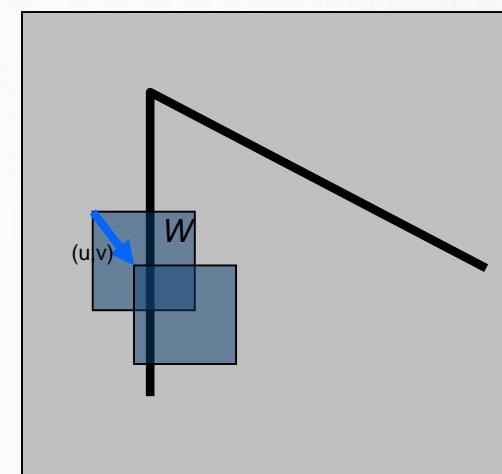
“corner”:  
significant change  
in all directions



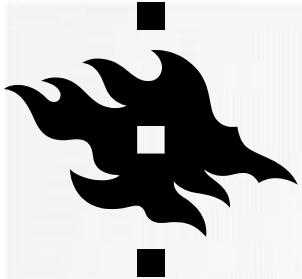
# HARRIS CORNER DETECTION: THE MATH

- Consider shifting the window  $W$  by  $(u, v)$ 
  - how do the pixels in  $W$  change?
  - compare each pixel before and after by summing up the squared differences (SSD)
  - this defines an SSD “error”  $E(u, v)$ :

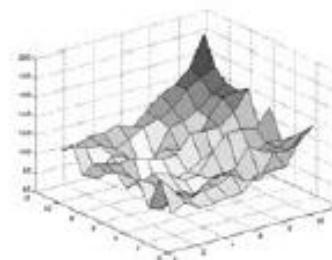
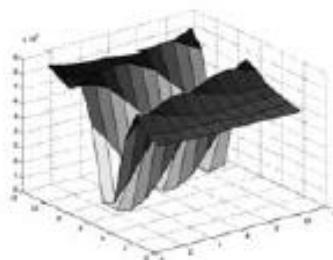
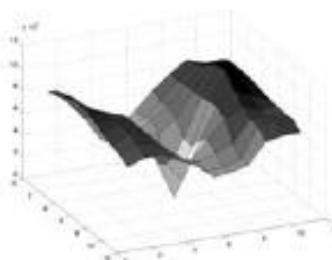
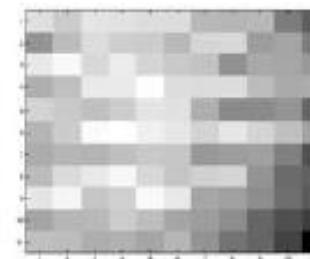
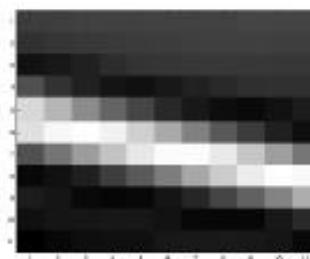
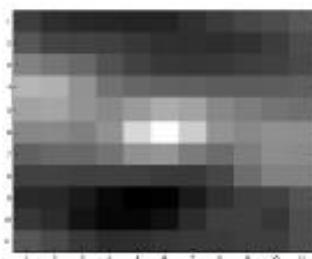
$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

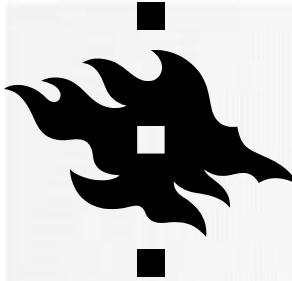


- We are happy if this error is high
- Slow to compute exactly for each pixel and each offset  $(u, v)$



(a)





# SMALL MOTION ASSUMPTION

Taylor Series expansion of  $I$ :

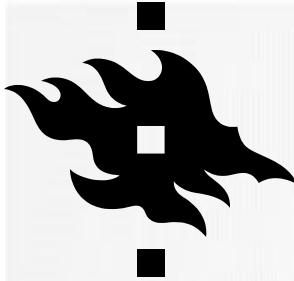
$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

If the motion  $(u, v)$  is small, then first order approximation is good

$$\begin{aligned} I(x + u, y + v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

shorthand:  $I_x = \frac{\partial I}{\partial x}$

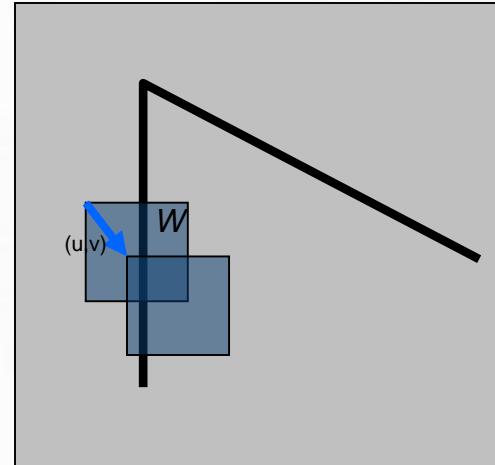
Plugging this into the formula on the previous slide...



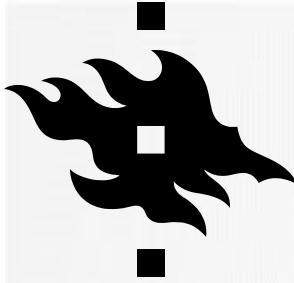
# CORNER DETECTION: THE MATH

Consider shifting the window  $W$  by  $(u, v)$

- define an SSD “error”  $E(u, v)$ :



$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \end{aligned}$$

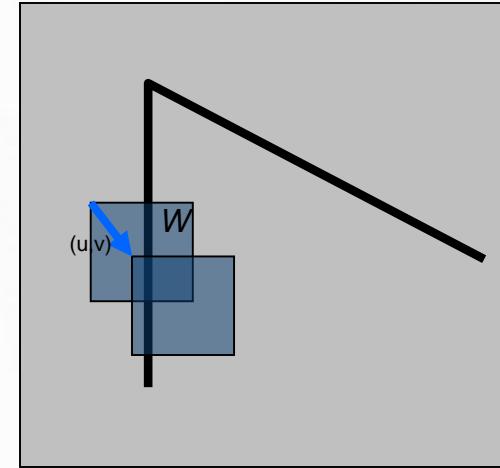


# CORNER DETECTION: THE MATH

Consider shifting the window  $W$  by  $(u, v)$

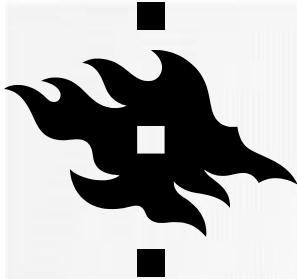
- define an SSD “error”  $E(u, v)$ :

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} [I_x u + I_y v]^2 \\ &\approx Au^2 + 2Buv + Cv^2 \end{aligned}$$



$$A = \sum_{(x,y) \in W} I_x^2 \quad B = \sum_{(x,y) \in W} I_x I_y \quad C = \sum_{(x,y) \in W} I_y^2$$

- Thus,  $E(u, v)$  is locally approximated as a quadratic error function



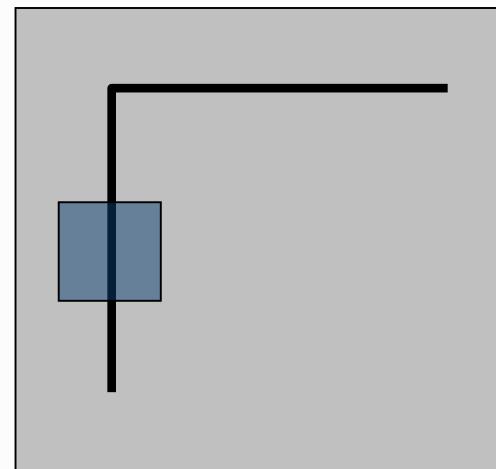
$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$



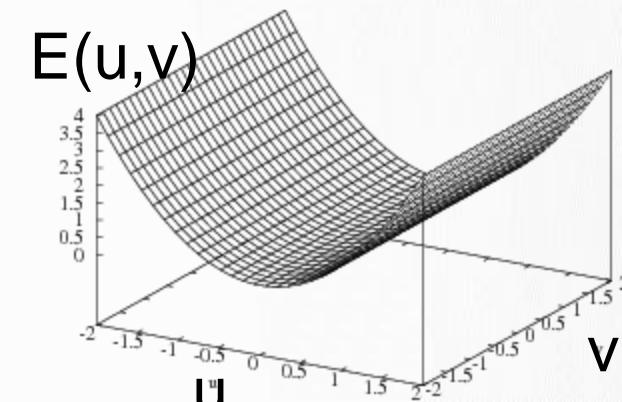
$$A = \sum_{(x,y) \in W} I_x^2$$

$$B = \sum_{(x,y) \in W} I_x I_y$$

$$C = \sum_{(x,y) \in W} I_y^2$$

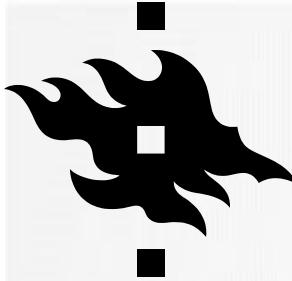


$$H = \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix}$$



Vertical edge:

$$I_y = 0$$

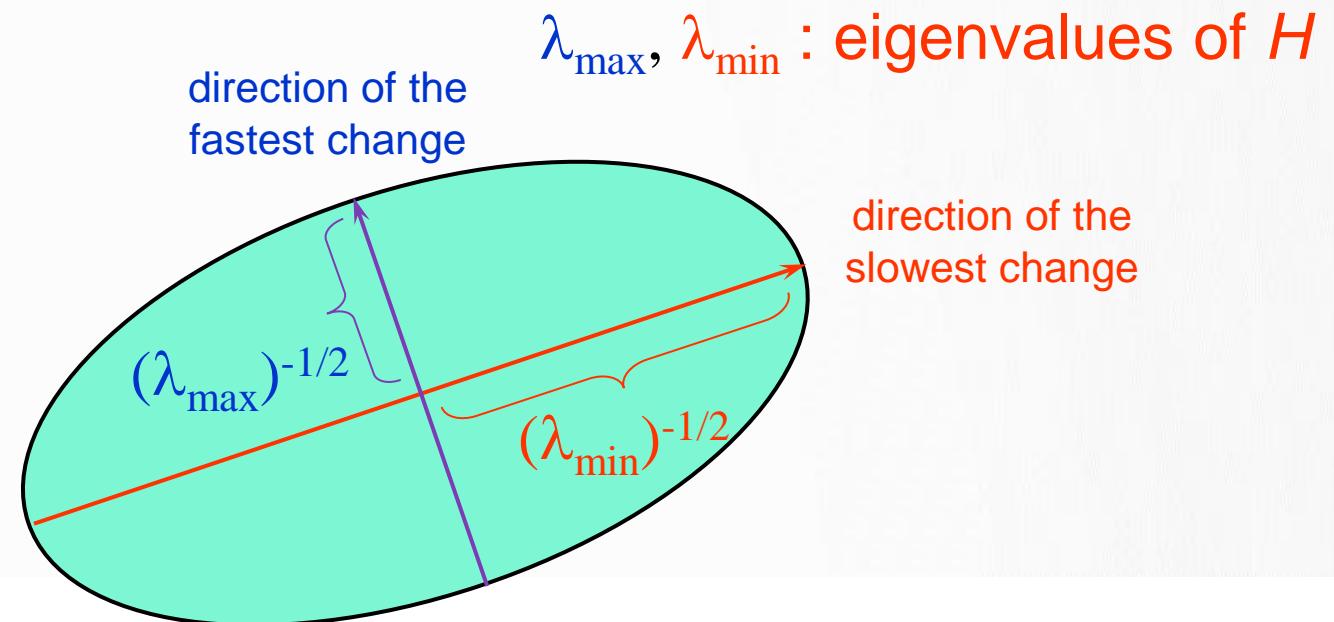


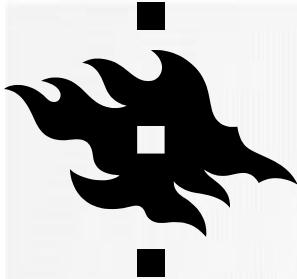
# GENERAL CASE

We can visualize  $H$  as an ellipse with axis lengths determined by the *eigenvalues* of  $H$  and orientation determined by the *eigenvectors* of  $H$

Ellipse equation:

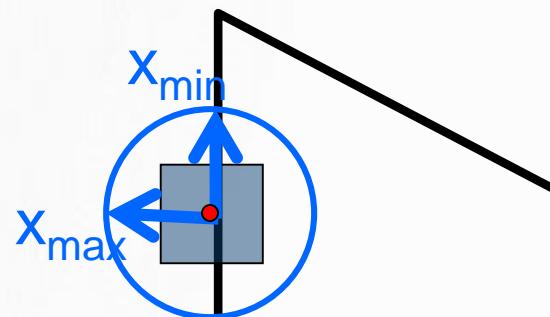
$$[u \ v] H \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$





# CORNER DETECTION: THE MATH

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \underbrace{\begin{bmatrix} A & B \\ B & C \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$



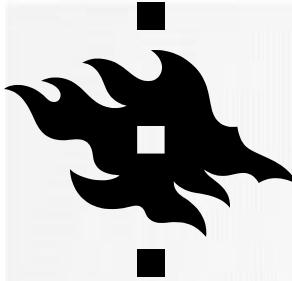
$$H$$

$$Hx_{\max} = \lambda_{\max} x_{\max}$$

$$Hx_{\min} = \lambda_{\min} x_{\min}$$

Eigenvalues and eigenvectors of  $H$

- Define shift directions with the smallest and largest change in error
- $x_{\max}$  = direction of largest increase in  $E$
- $\lambda_{\max}$  = amount of increase in direction  $x_{\max}$
- $x_{\min}$  = direction of smallest increase in  $E$
- $\lambda_{\min}$  = amount of increase in direction  $x_{\min}$



# CORNER DETECTION: THE MATH

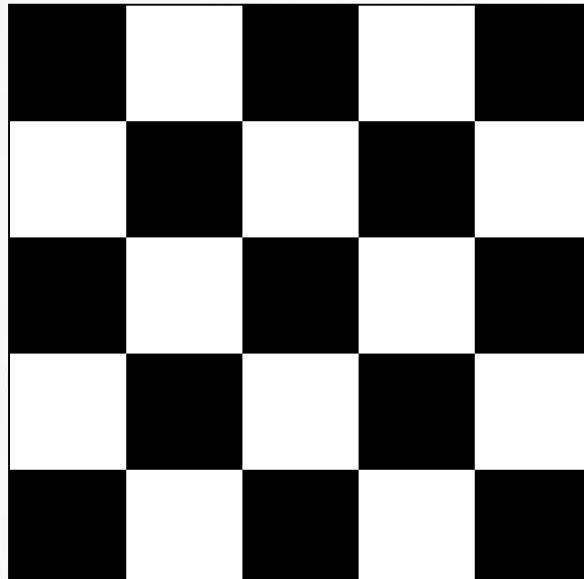


How are  $\lambda_{\max}$ ,  $x_{\max}$ ,  $\lambda_{\min}$ , and  $x_{\min}$  relevant for feature detection?

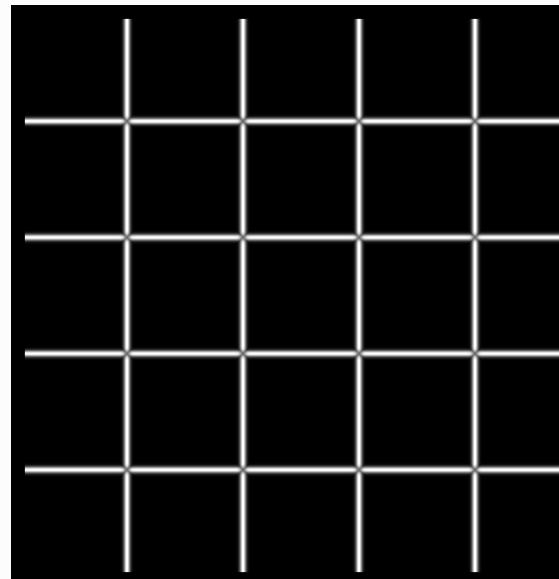
- What's our feature scoring function?

Want  $E(u, v)$  to be large for small shifts in all directions

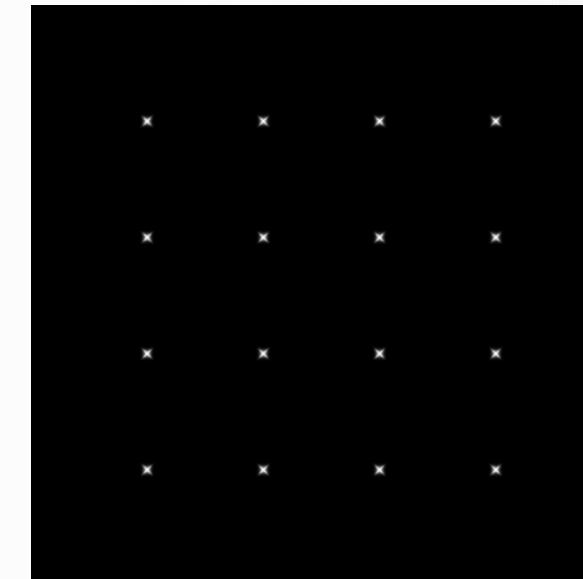
- the minimum of  $E(u, v)$  should be large, over all unit vectors  $[u \ v]$
- this minimum is given by the smaller eigenvalue ( $\lambda_{\min}$ ) of  $H$



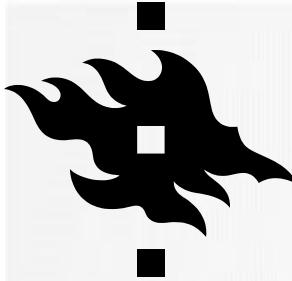
$I$



$\lambda_{\max}$



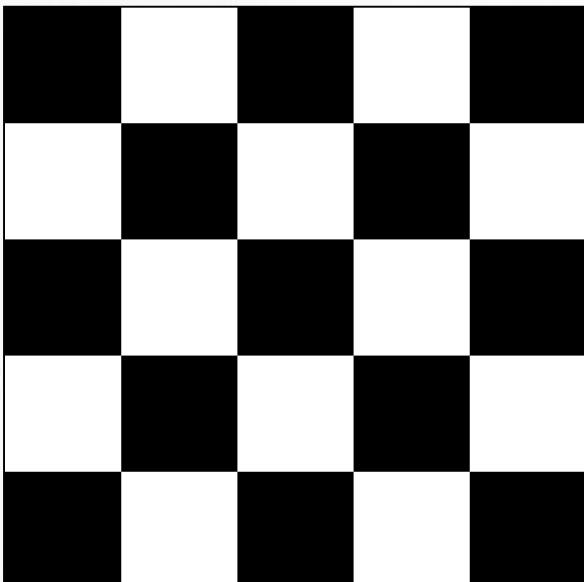
$\lambda_{\min}$



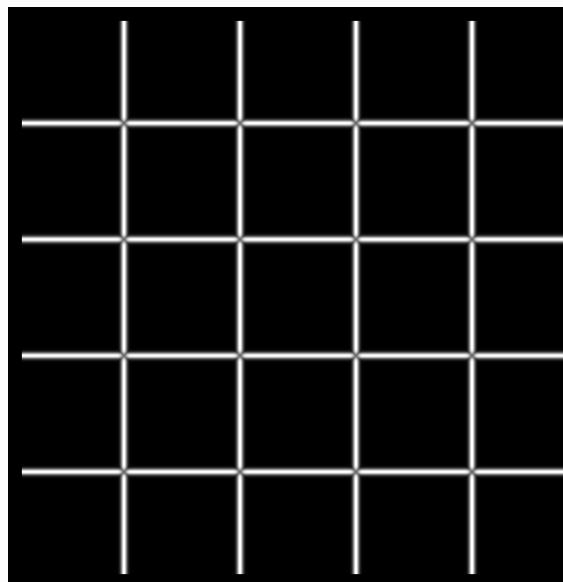
# CORNER DETECTION SUMMARY

Here's what you do

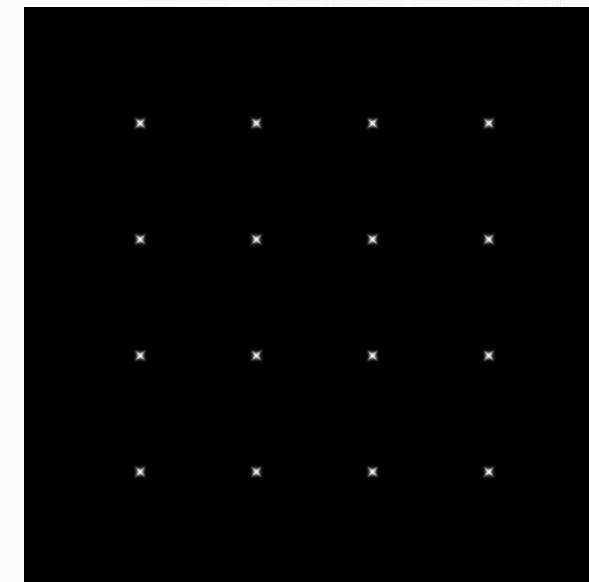
- Compute the gradient at each point in the image
- Create the  $H$  matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ( $\lambda_{\min} > \text{threshold}$ )
- Choose those points where  $\lambda_{\min}$  is a local maximum as features



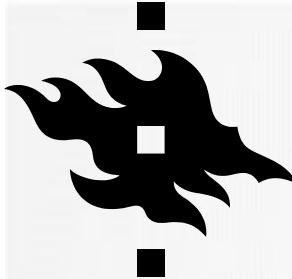
$I$



$\lambda_{\max}$



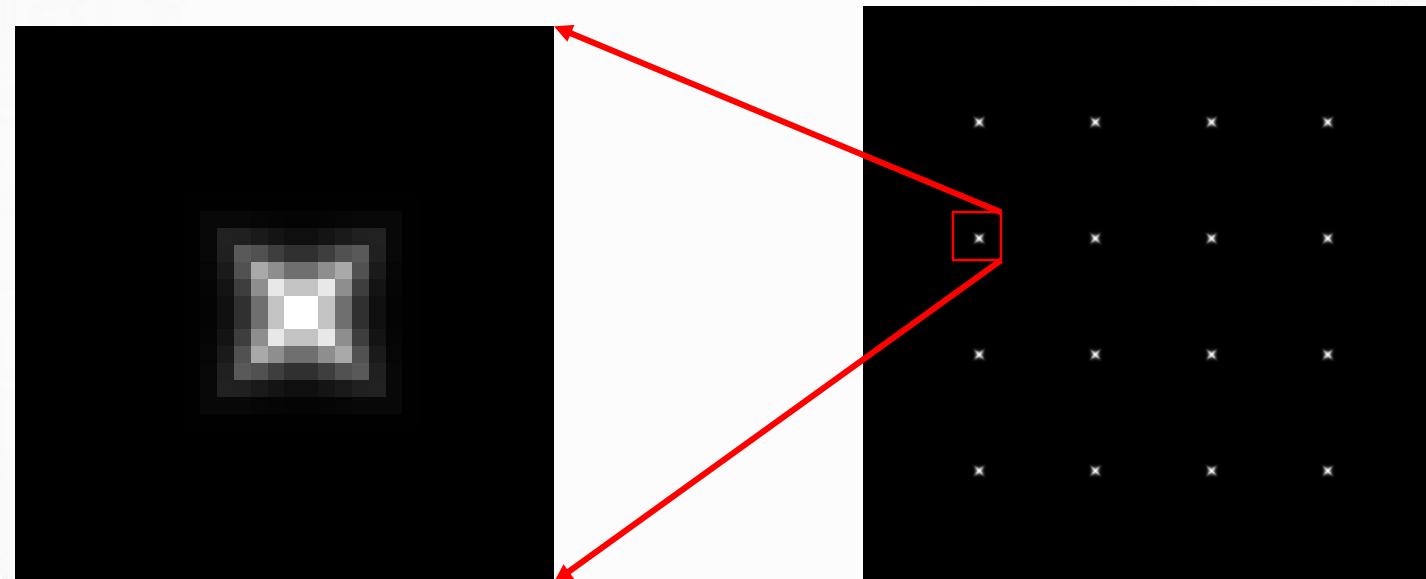
$\lambda_{\min}$

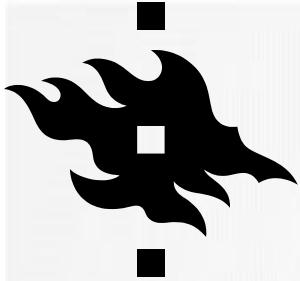


# CORNER DETECTION SUMMARY



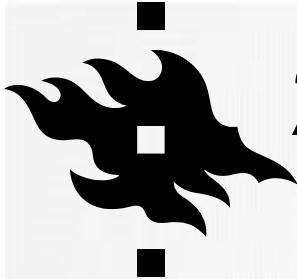
- Compute the gradient at each point in the image
- Create the  $H$  matrix from the entries in the gradient
- Compute the eigenvalues.
- Find points with large response ( $\lambda_{\min} > \text{threshold}$ )
- Choose those points where  $\lambda_{\min}$  is a local maximum as features



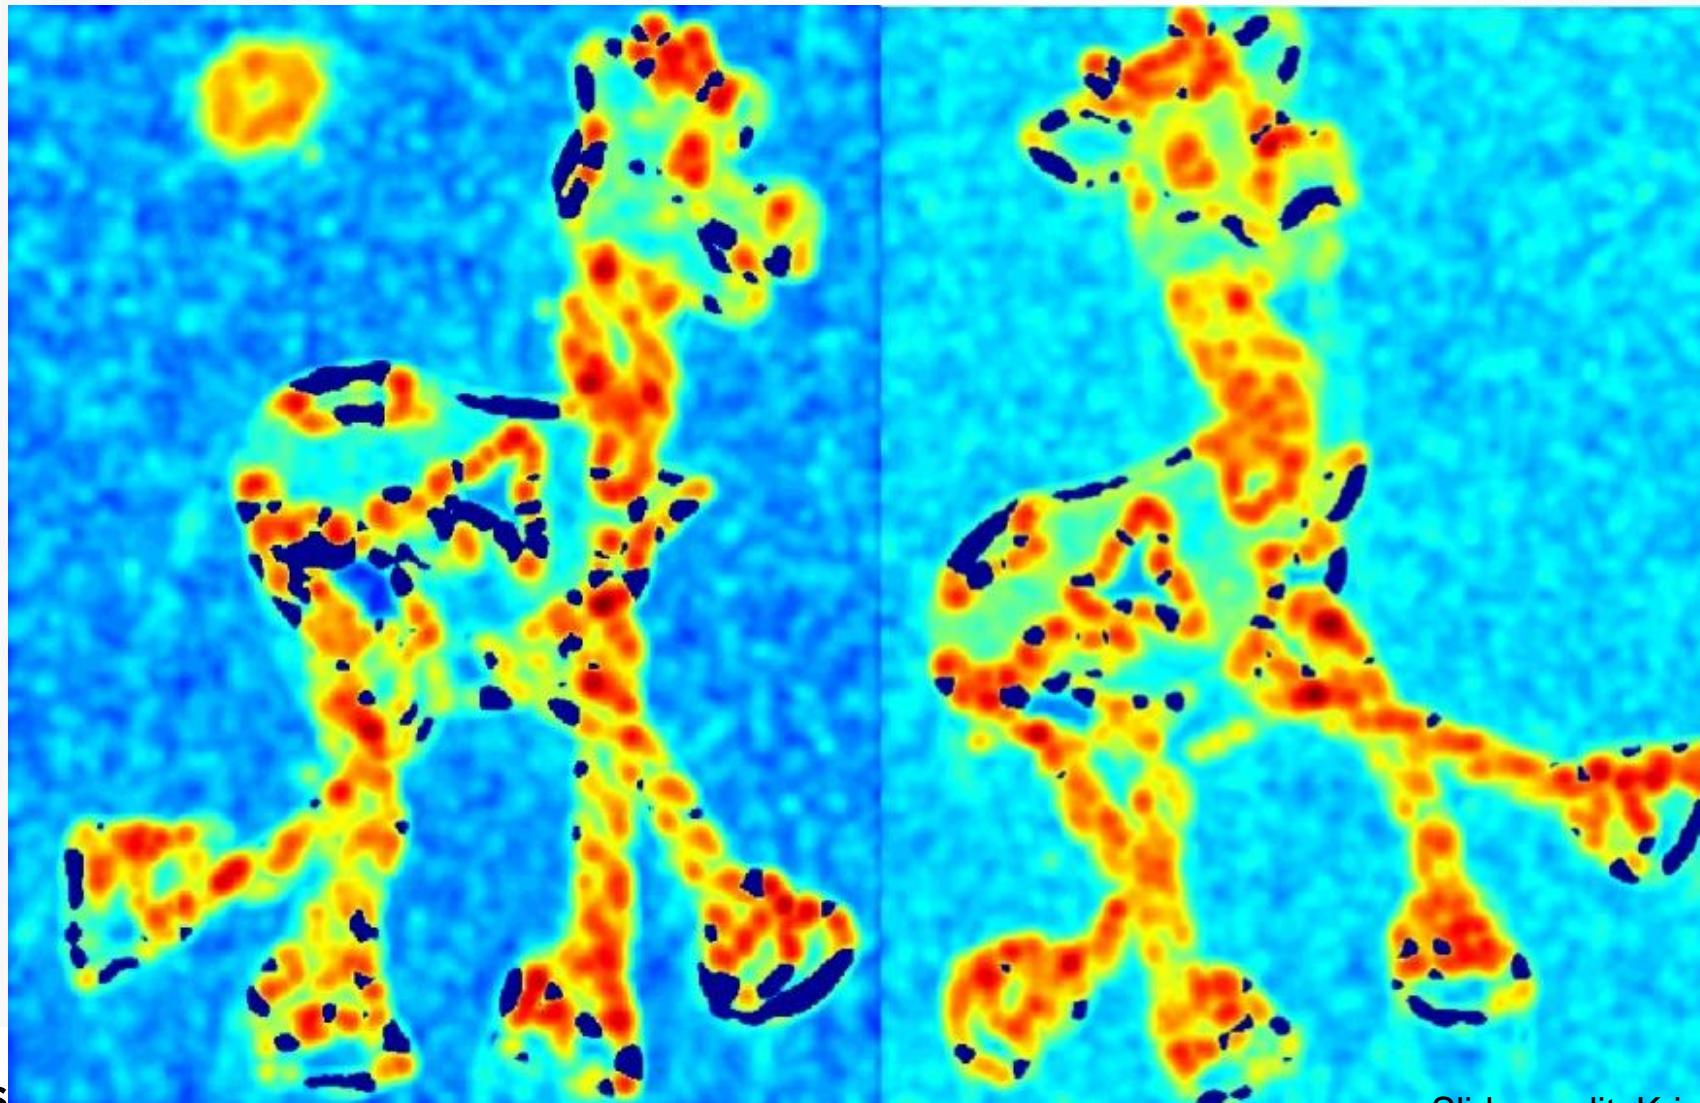


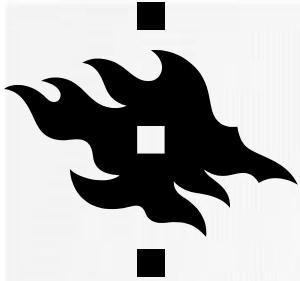
# HARRIS DETECTOR EXAMPLE



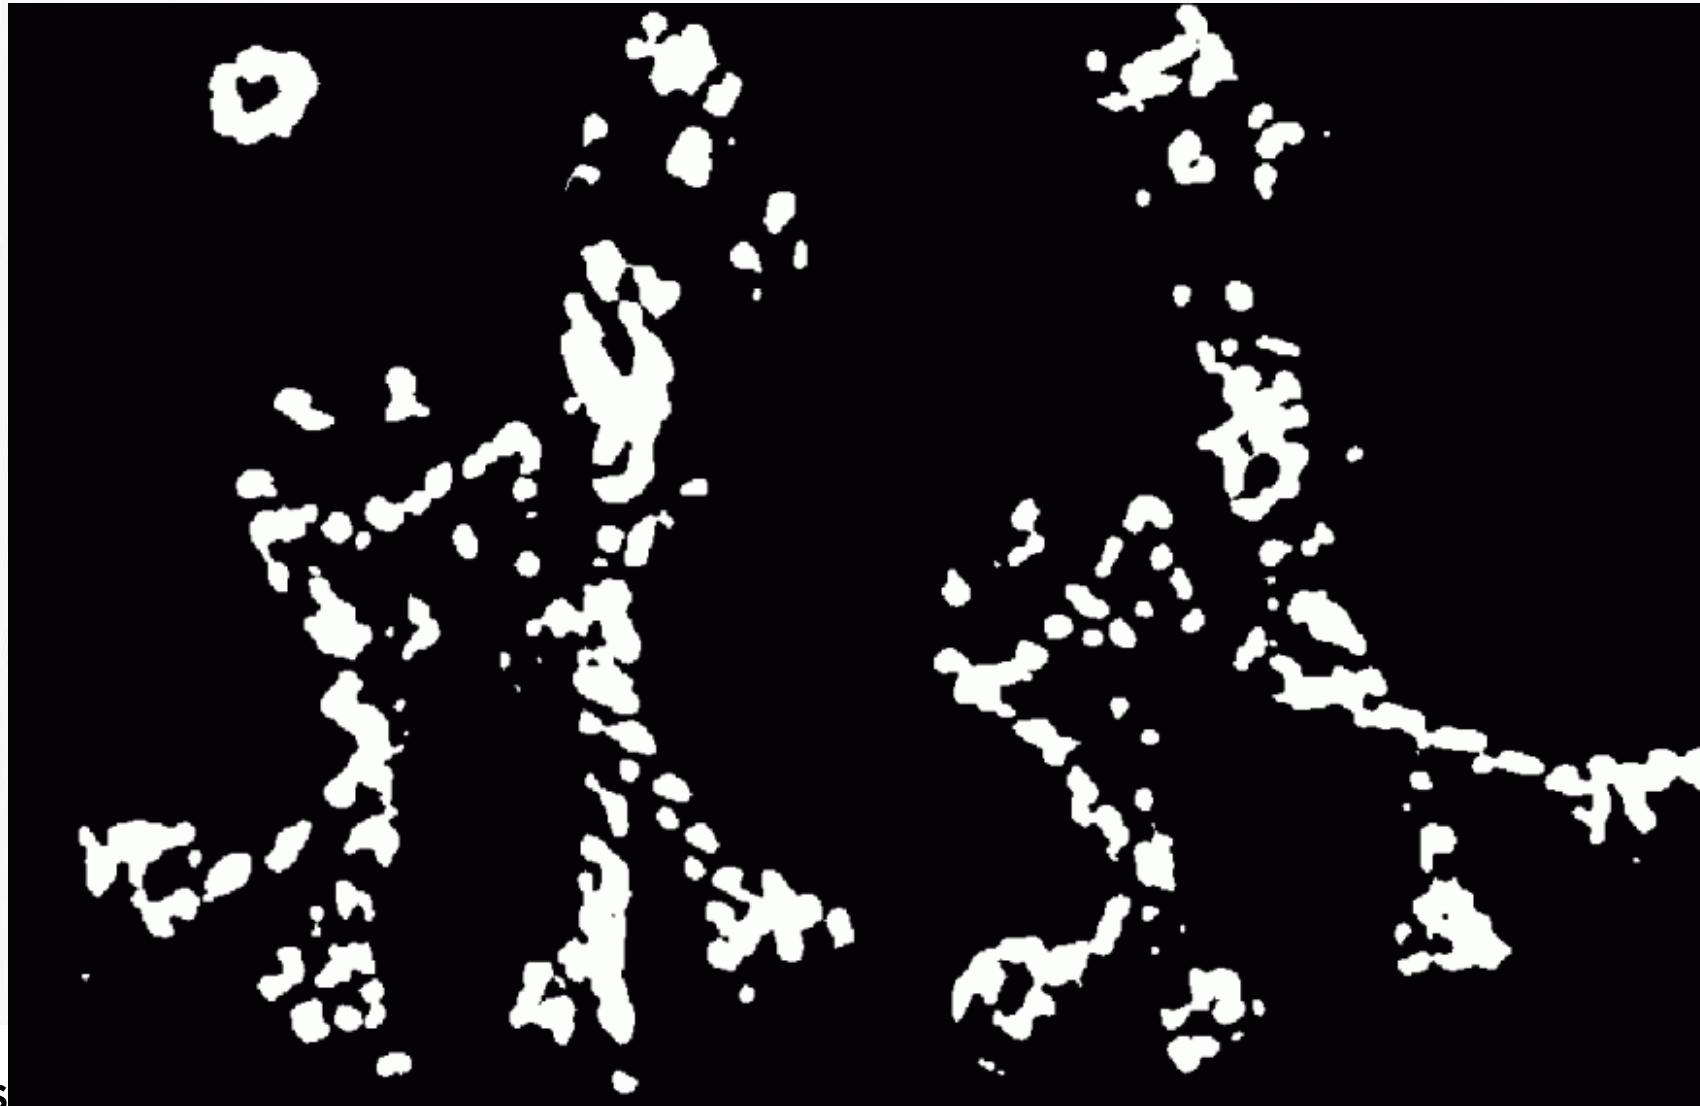


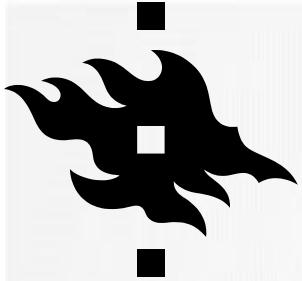
# $\lambda_{\text{MIN}}$ VALUE (RED HIGH, BLUE LOW)



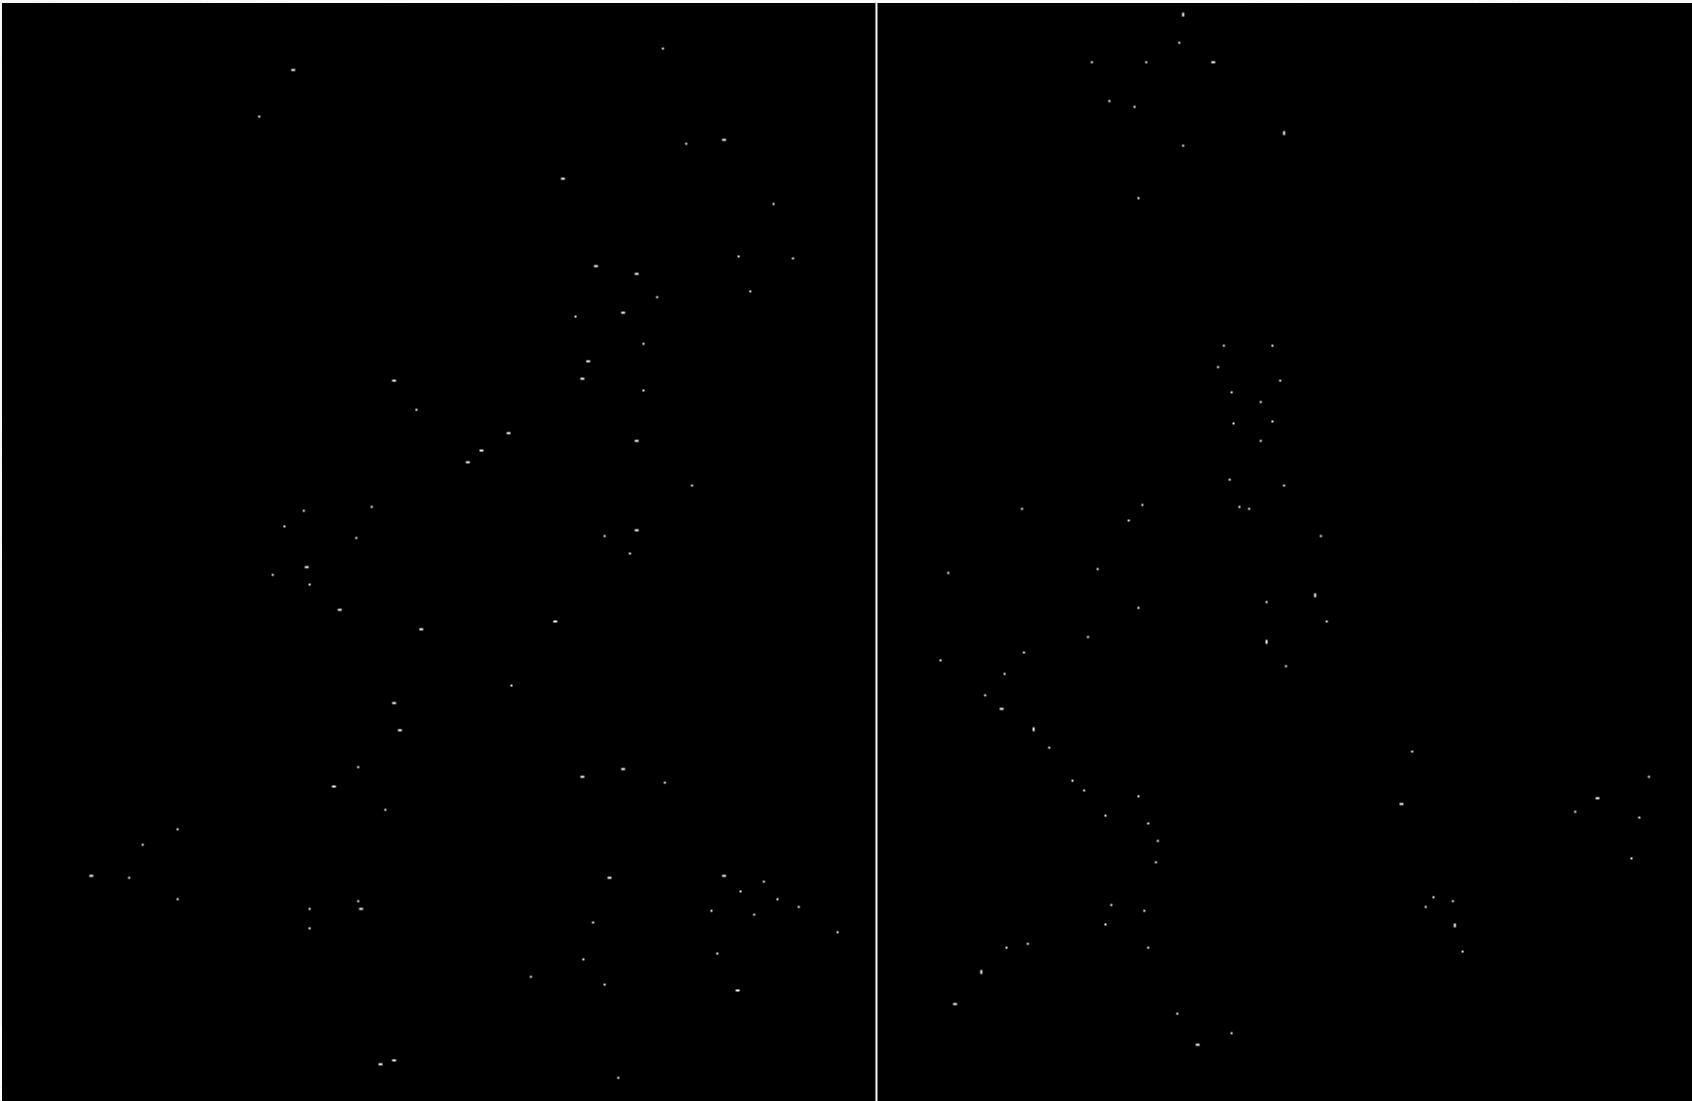


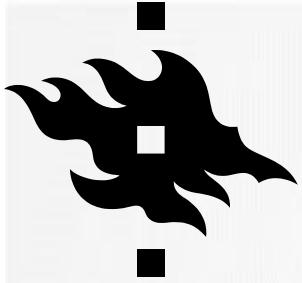
# THRESHOLD ( $\lambda_{\text{MIN}} > \text{VALUE}$ )





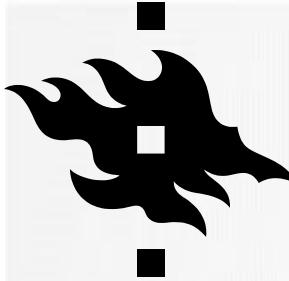
# FIND LOCAL MAXIMA OF $\lambda_{\text{MIN}}$





# HARRIS FEATURES (IN RED)





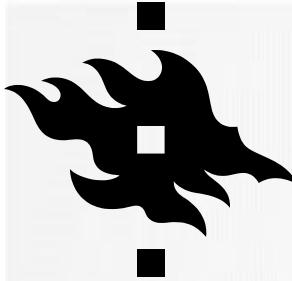
# INVARIANCE AND EQUIVARIANCE

- We want corner locations to be *invariant* to photometric transformations and *equivariant* to geometric transformations
  - **Invariance:** image is transformed and corner locations do not change
  - **Equivariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

(Sometimes “invariant” and “equivariant” are both referred to as “invariant”)

(Sometimes “equivariant” is called “covariant”)

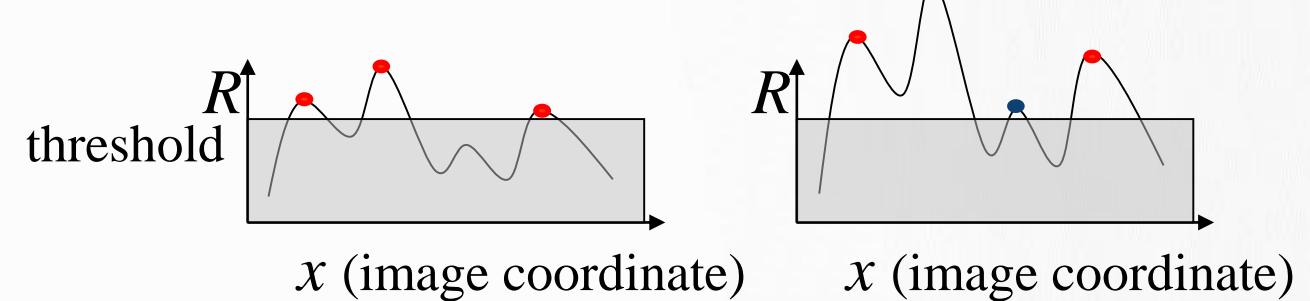
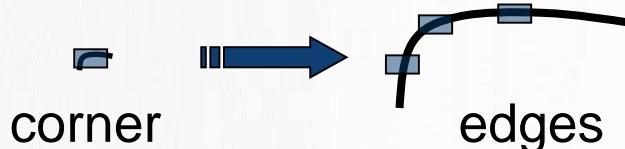


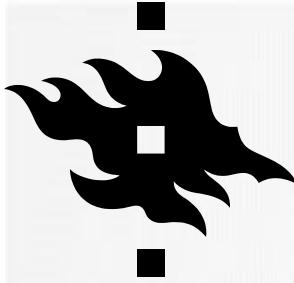


# HARRIS DETECTOR: INVARIANCE



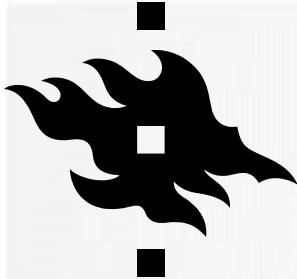
- Detector is invariant to
  - translation: derivatives and window function are equivariant
  - rotation: second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same
- Partially invariant to:
  - affine intensity change
- Is not invariant to scaling!





# FEATURE DESCRIPTORS

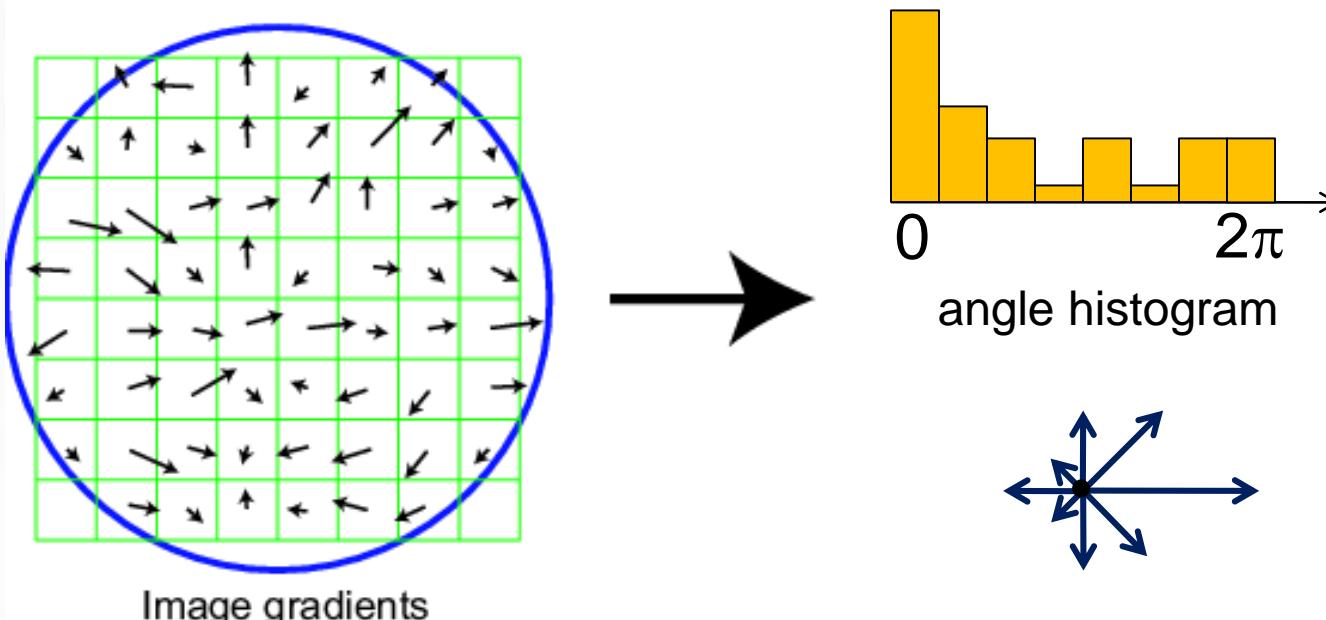
- A good feature descriptor should
  - be invariant: it shouldn't change even if image is transformed
  - be highly unique for each point
- SIFT is a very popular descriptor
- Other, e.g.
  - HOG: Histogram of Gradients (HOG)
  - FREAK: Fast Retina Keypoint, used in Visual SLAM
  - LIFT: Learned Invariant Feature Transform, learned via deep learning  
<https://arxiv.org/abs/1603.09114>

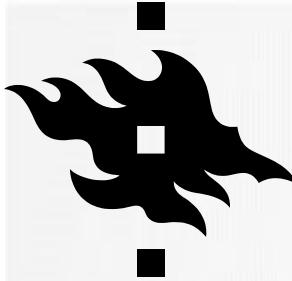


# SCALE INVARIANT FEATURE TRANSFORM

Basic idea:

- Take 16x16 square window around detected feature
- Compute edge orientation (angle of the gradient -  $90^\circ$ ) for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations

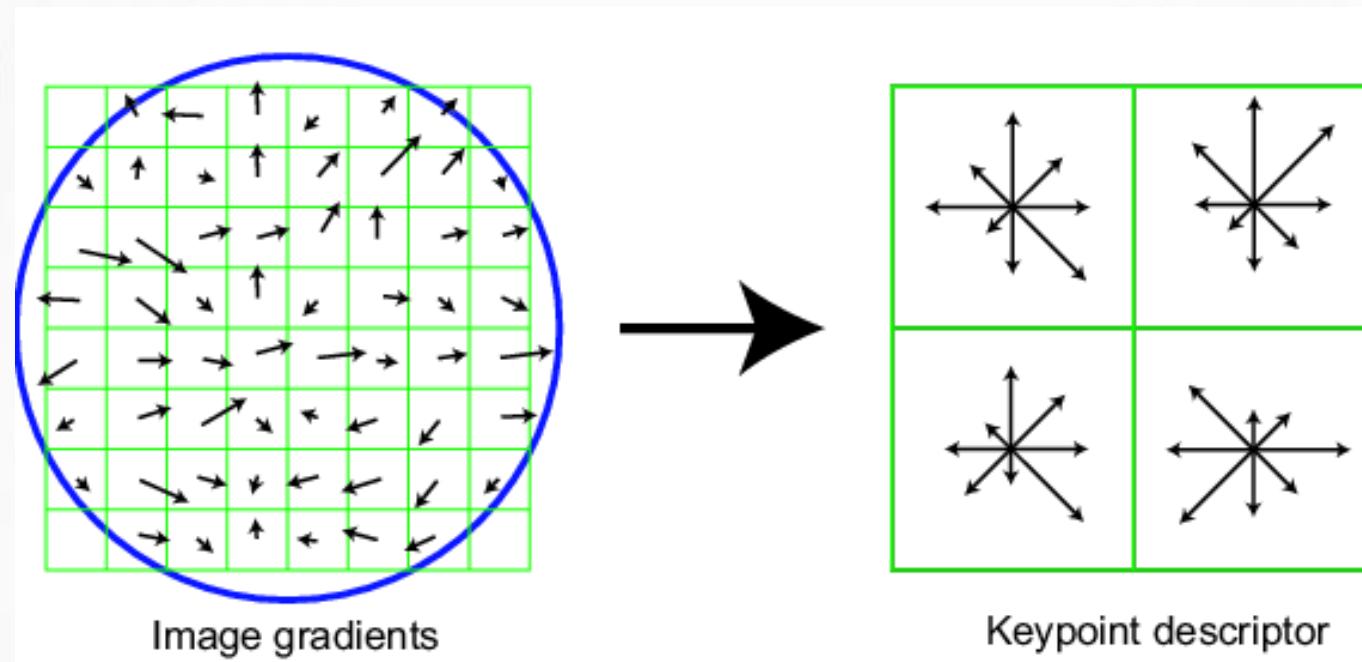


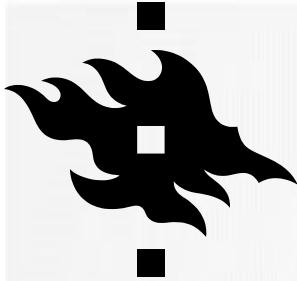


# SIFT DESCRIPTOR

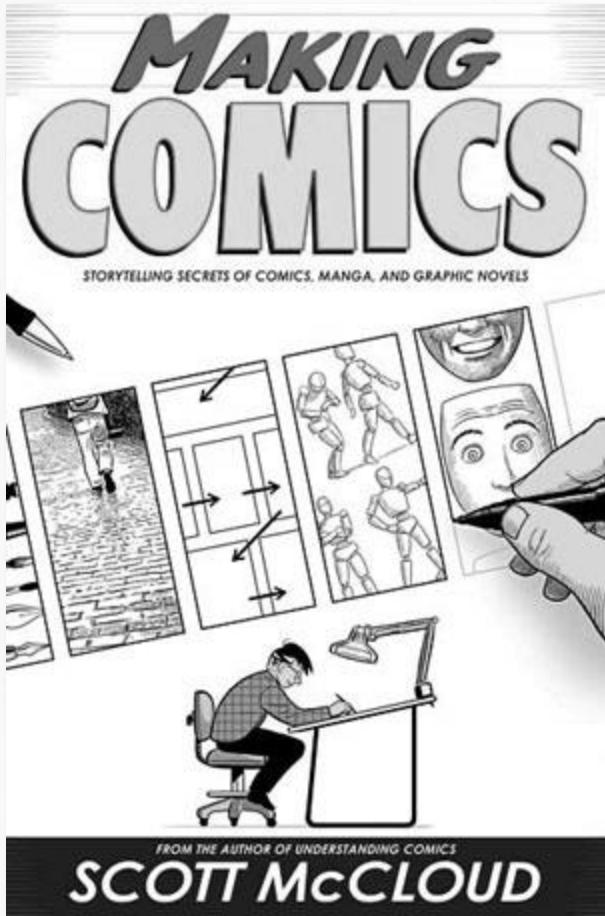
## Full version

- Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- 16 cells \* 8 orientations = 128 dimensional descriptor

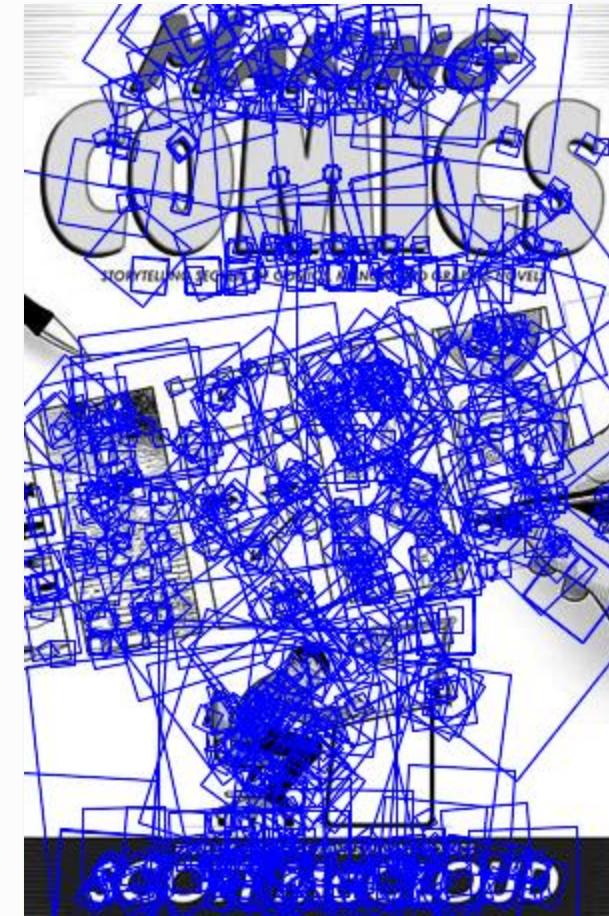




# SIFT EXAMPLE

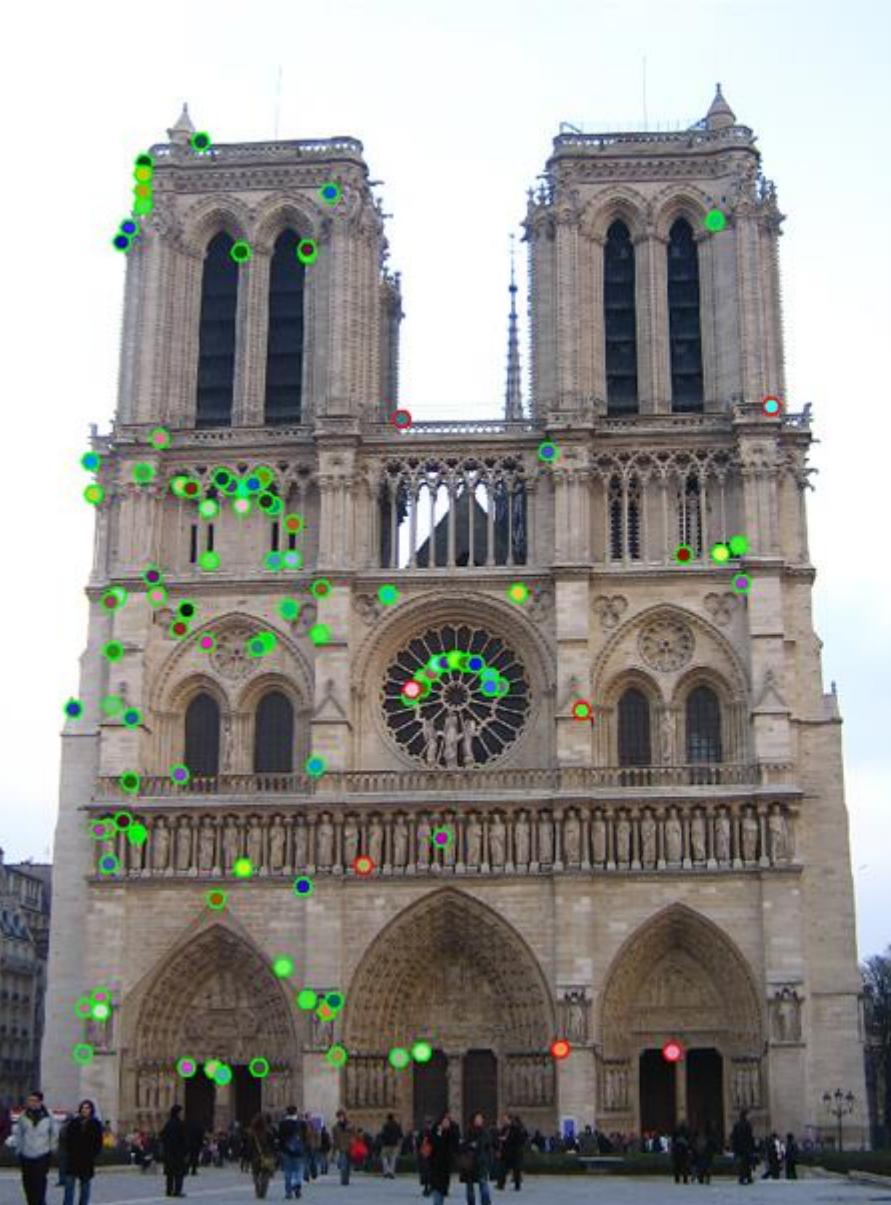
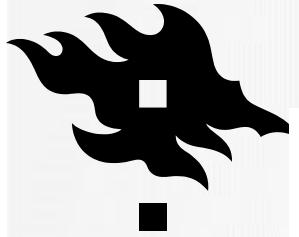


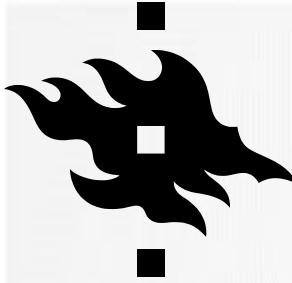
sift



868 SIFT features

# ■ WHICH FEATURES MATCH?

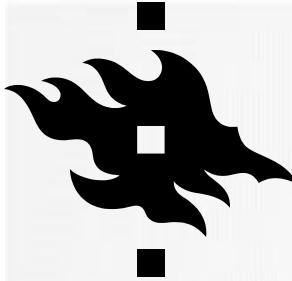




# FEATURE MATCHING



- Given a feature in Image ( $I_0$ ), how to find the best match in ( $I_1$ )
  1. Define distance function that compares two descriptors
  2. Test all the features in  $I_1$ , find the one with min distance



# FEATURE DISTANCE

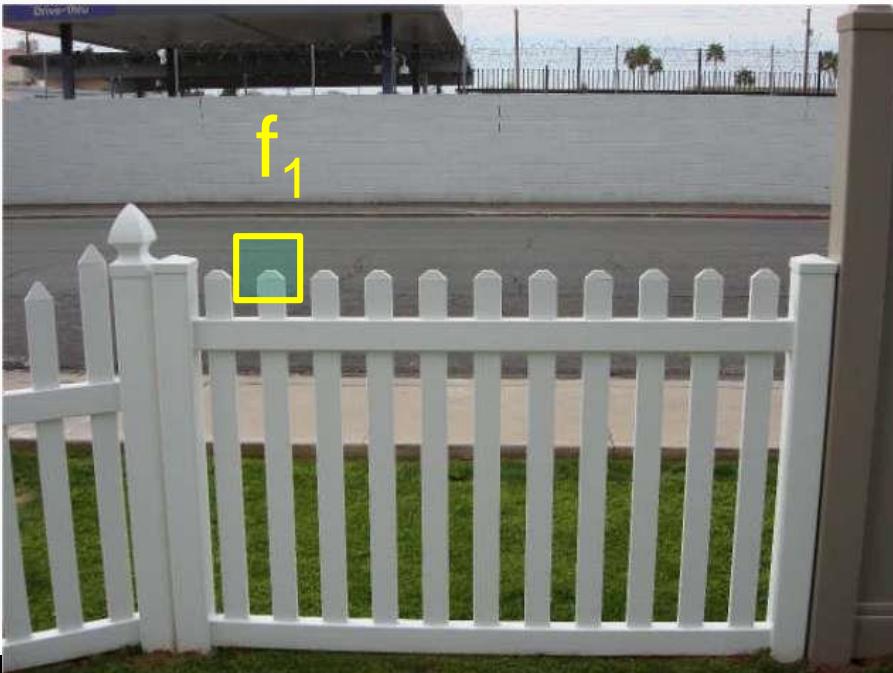


How to define the difference between two features  $f_0, f_1$ ?

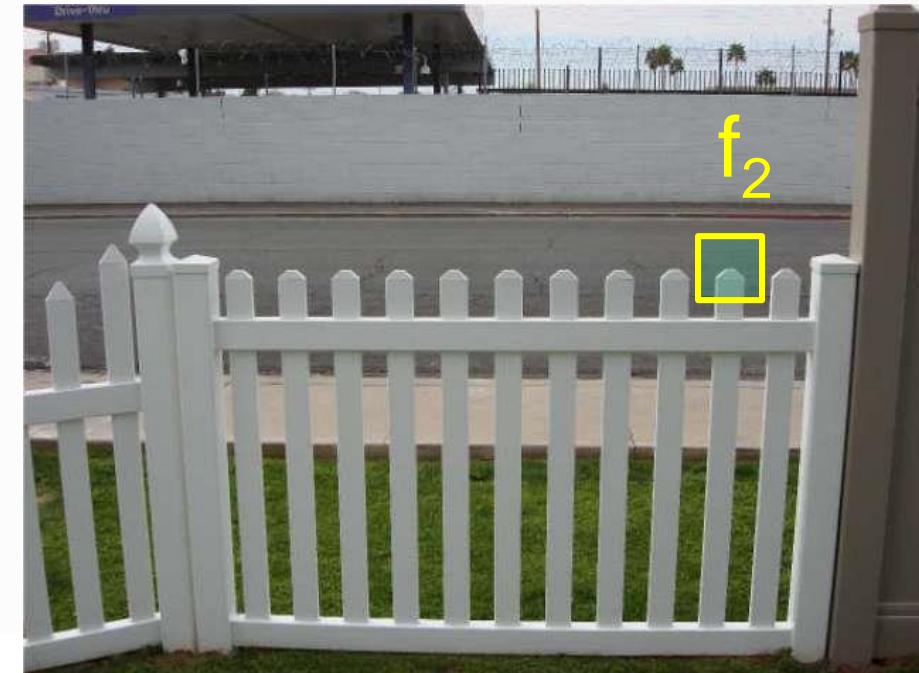
Simple approach: L<sub>2</sub> distance,  $\|f_1 - f_2\|$

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2}.$$

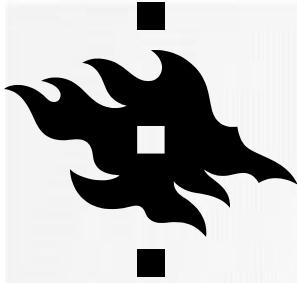
can give small distances for ambiguous (incorrect) matches



$I_1$

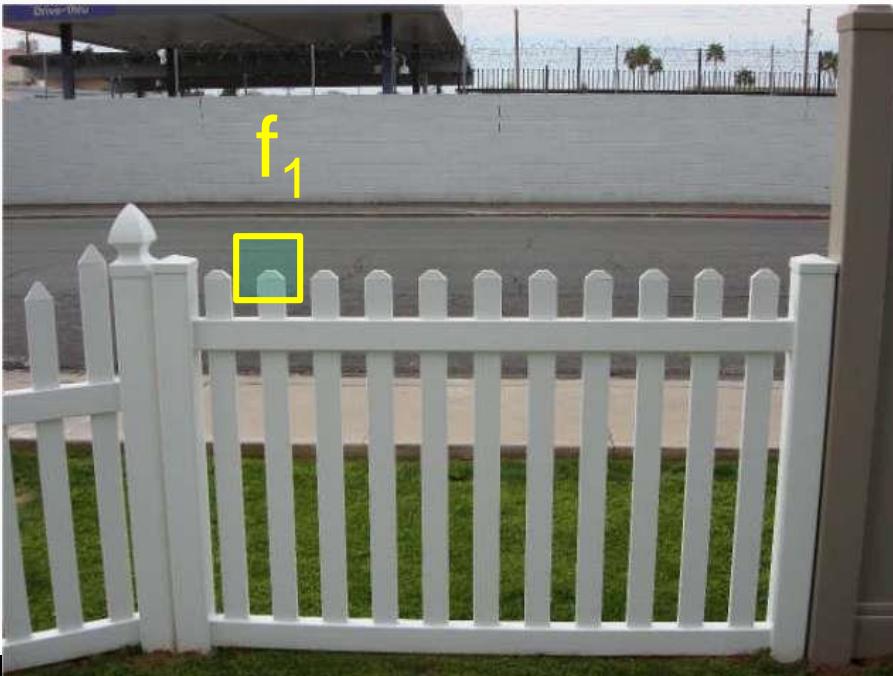


$I_2$



# Feature distance

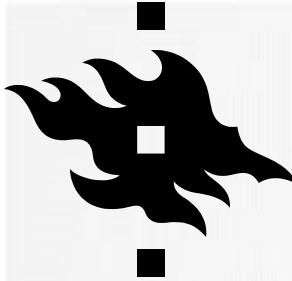
- Better approach: ratio distance =  $\|f_1 - f_2\| / \|f_1 - f_2'\|$ 
  - $f_2$  is best SSD match to  $f_1$  in  $I_2$
  - $f_2'$  is 2<sup>nd</sup> best SSD match to  $f_1$  in  $I_2$
  - gives large values for ambiguous matches



$I_1$



$I_2$  Slide credit: Kris Kitani

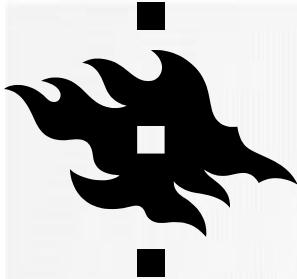


# FEATURE MATCHING EXAMPLE 1

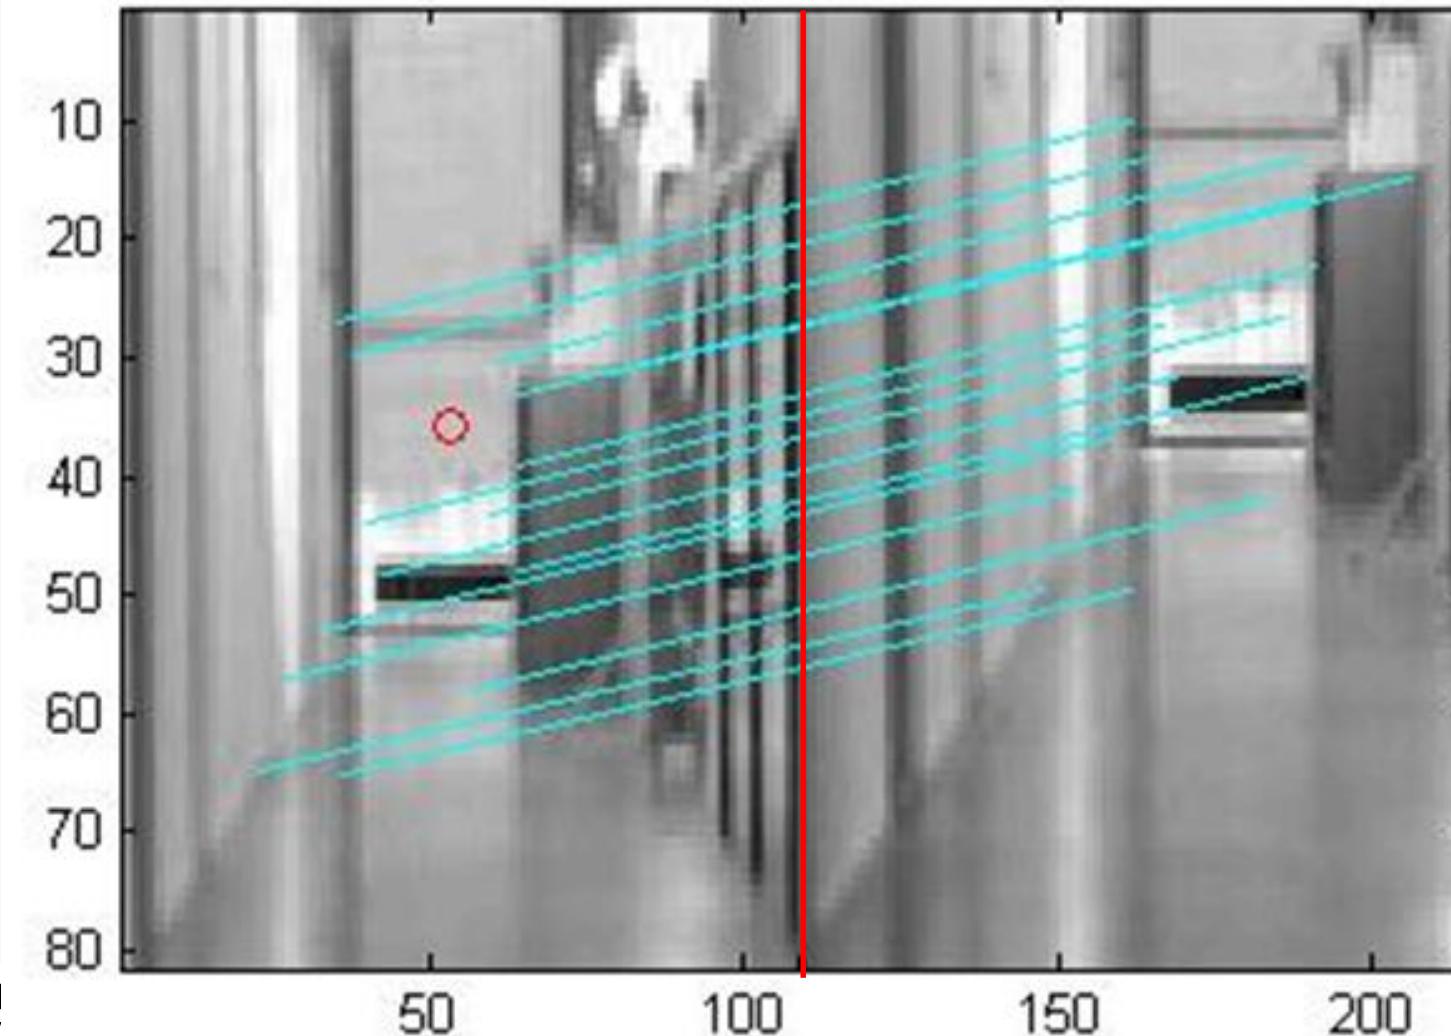
We'll deal  
with  
**outliers**  
later



51 matches (thresholded by ratio score)

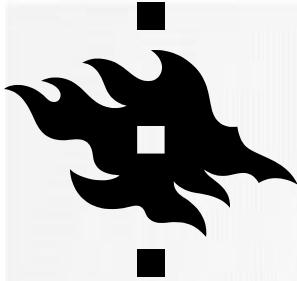


## FEATURE MATCHING EXAMPLE 2



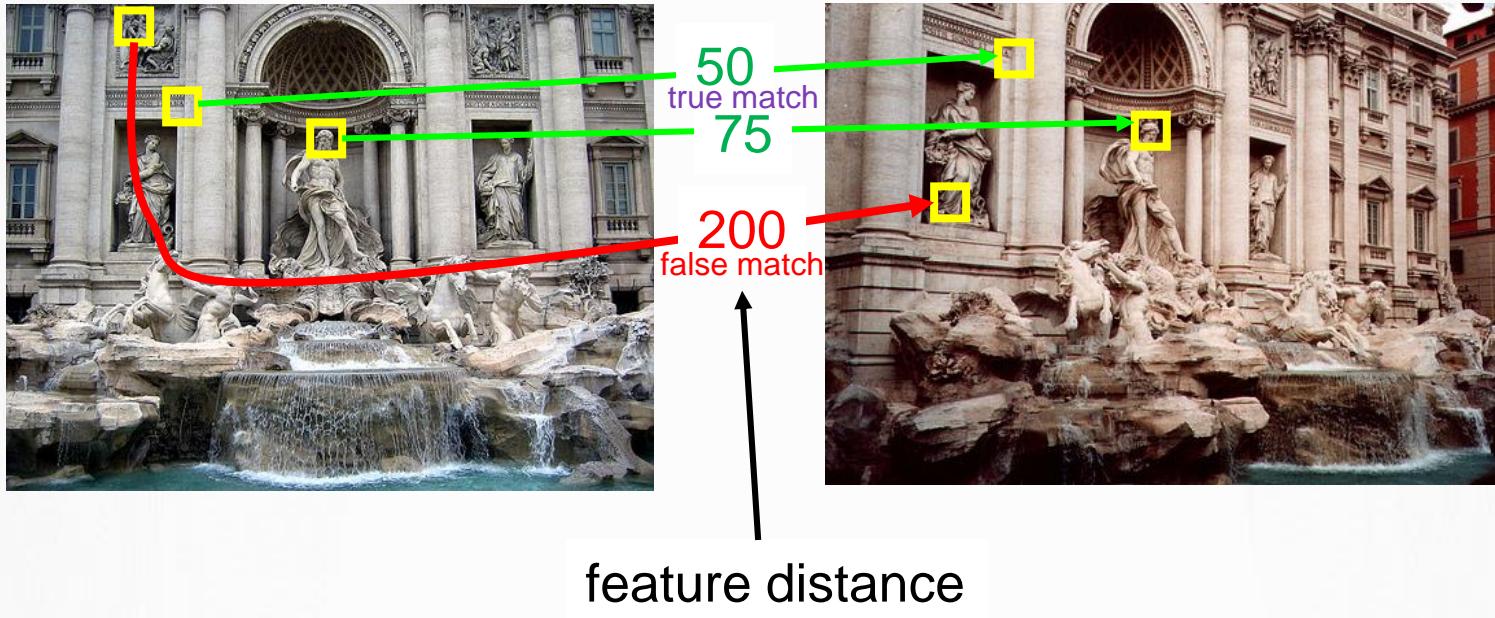
< 20 features  
matched,

Computer vision  
methods are  
very often tested  
and published  
by using  
favorable  
images



# TRUE/FALSE POSITIVES

How can we measure the performance of a feature matcher?



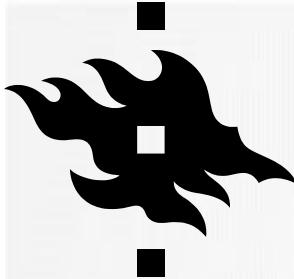
The distance threshold affects performance

True positives = # of detected matches that are correct

- Suppose we want to maximize these—how to choose threshold?

False positives = # of detected matches that are incorrect

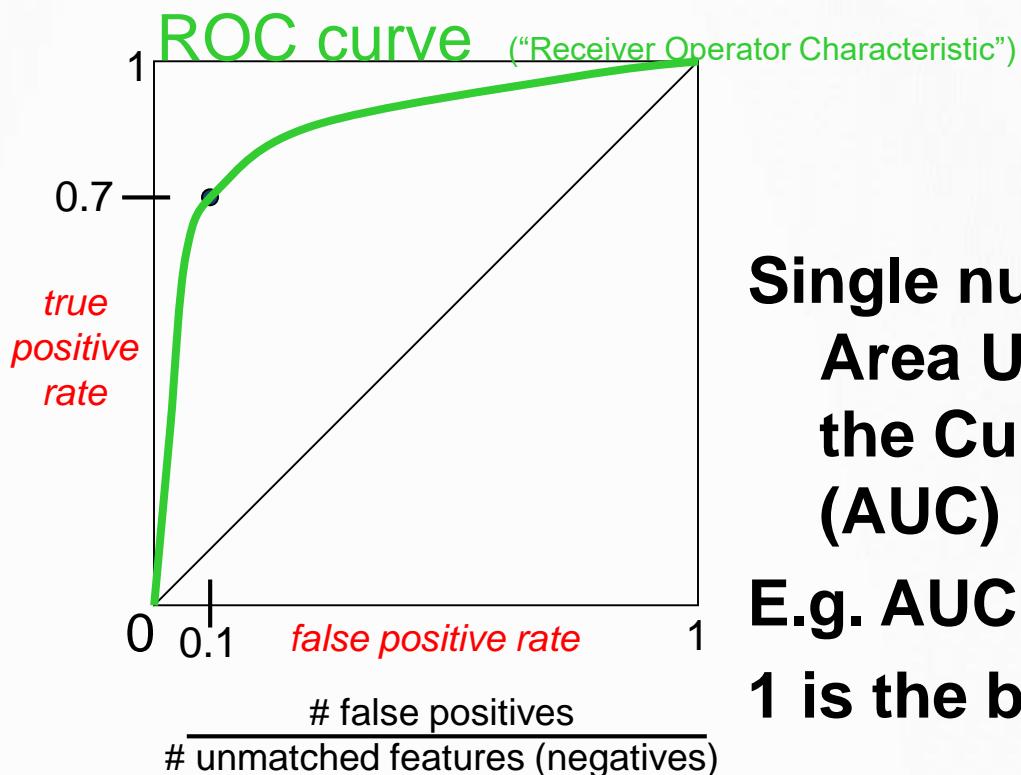
- Suppose we want to minimize these—how to choose threshold? Slide credit: Kris Kitani



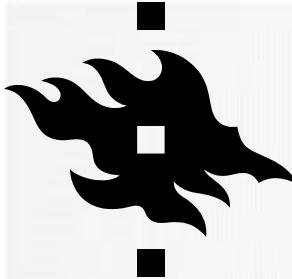
# EVALUATING THE RESULTS

How can we measure the performance of a feature matcher?

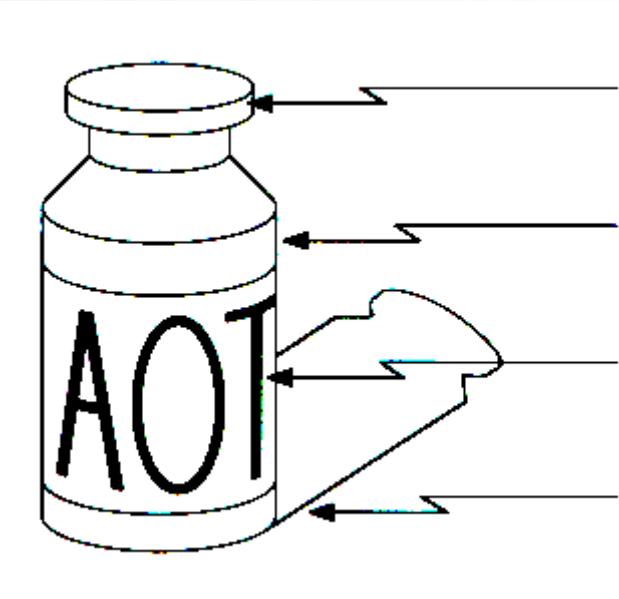
$$\text{recall} = \frac{\# \text{ true positives}}{\# \text{ matching features (positives)}}$$



**Single number:  
Area Under  
the Curve  
(AUC)**  
E.g. AUC = 0.87  
**1 is the best**



# EDGE DETECTION



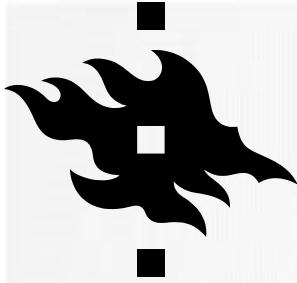
surface normal discontinuity

depth discontinuity

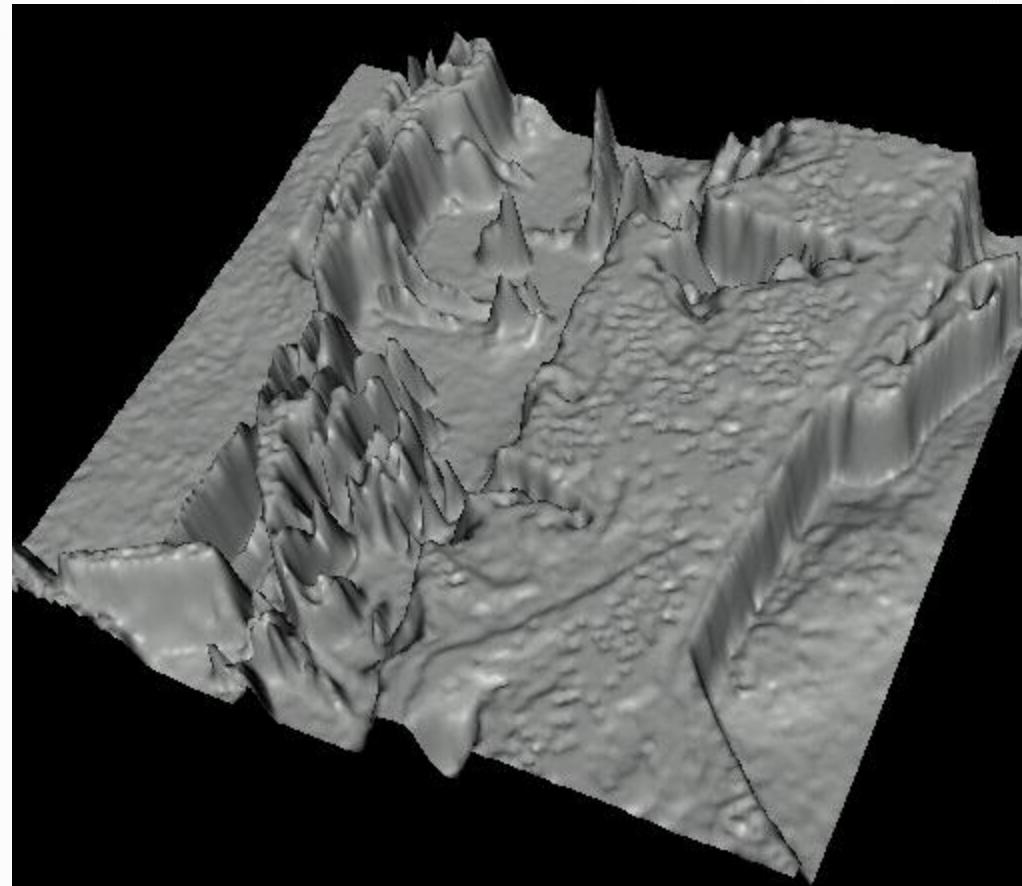
surface color discontinuity

illumination discontinuity

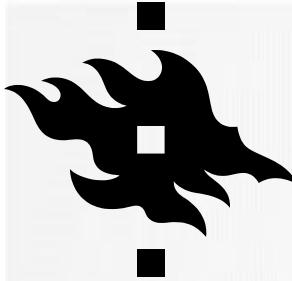
Edges are caused by a variety of factors



# IMAGES AS FUNCTIONS...

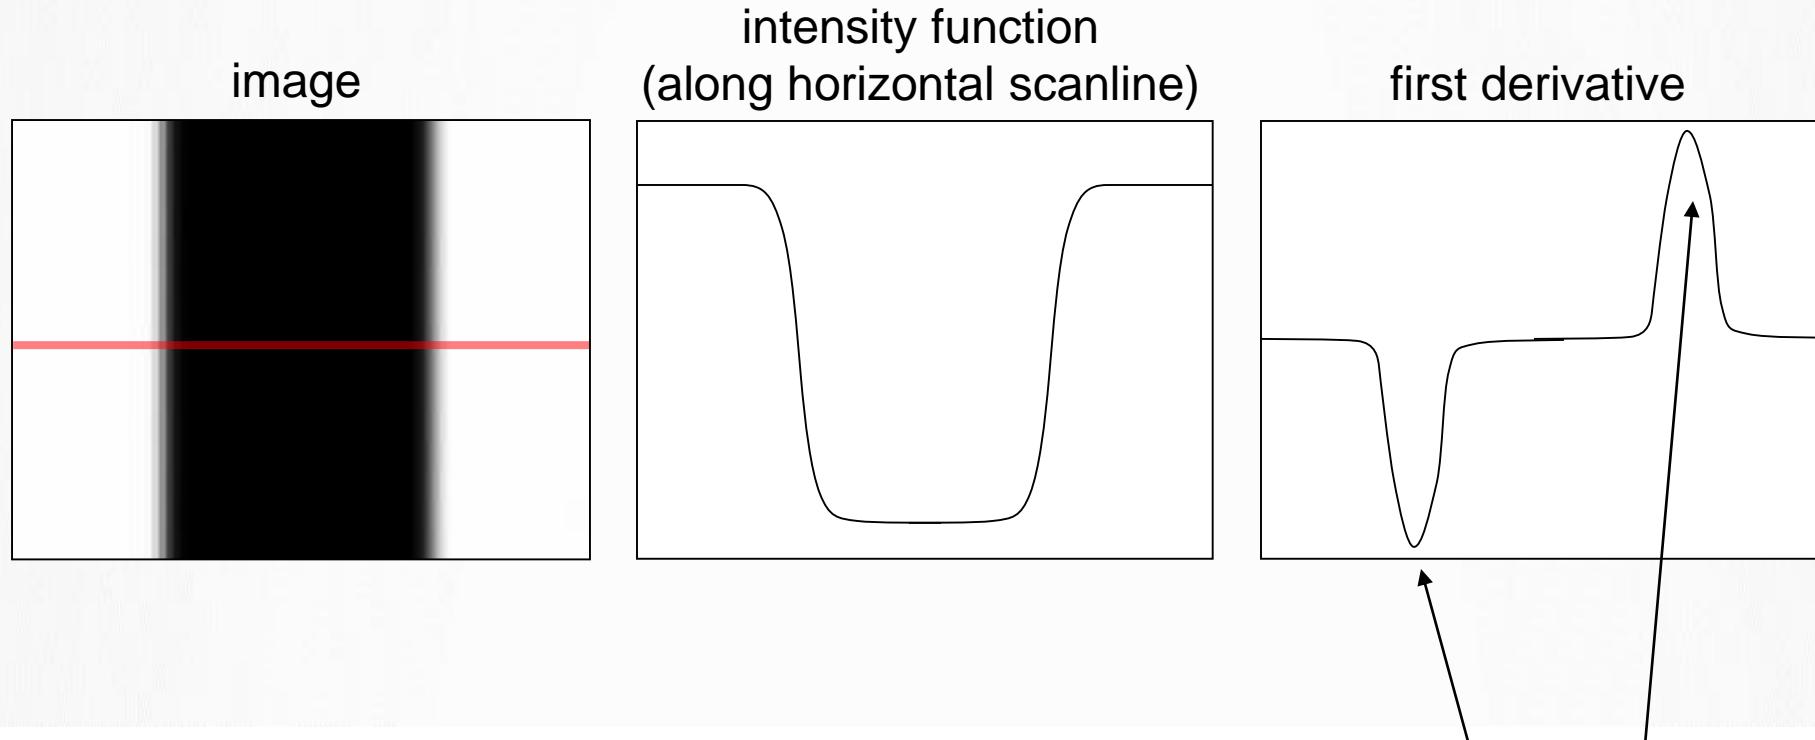


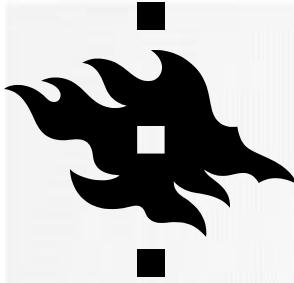
Edges look like steep cliffs



# CHARACTERIZING EDGES

- An edge is a place of *rapid change* in the image intensity function





# IMAGE GRADIENT

The *gradient* of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

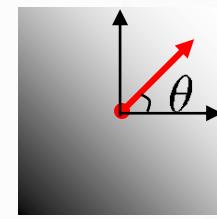
The gradient points in the direction of most rapid increase in intensity



$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$



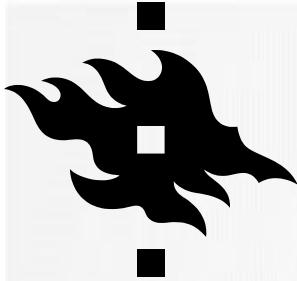
$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



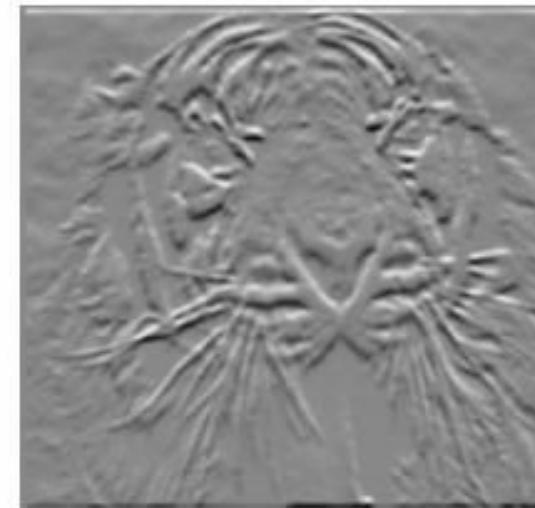
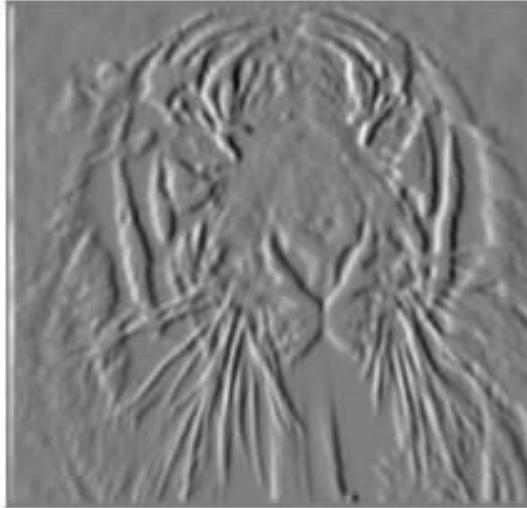
The *edge strength* is given by the gradient magnitude:

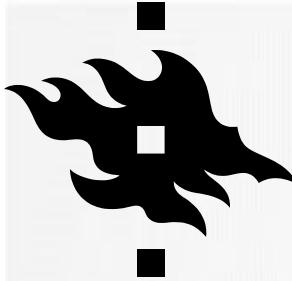
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

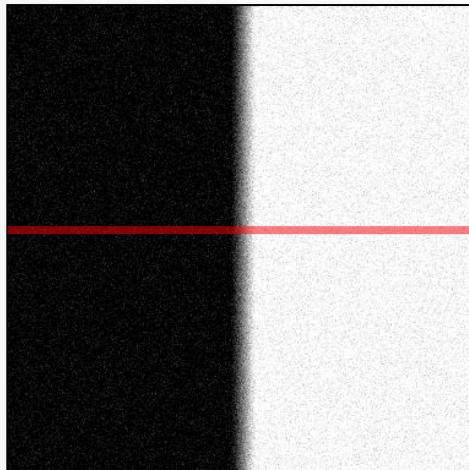


# IMAGE GRADIENT

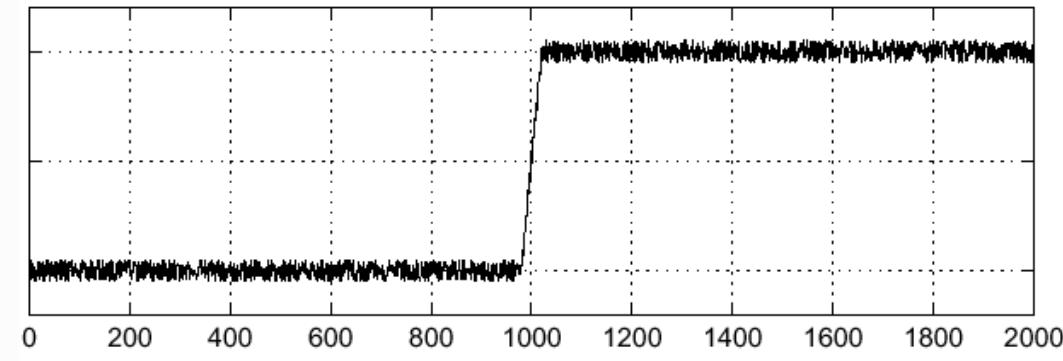




# EFFECTS OF NOISE

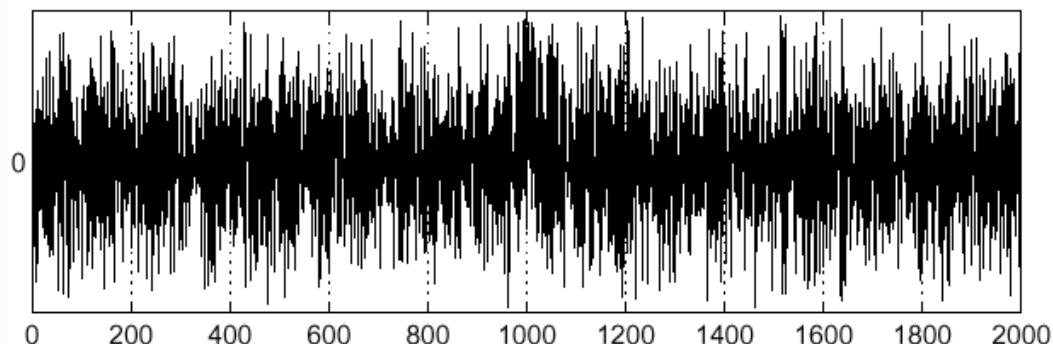


$$f(x)$$

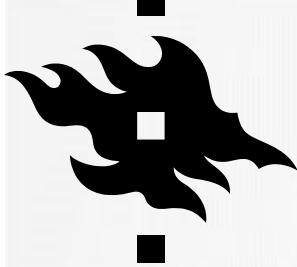


Noisy input image

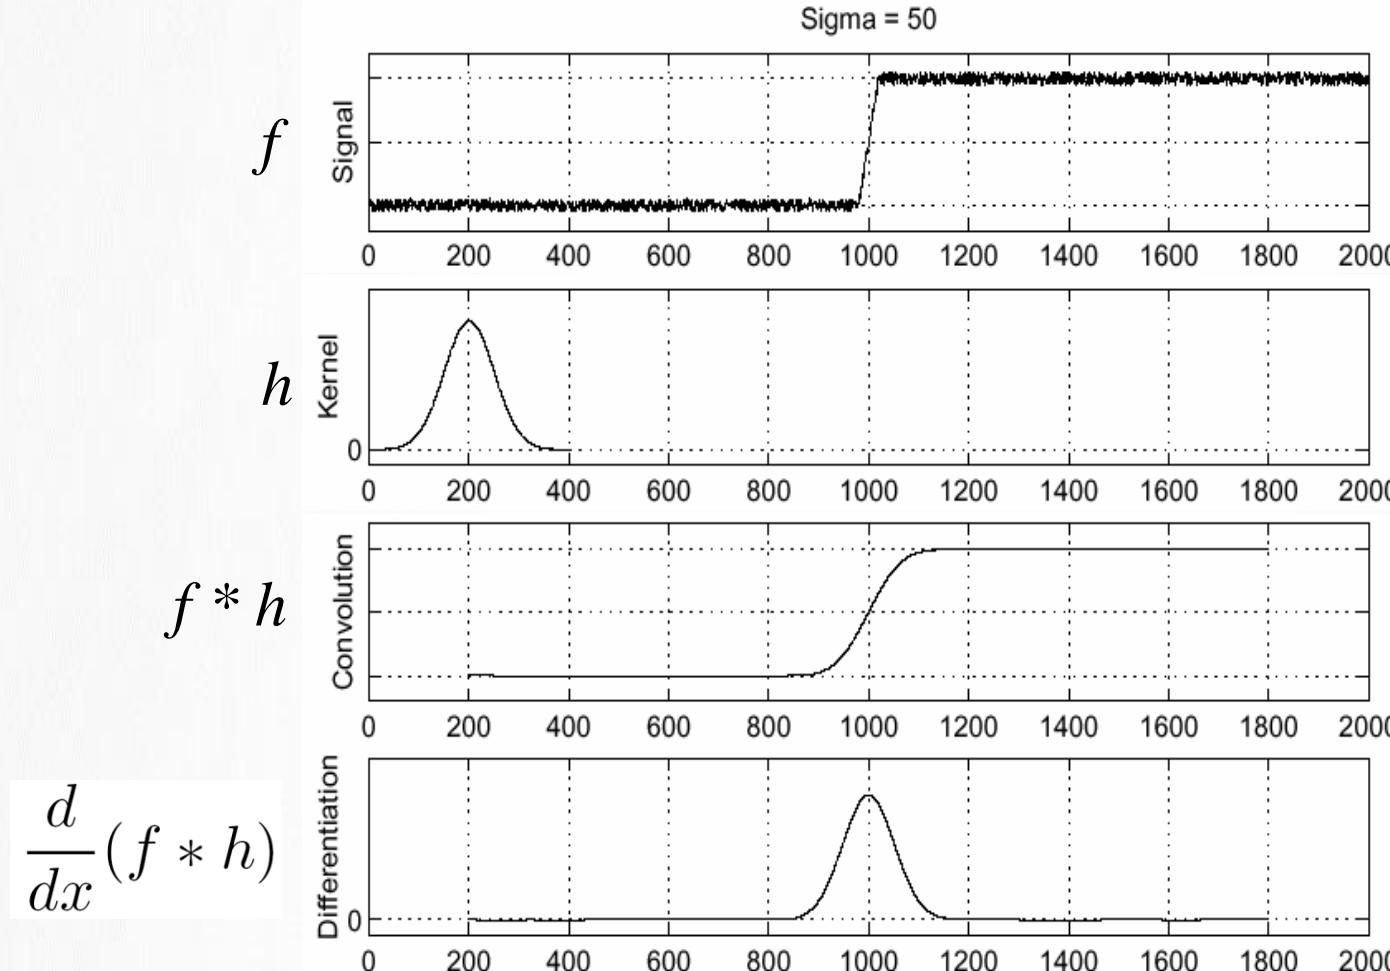
$$\frac{d}{dx}f(x)$$



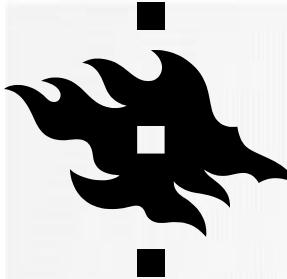
Where is the edge?



# SOLUTION: SMOOTH FIRST

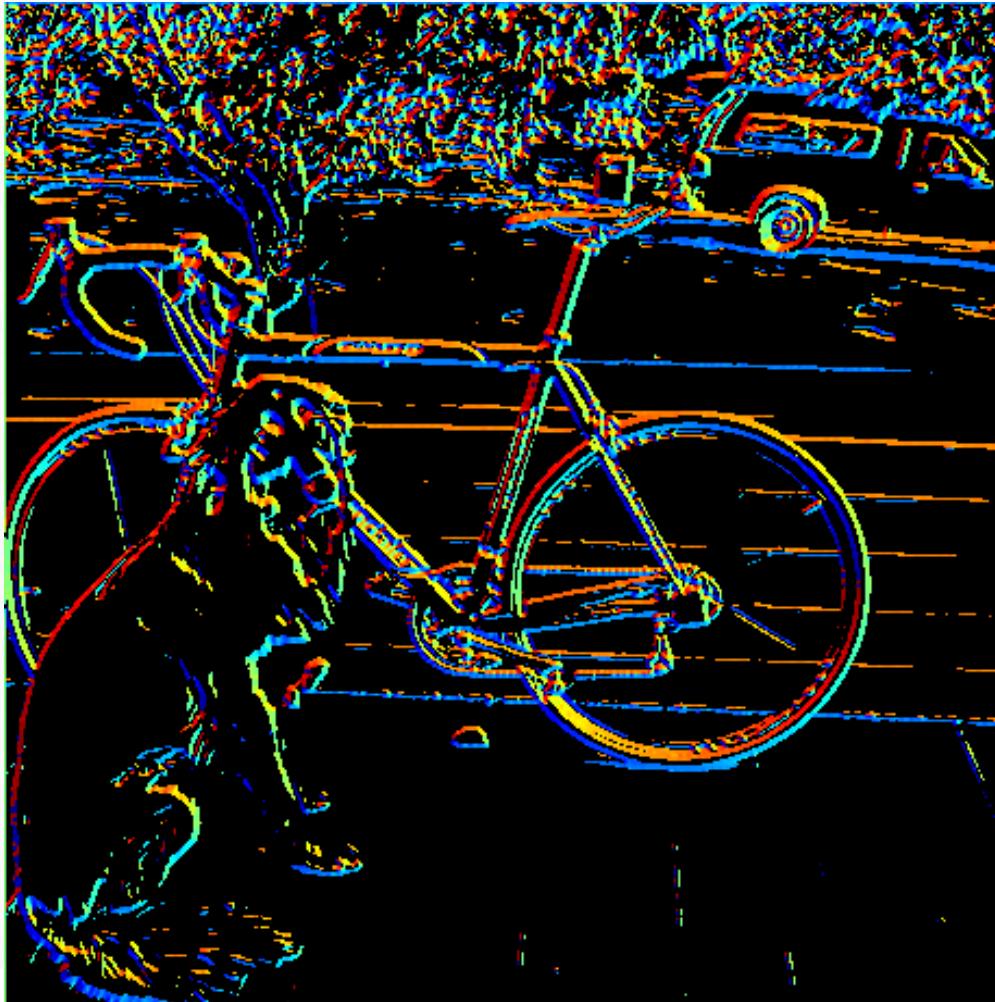


To find edges, look for peaks in  $\frac{d}{dx}(f * h)$



# GET ORIENTATION AT EACH PIXEL

Get orientation (below, threshold at minimum gradient magnitude)

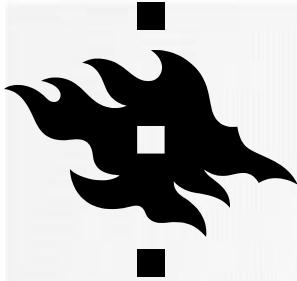


$$\theta = \text{atan2}(g_y, g_x)$$

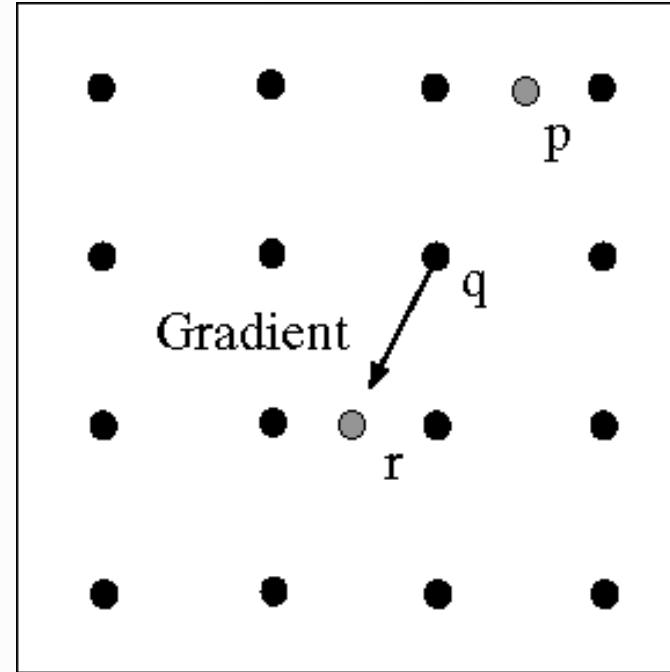
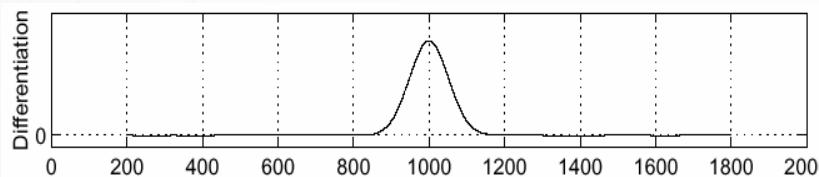
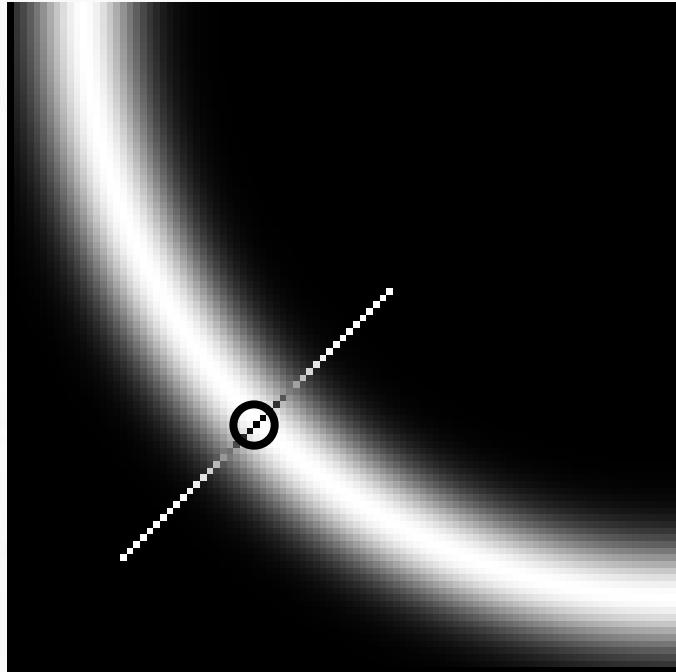
360



**Gradient orientation angle**

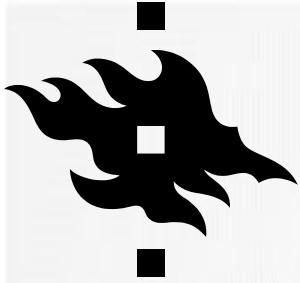


# NON-MAXIMUM SUPPRESSION



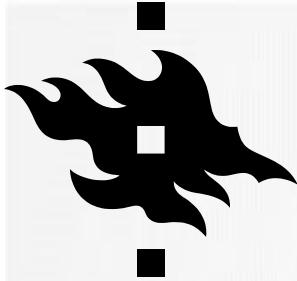
Check if pixel is local maximum along gradient direction

requires *interpolating* pixels p and r



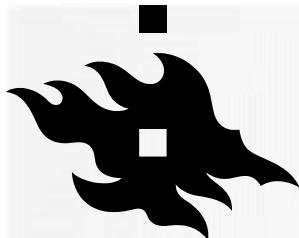
# BEFORE NON-MAX SUPPRESSION





# AFTER NON-MAX SUPPRESSION

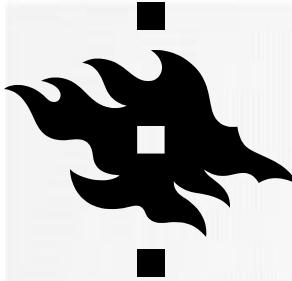




# THRESHOLDING EDGES

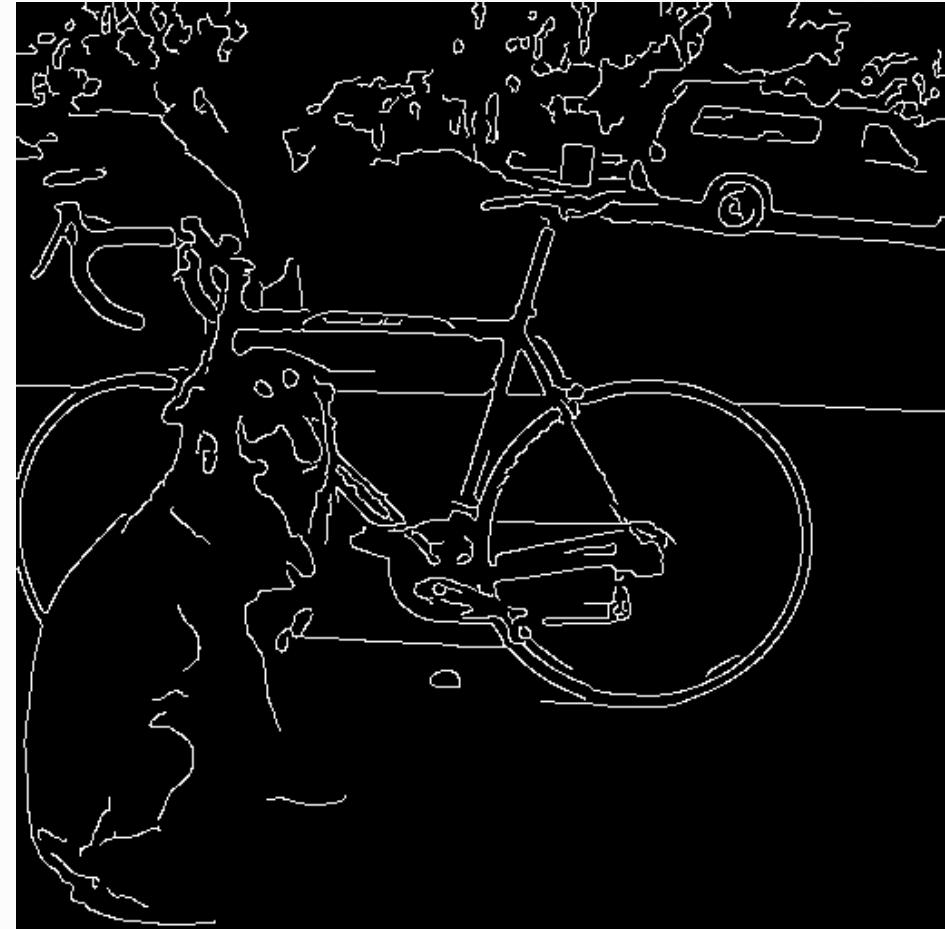
- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
  - $R > T$ : strong edge
  - $R < T$  but  $R > t$ : weak edge
  - $R < t$ : no edge
- Why two thresholds?

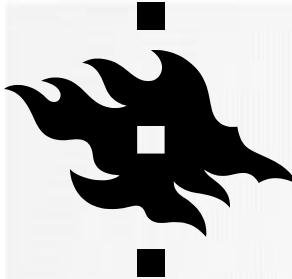




# CONNECTING EDGES

- Strong edges are edges!
- Weak edges are edges iff they connect to strong
- Look in some neighborhood (usually 8 closest)





# CANNY EDGE DETECTOR

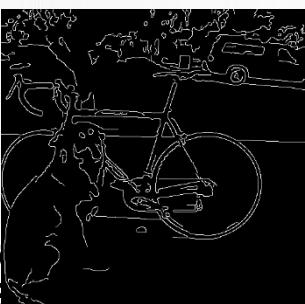
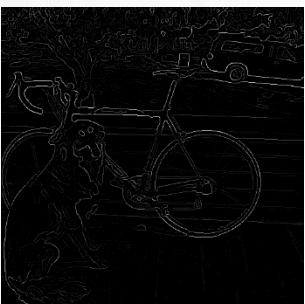
MATLAB: `edge(image, 'canny')`

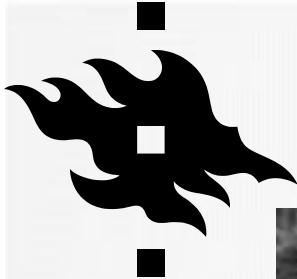


1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression
4. Linking and thresholding (hysteresis):

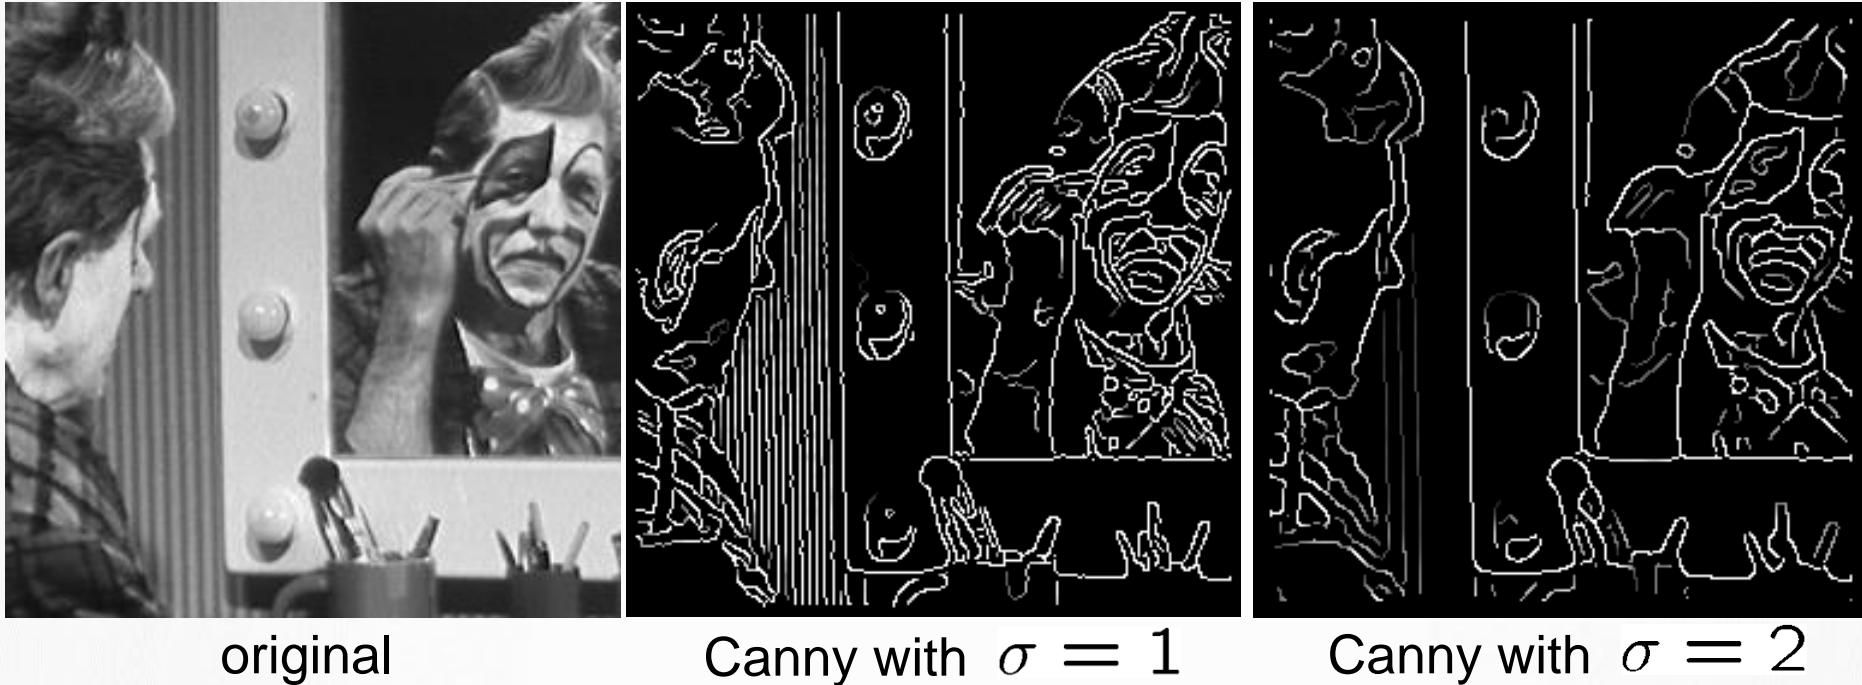
Define two thresholds: low and high

Use the high threshold to start edge curves and the low threshold to continue them



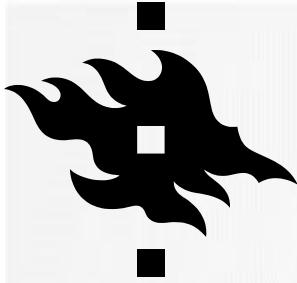


# CANNY EDGE DETECTOR



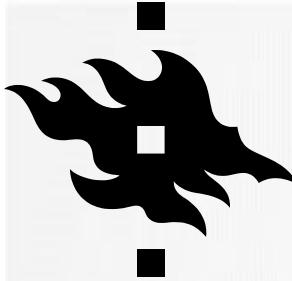
- The choice of  $\sigma$  depends on desired behavior
  - large  $\sigma$  detects “large-scale” edges
  - small  $\sigma$  detects fine edges





# EDGES FROM AN IMAGE



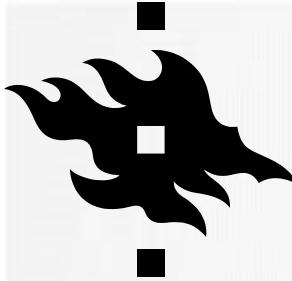


# HOUGH TRANSFORM



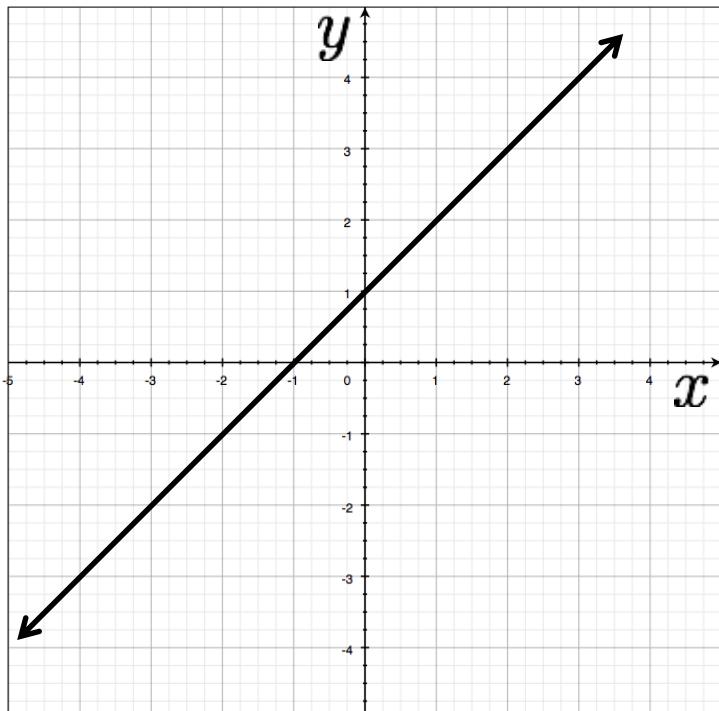
- Detecting lines from all edges
  - Generic framework for detecting a parametric model
  - Edges don't have to be connected
  - Lines can be occluded
  - Key idea: edges **vote** for the possible models





# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters



a line  
becomes  
a point

variables  
 $y - mx = b$   
parameters

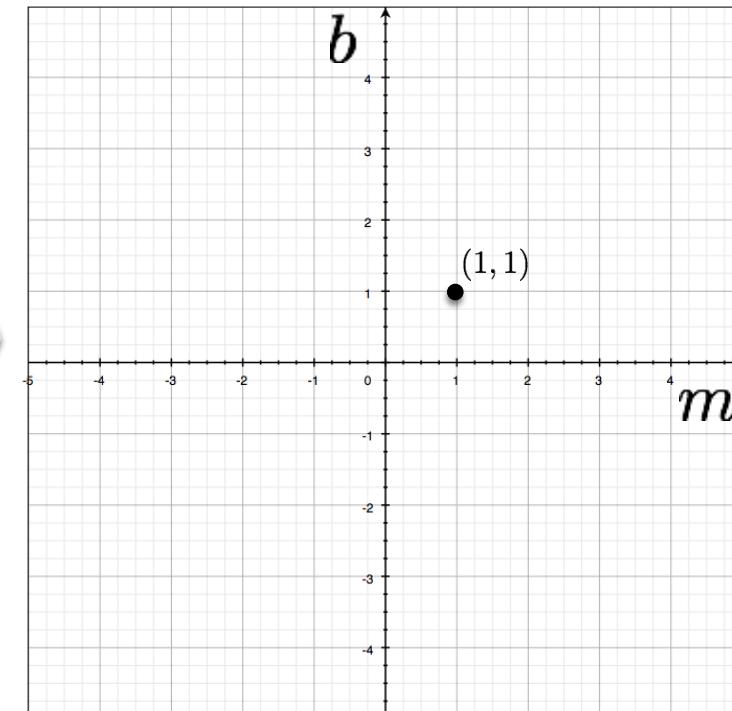
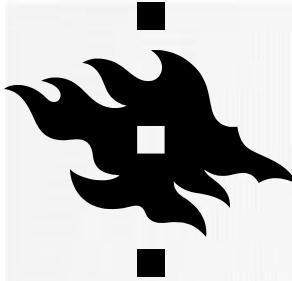


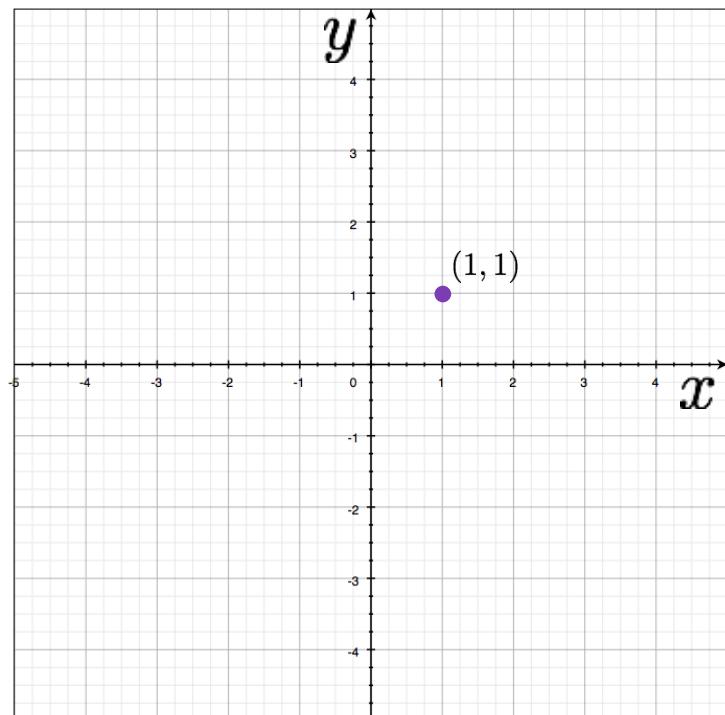
Image space

Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters



a point becomes a line

variables  
 $y - mx = b$   
parameters

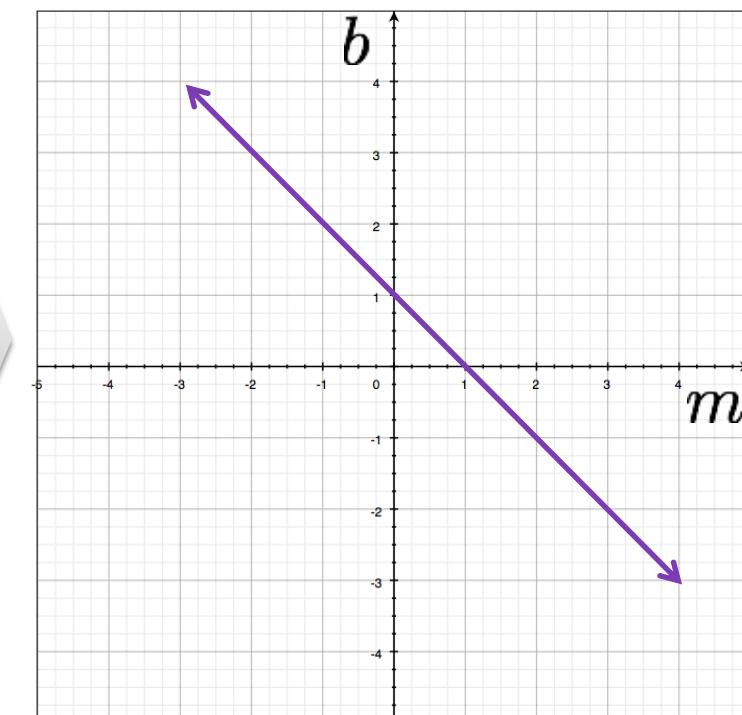
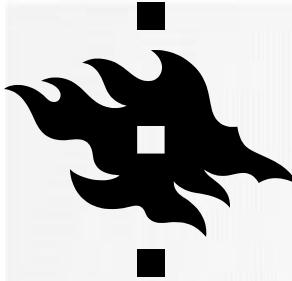


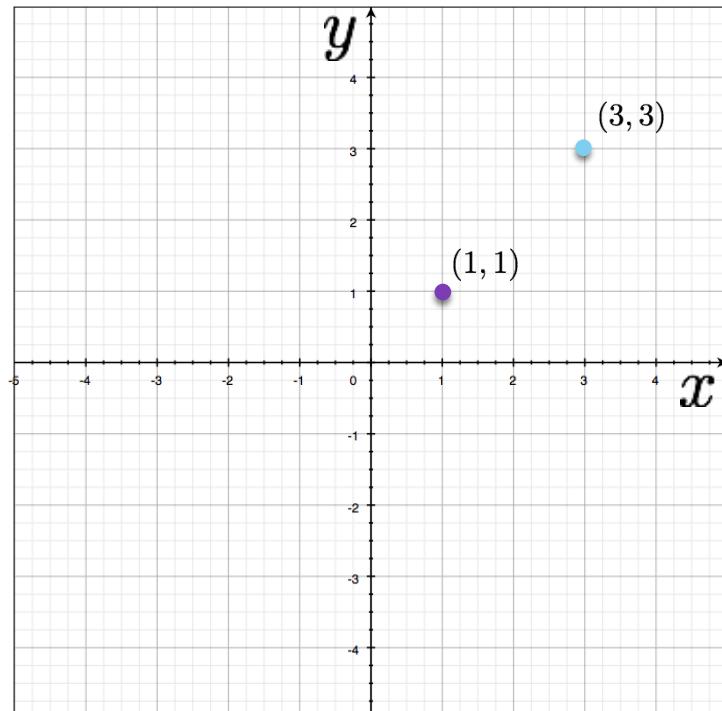
Image space

Parameter space



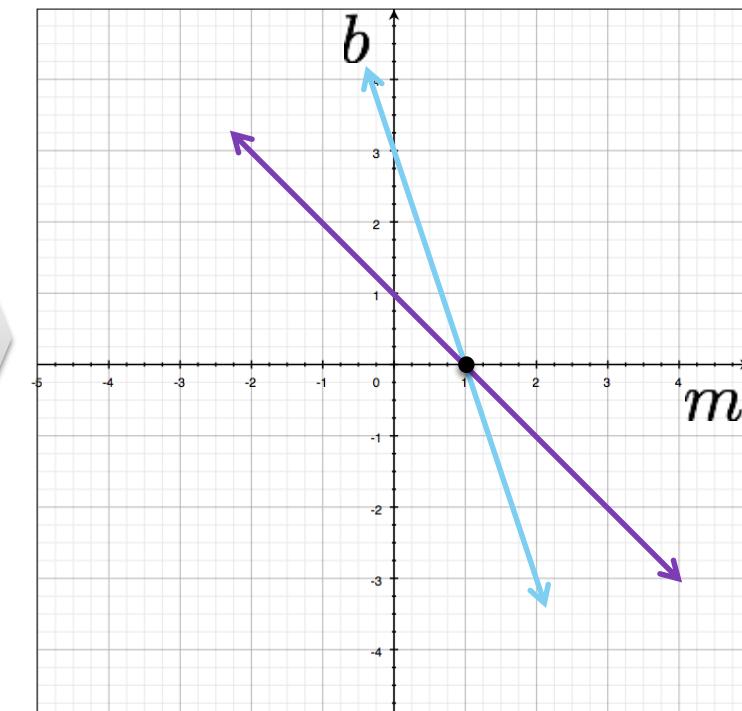
# IMAGE AND PARAMETER SPACE

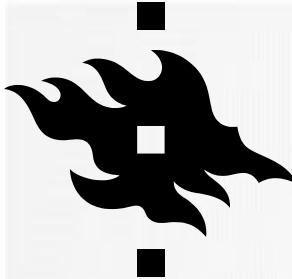
variables  
 $y = mx + b$   
parameters



two points  
become  
?

variables  
 $y - mx = b$   
parameters

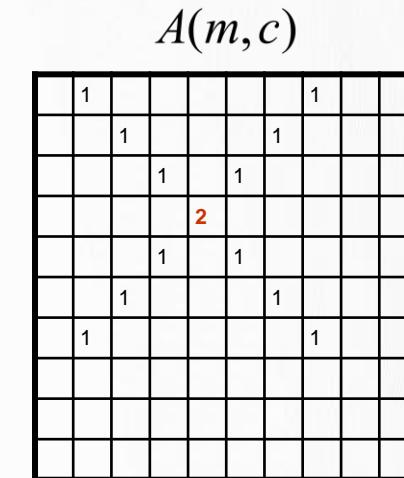
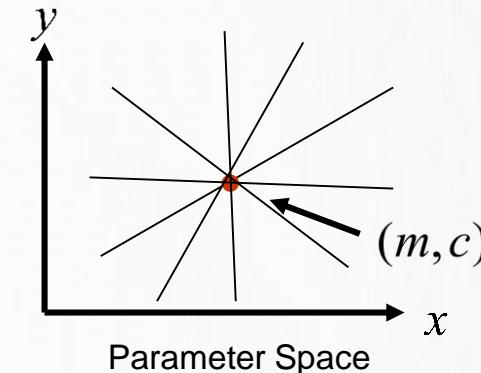


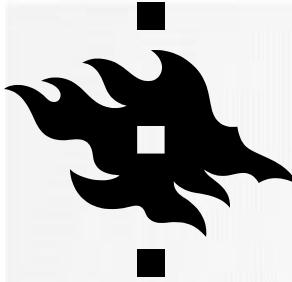


# Line Detection by Hough Transform

## Algorithm:

1. Quantize Parameter Space  $(m, c)$
2. Create Accumulator Array  $A(m, c)$
3. Set  $A(m, c) = 0 \quad \forall m, c$
4. For each image edge  $(x_i, y_i)$   
For each element in  $A(m, c)$   
If  $(m, c)$  lies on the line:  $c = -x_i m + y_i$   
Increment  $A(m, c) = A(m, c) + 1$
5. Find local maxima in  $A(m, c)$



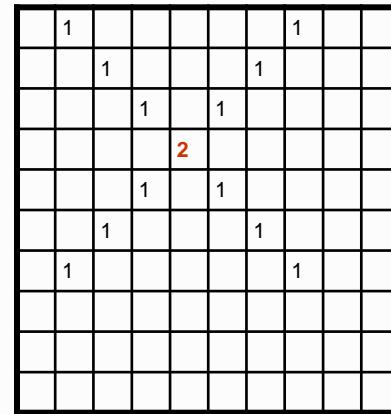


# PROBLEMS WITH PARAMETERIZATION



*How big does the accumulator need to be for the parameterization  $(m, c)$ ?*

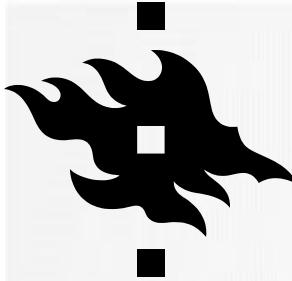
$A(m, c)$



The space of  $m$  is huge!      The space of  $c$  is huge!

$$-\infty \leq m \leq \infty$$

$$-\infty \leq C \leq \infty$$



# BETTER PARAMETERIZATION

Use polar coordinates:

$$x \cos \theta + y \sin \theta = \rho$$

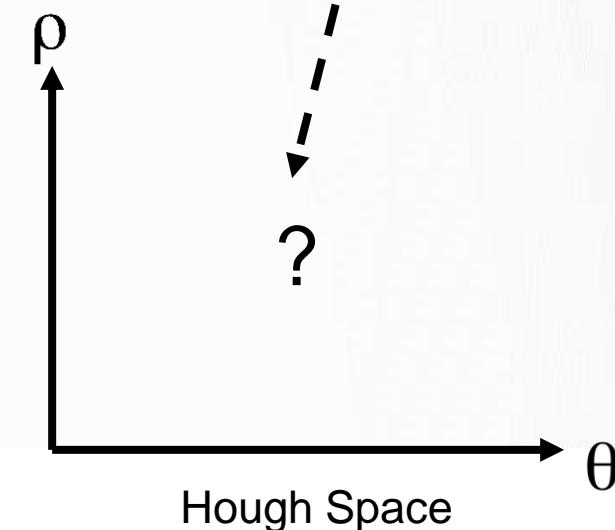
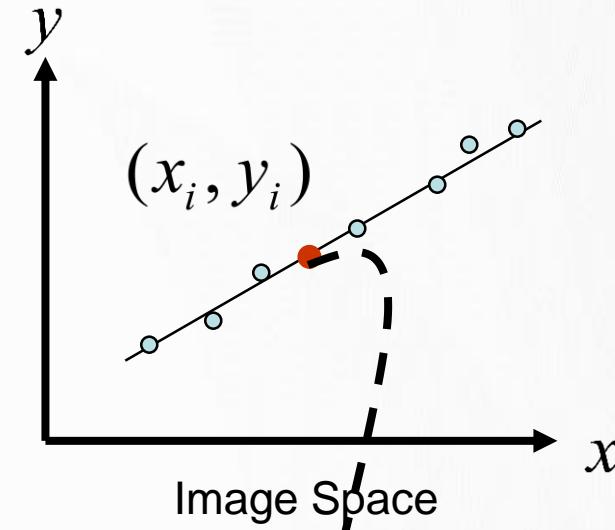
Given points  $(x_i, y_i)$  find  $(\rho, \theta)$

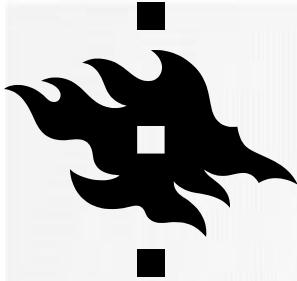
Hough Space Sinusoid

$$0 \leq \theta \leq 2\pi$$

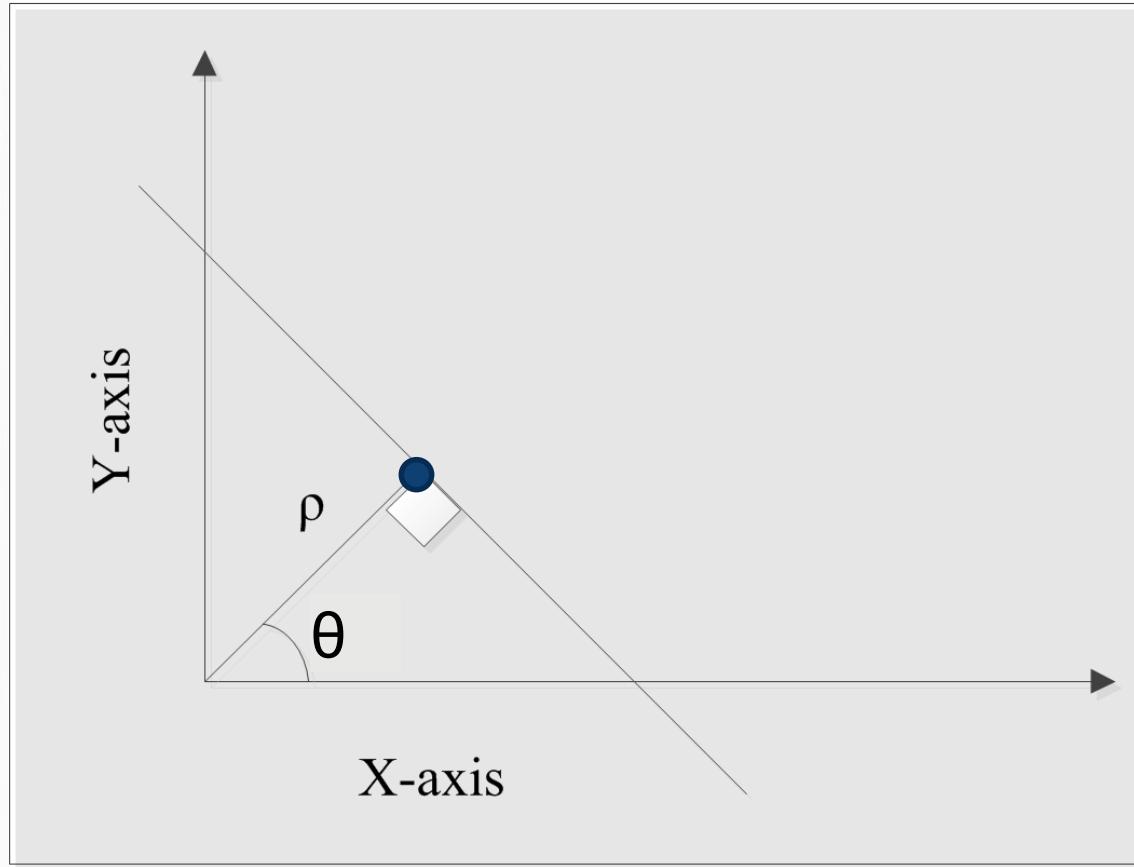
$$0 \leq \rho \leq \rho_{\max}$$

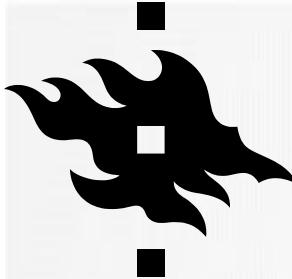
(Finite Accumulator Array Size)





# POLAR COORDINATES

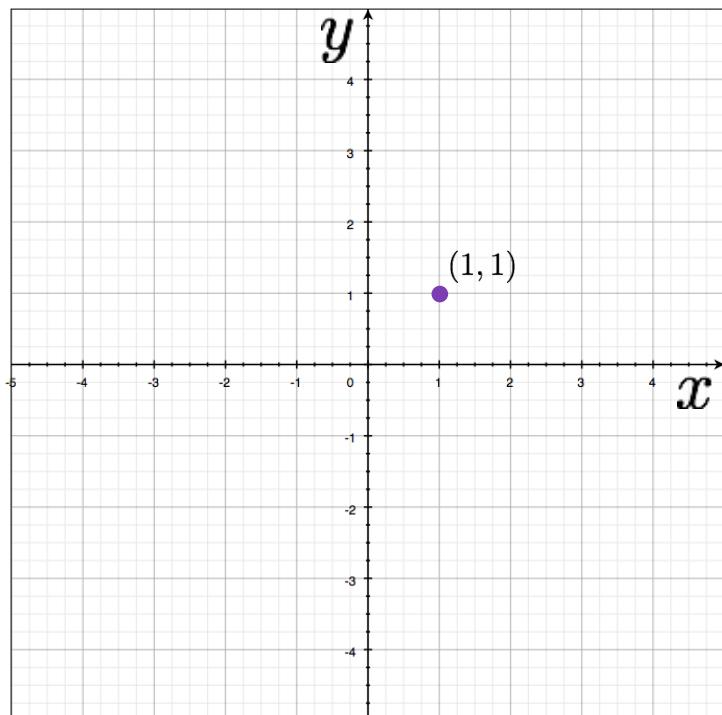




# IMAGE AND PARAMETER SPACE

$$y = mx + b$$

variables  
parameters

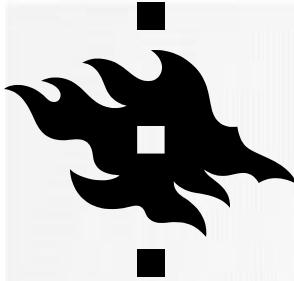


a point becomes a wave

$$x \cos \theta + y \sin \theta = \rho$$

parameters  
variables





# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

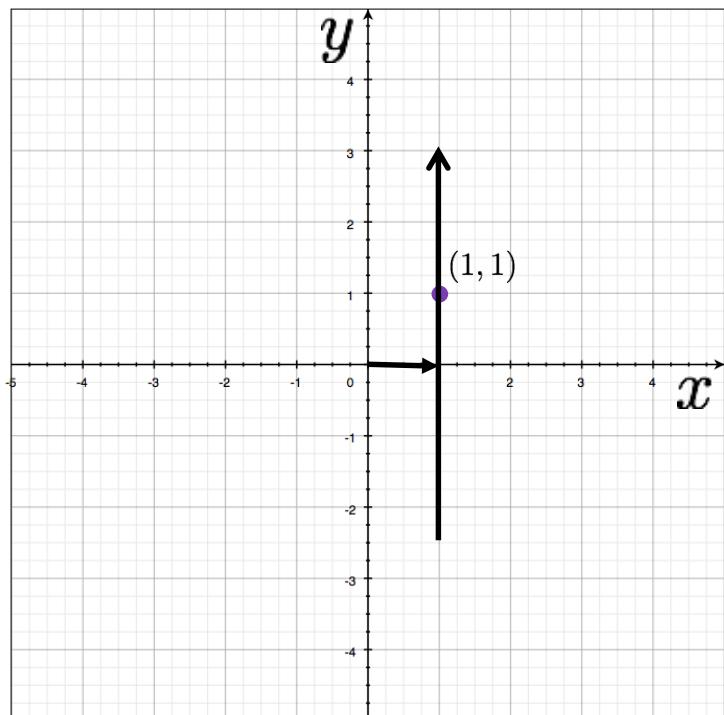
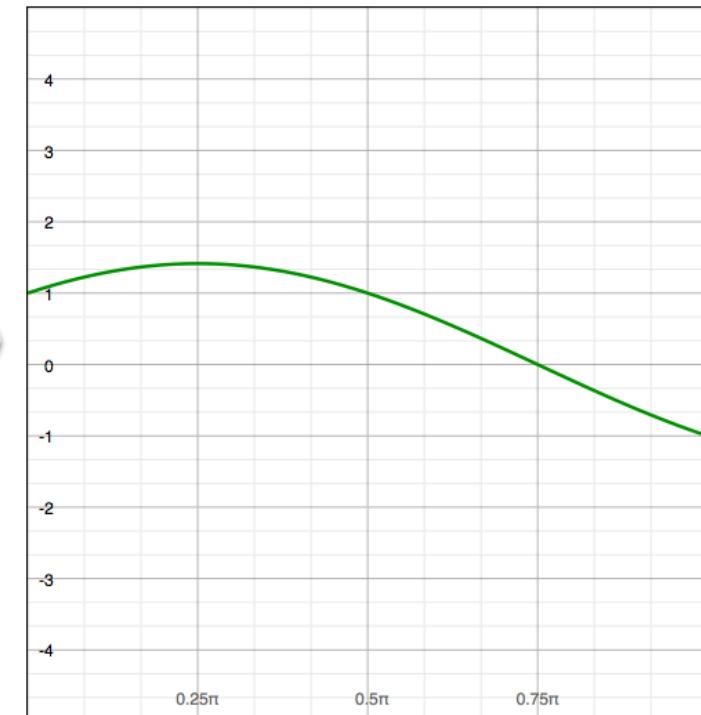


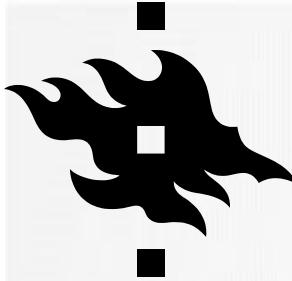
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line  
becomes?



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

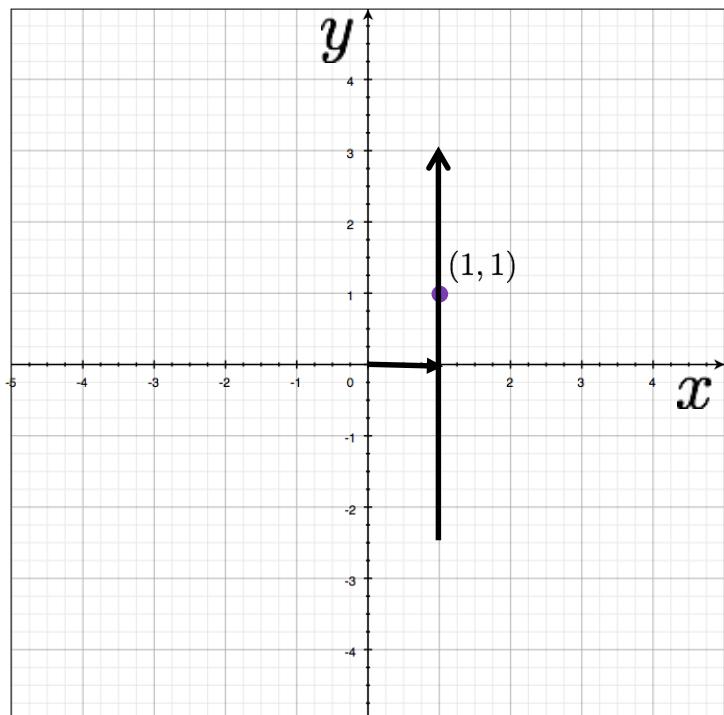
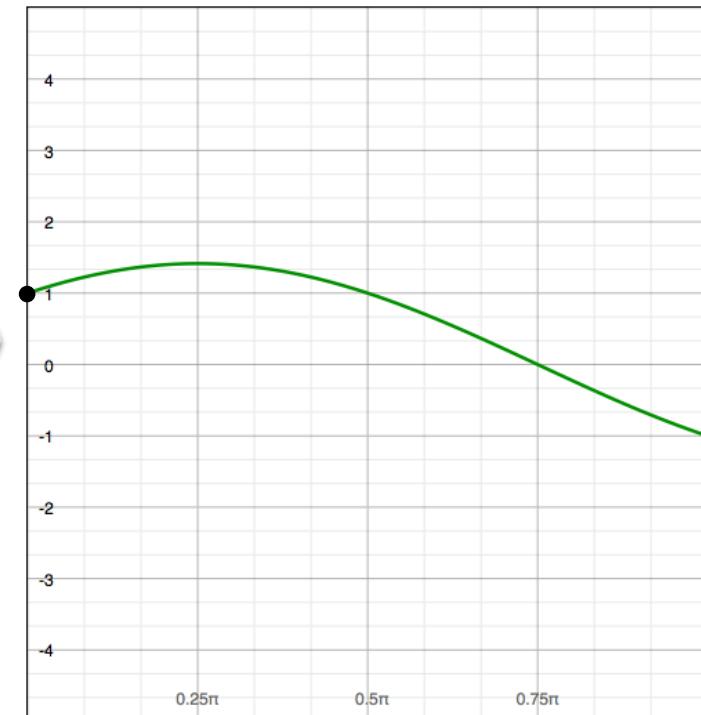


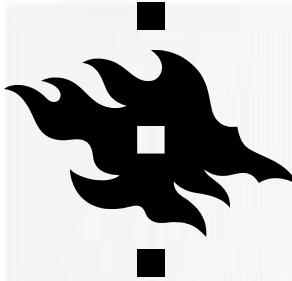
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line  
becomes  
a point

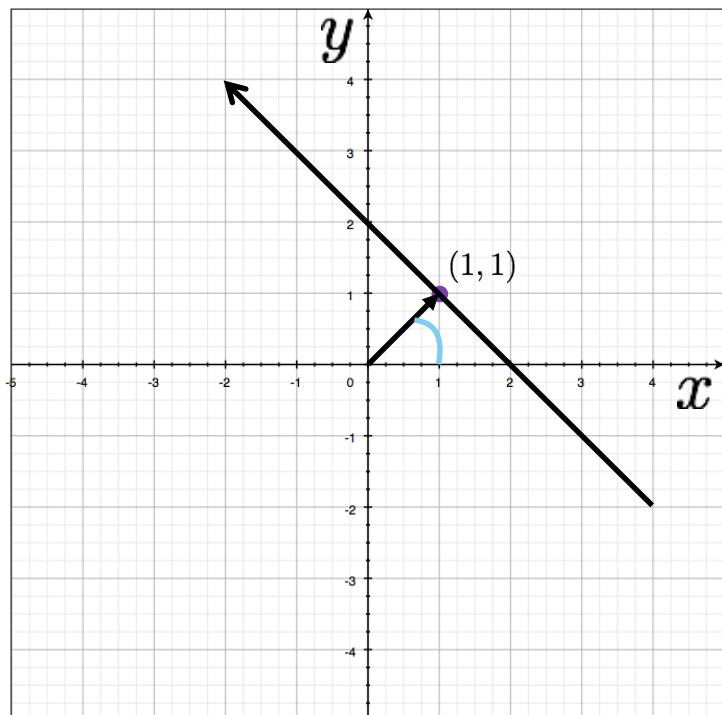


Parameter space



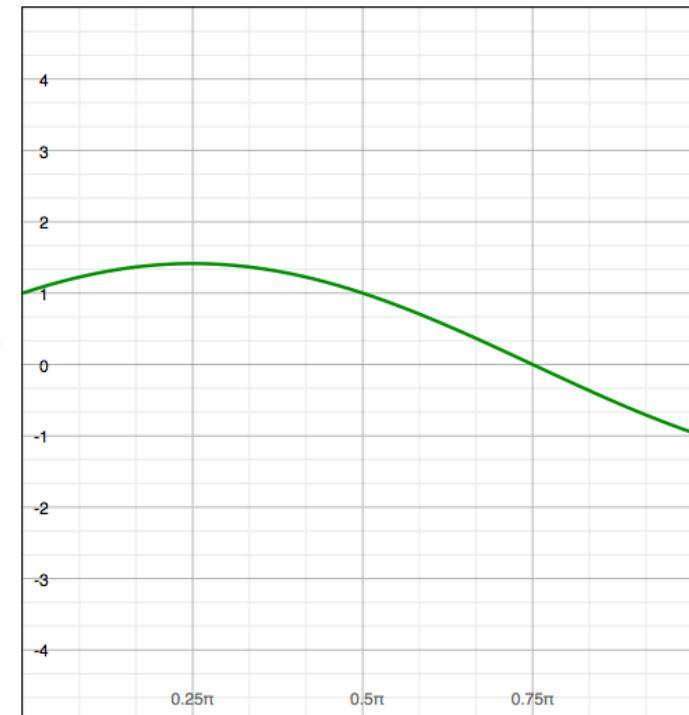
# IMAGE AND PARAMETER SPACE

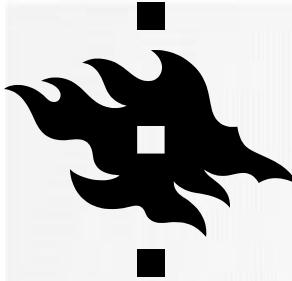
variables  
 $y = mx + b$   
parameters



a line  
becomes?

$$x \cos \theta + y \sin \theta = \rho$$





# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

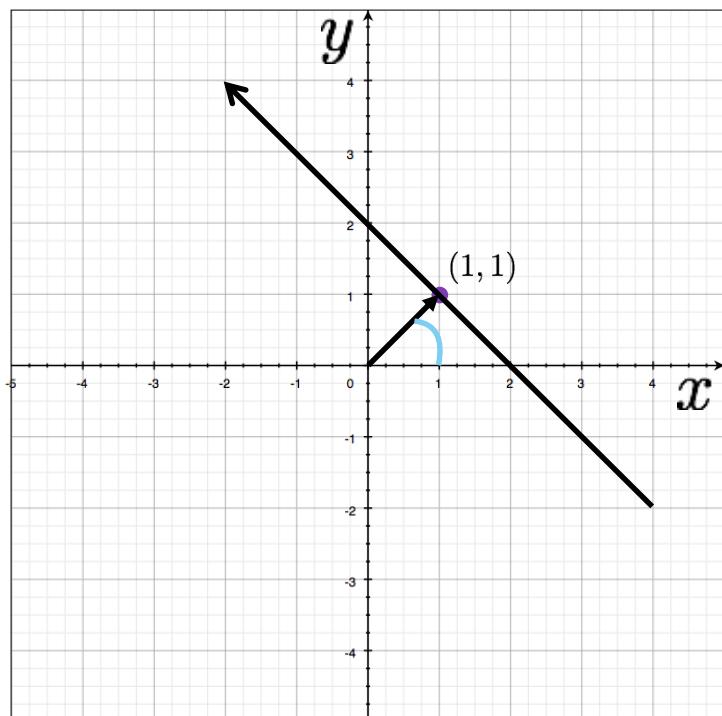
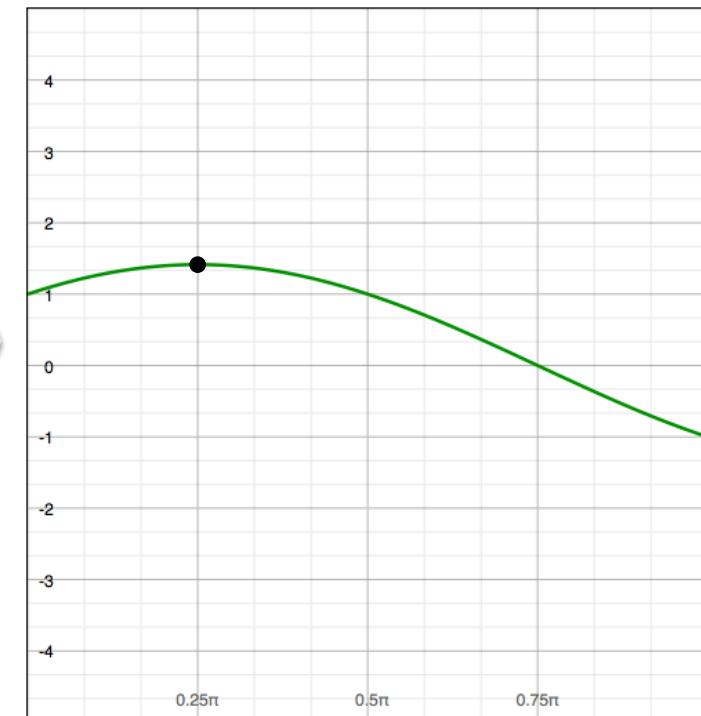


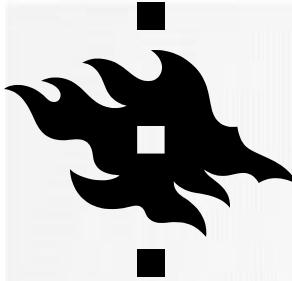
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line becomes a point



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

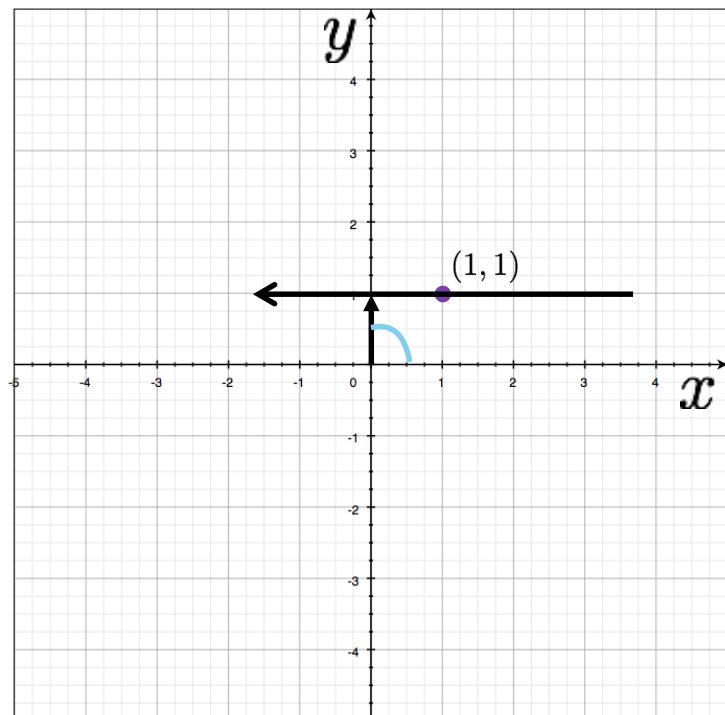
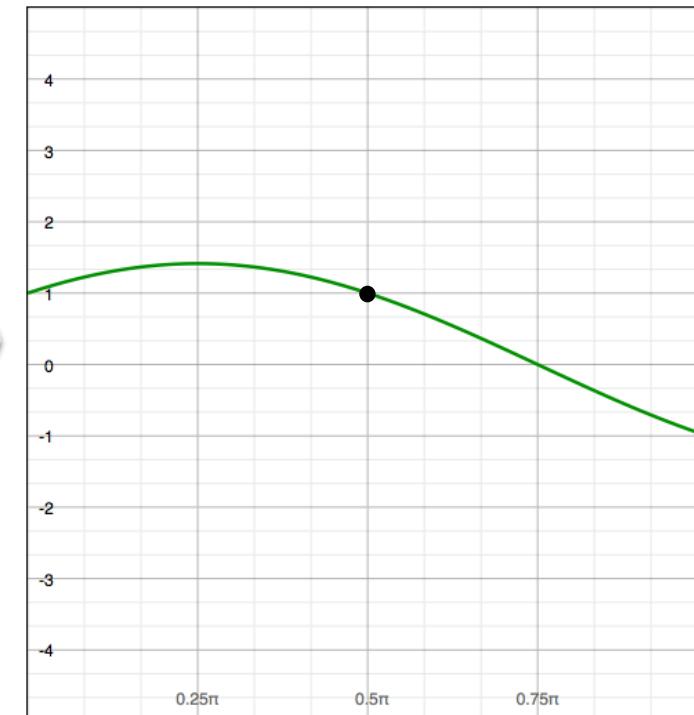


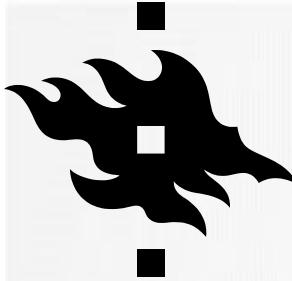
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line becomes a point



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

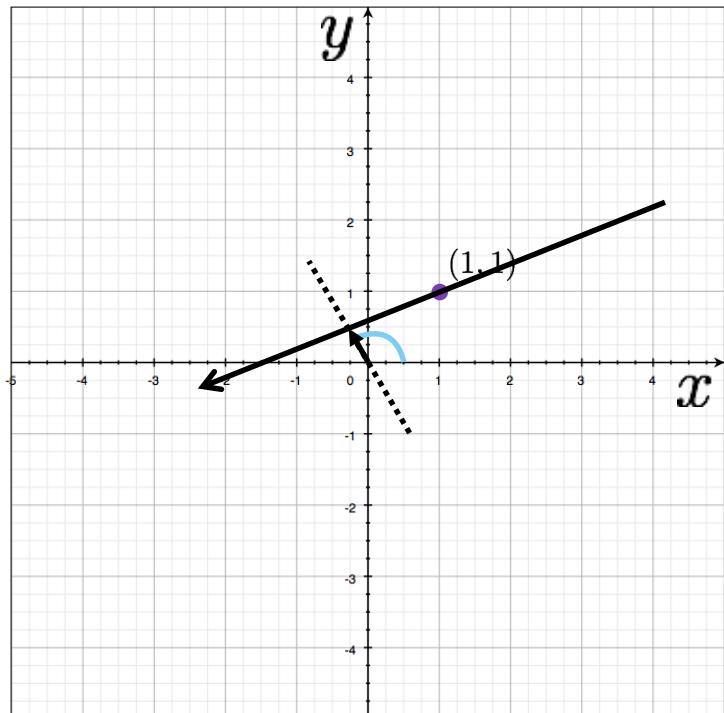
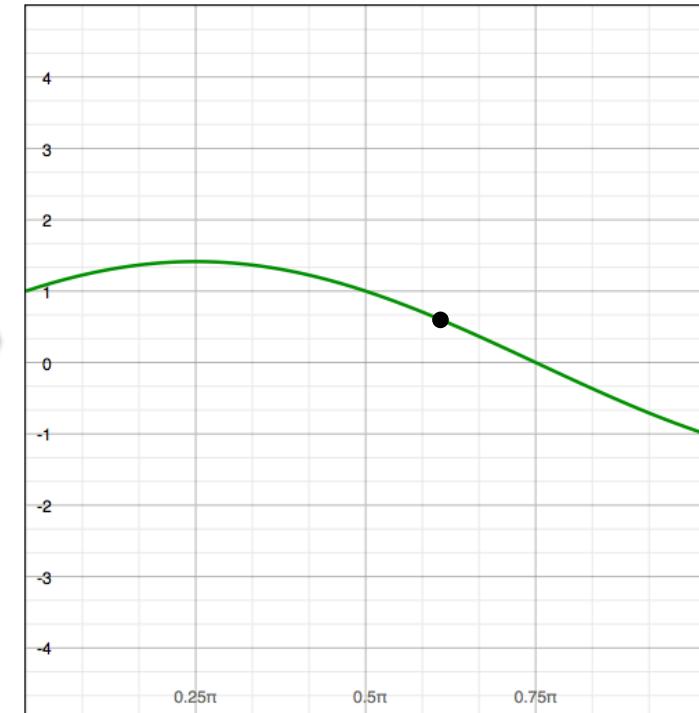


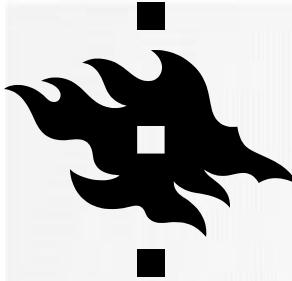
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line  
becomes  
a point



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

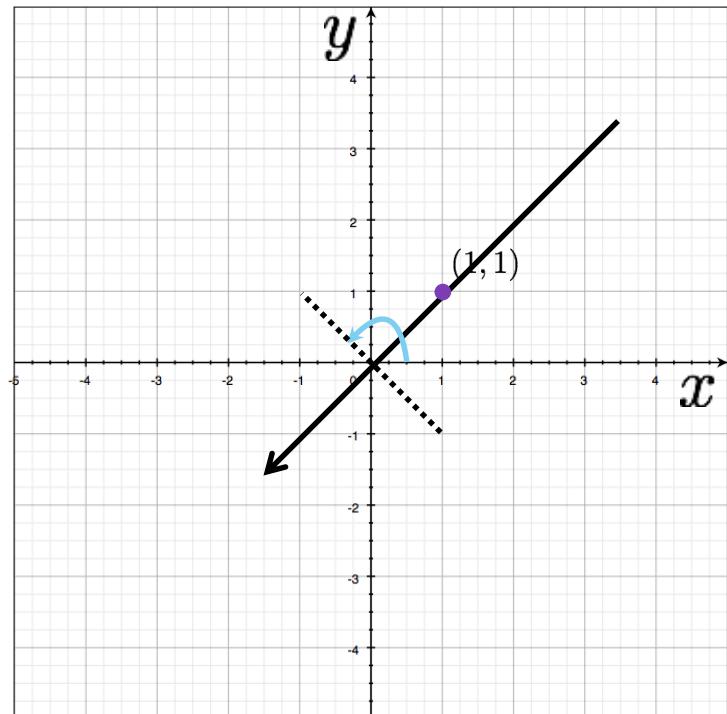
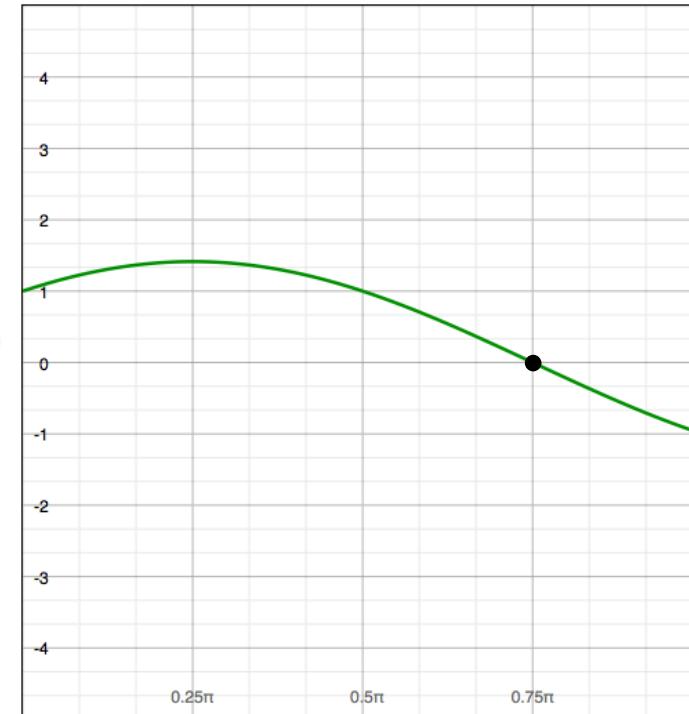


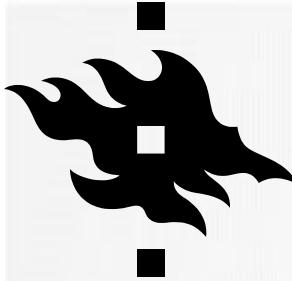
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line  
becomes  
a point



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

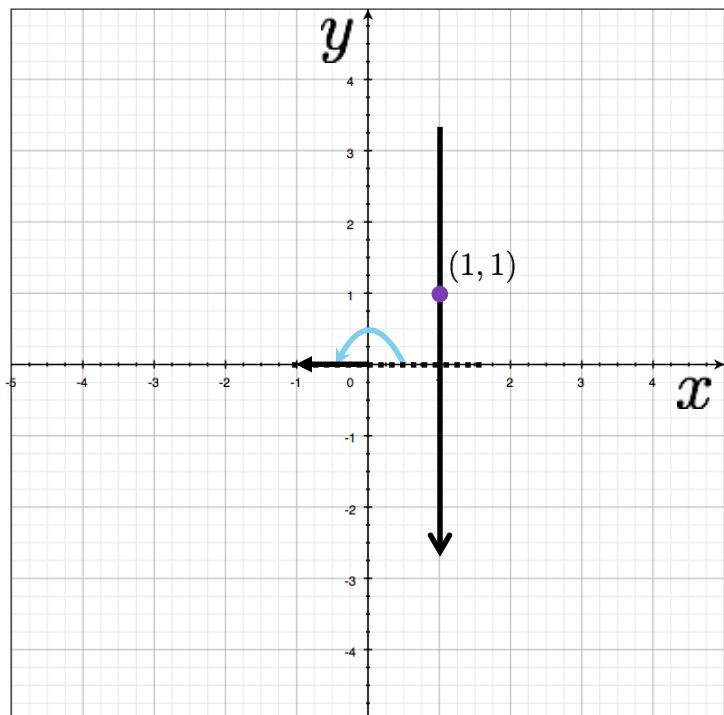
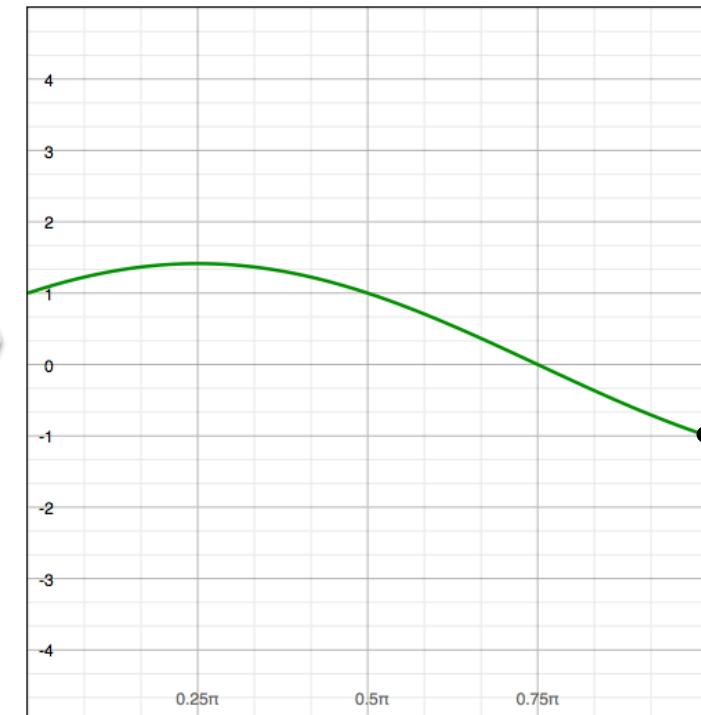


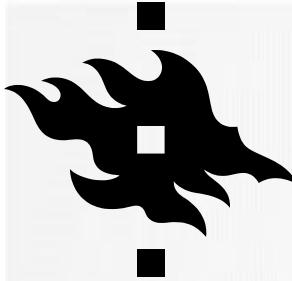
Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line becomes a point



Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters

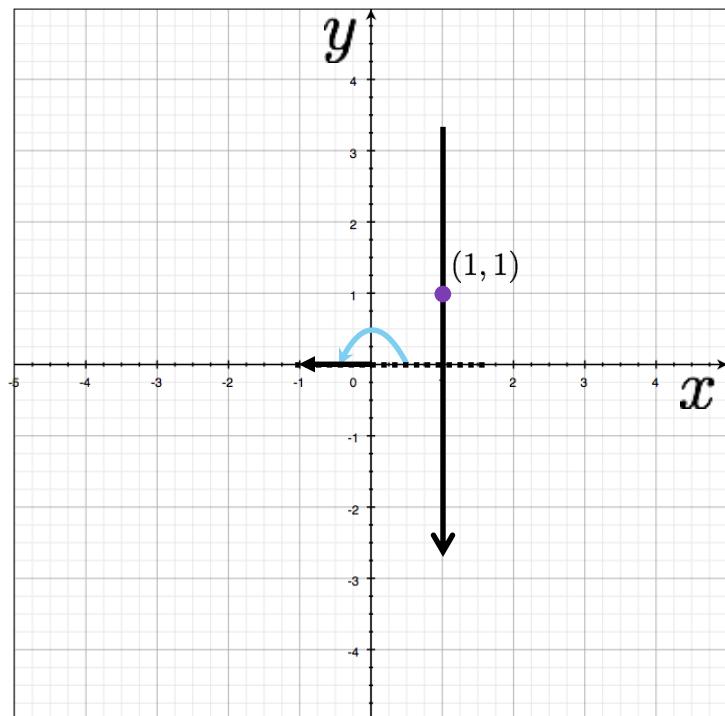
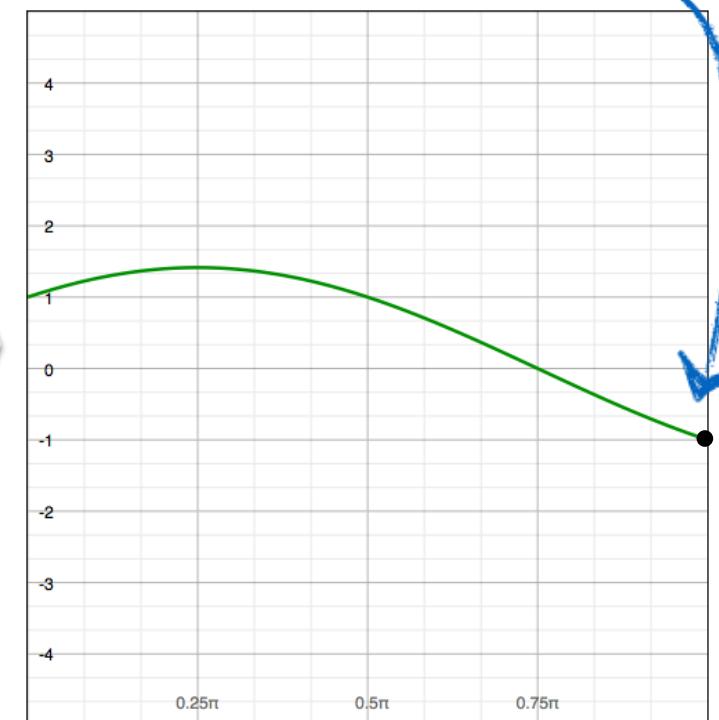


Image space

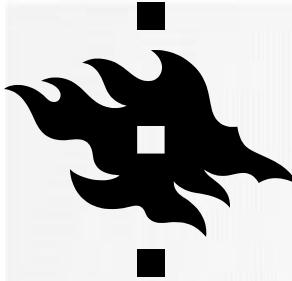
a line becomes a point

$$x \cos \theta + y \sin \theta = \rho$$

Wait ...why is rho negative?

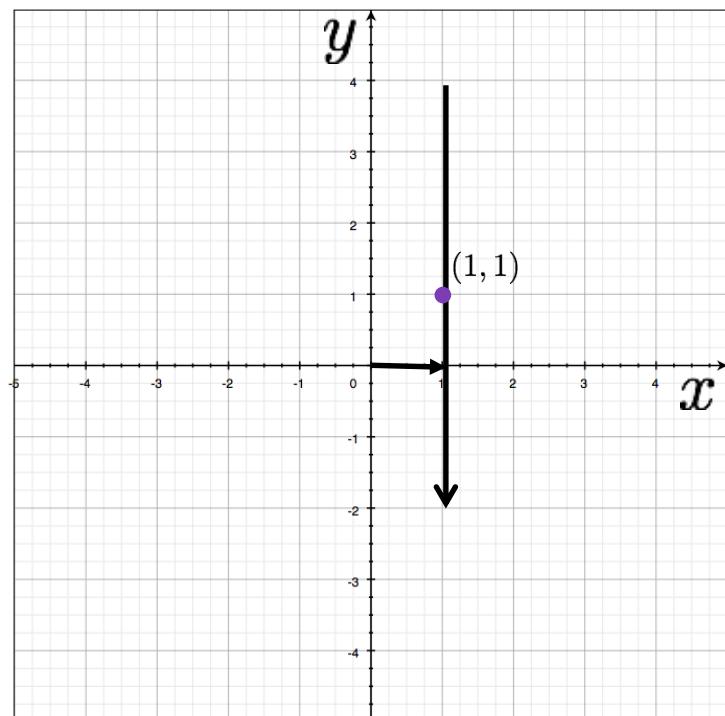


Parameter space



# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters



a line becomes a point

$$x \cos \theta + y \sin \theta = \rho$$

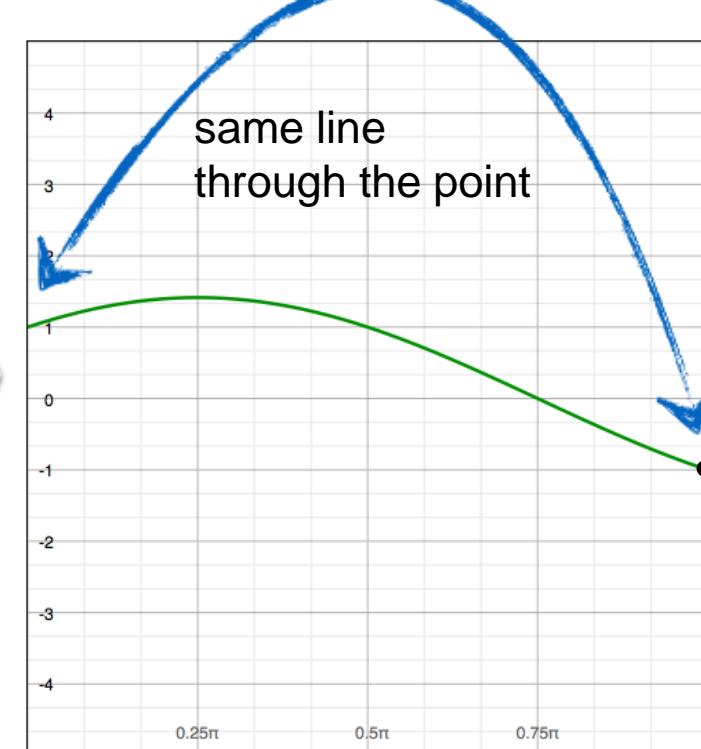
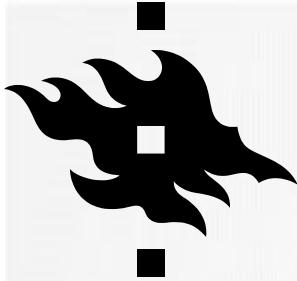


Image space

Parameter space



There are two ways to write the same line:

Positive rho version:

$$x \cos \theta + y \sin \theta = \rho$$

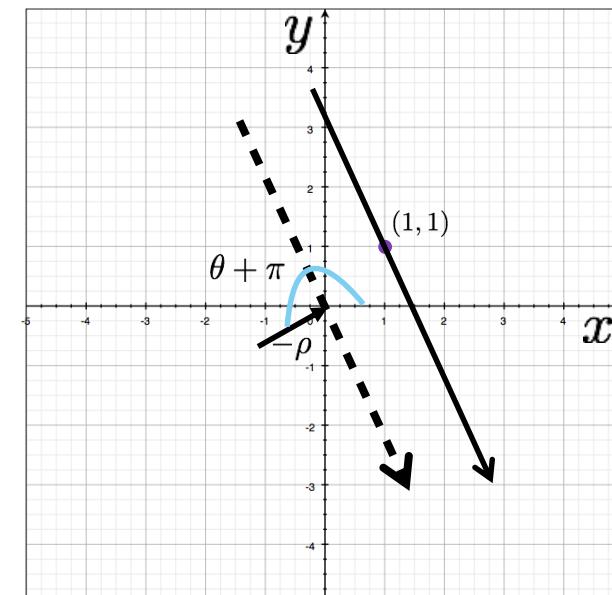
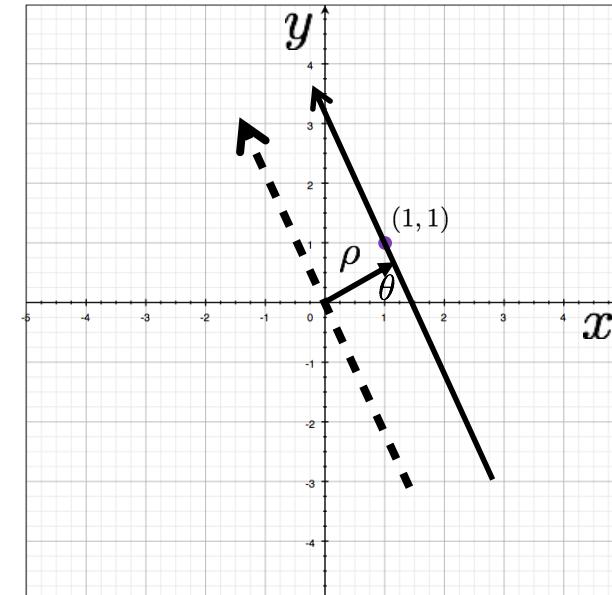
Negative rho version:

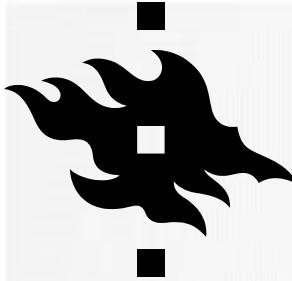
$$x \cos(\theta + \pi) + y \sin(\theta + \pi) = -\rho$$

Recall:

$$\sin(\theta) = -\sin(\theta + \pi)$$

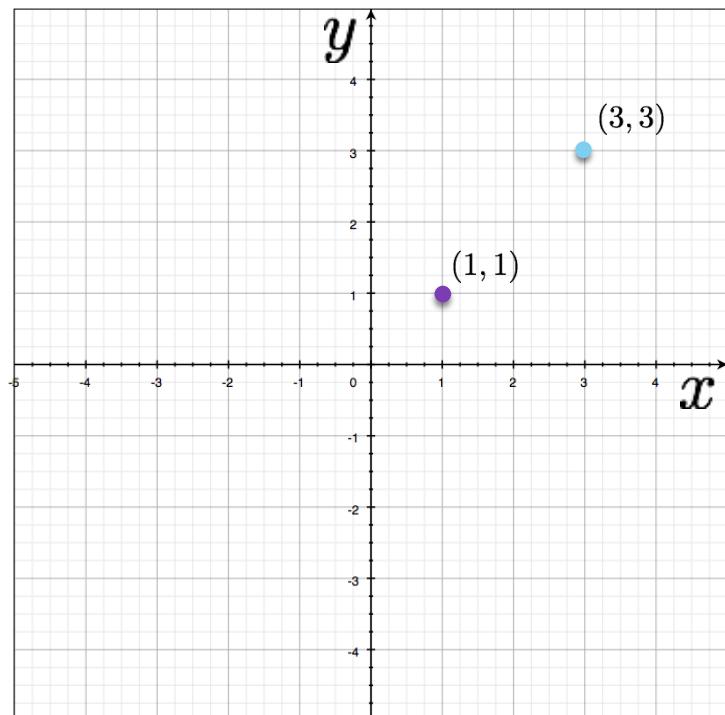
$$\cos(\theta) = -\cos(\theta + \pi)$$





# IMAGE AND PARAMETER SPACE

variables  
 $y = mx + b$   
parameters



two points  
become  
?

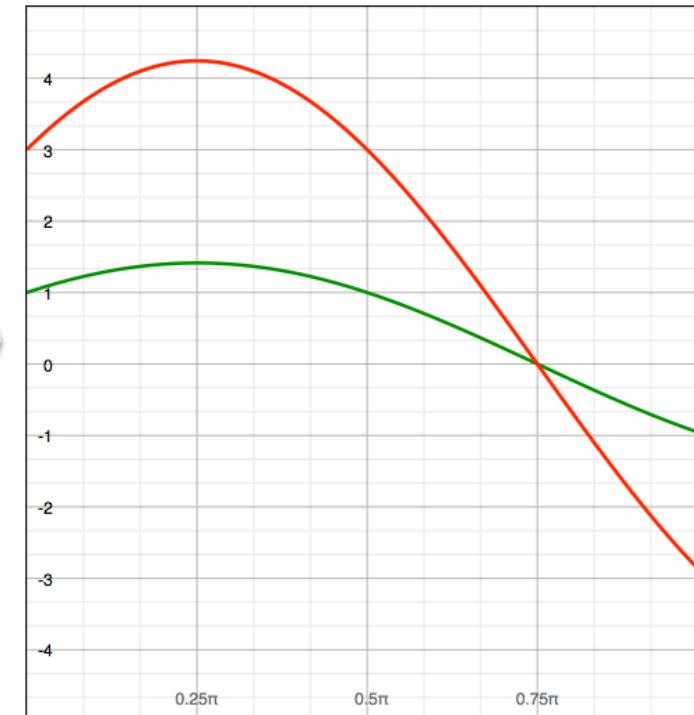
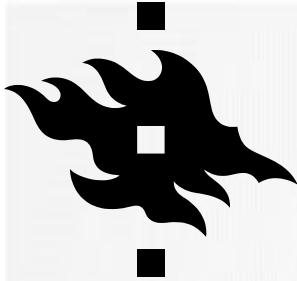


Image space

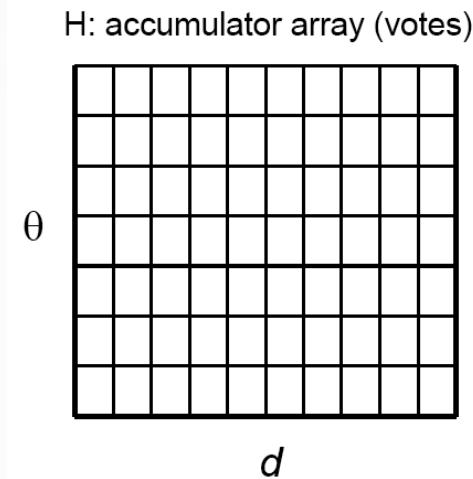
Parameter space



# IMPLEMENTATION



1. Initialize accumulator  $H$  to all zeros
2. For each edge point  $(x, y)$  in the image  
    For  $\theta = 0$  to 180  
         $\rho = x \cos \theta + y \sin \theta$   
         $H(\theta, \rho) = H(\theta, \rho) + 1$   
    end  
end
3. Find the value(s) of  $(\theta, \rho)$  where  $H(\theta, \rho)$  is a local maximum
4. The detected line in the image is given by  
$$\rho = x \cos \theta + y \sin \theta$$



NOTE: Watch your coordinates. Image origin is top left!

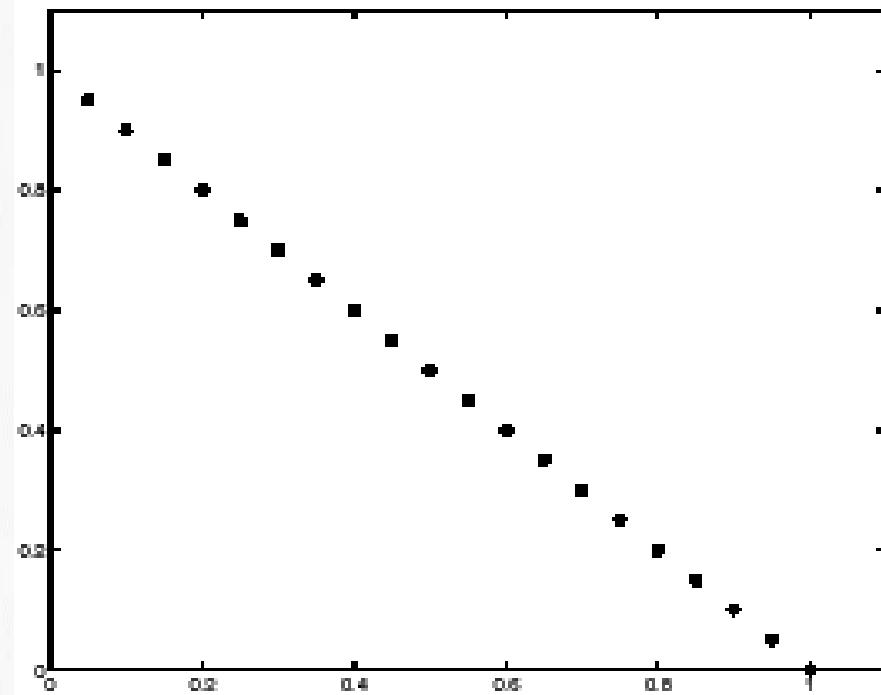
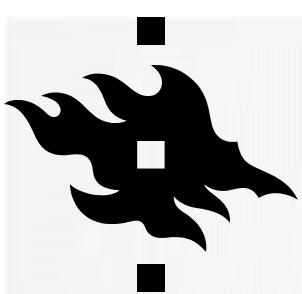
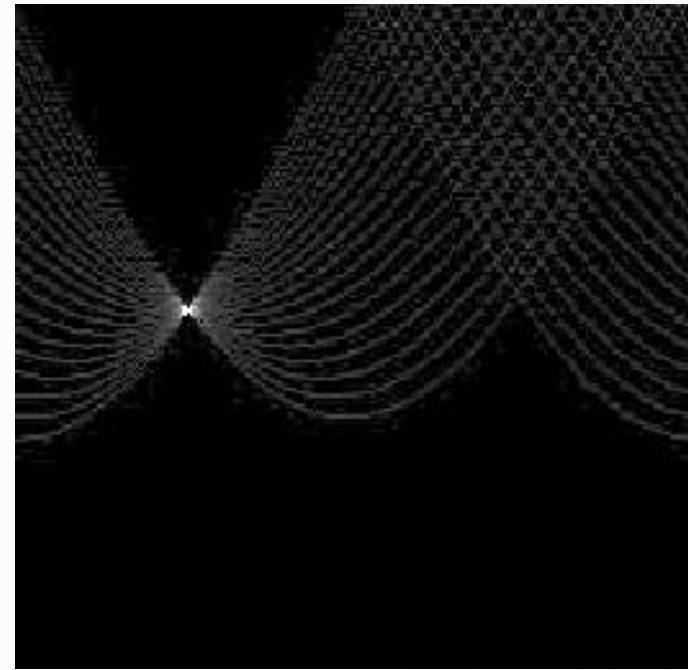
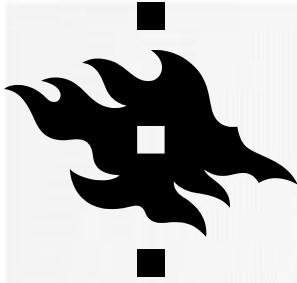


Image space

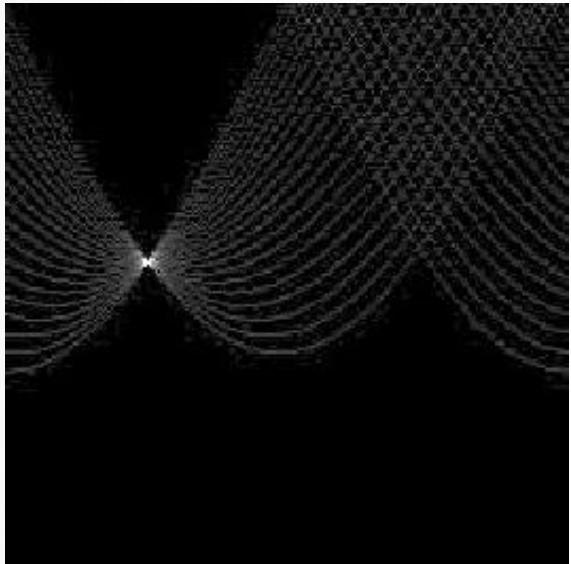


Votes

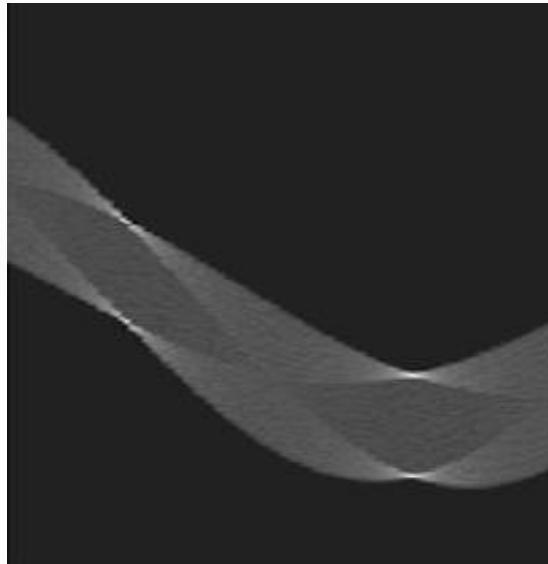


# Basic shapes

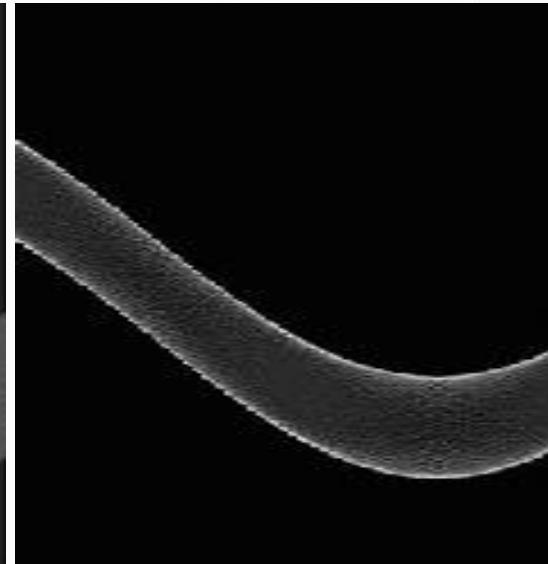
(in parameter space)



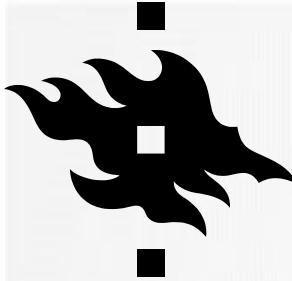
line



rectangle



circle

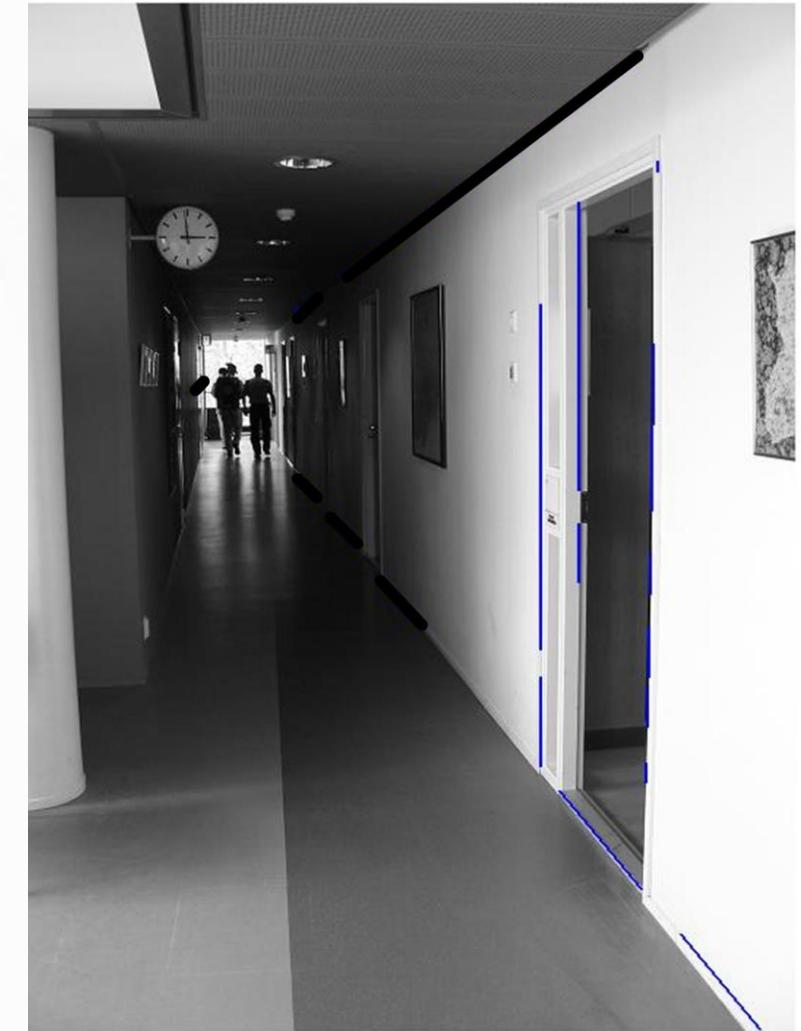
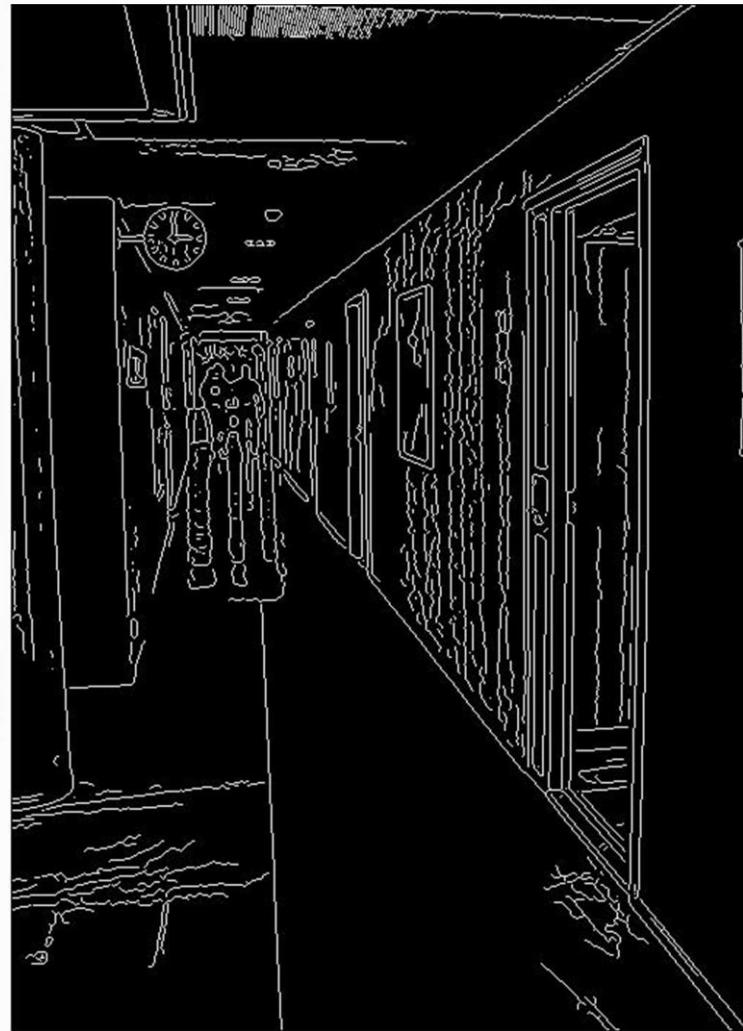
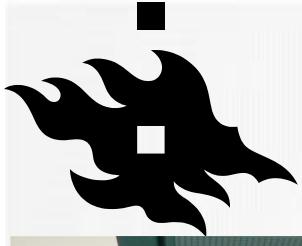


# HOUGH TRANSFORM



- Why don't we find all lines existing in the image?







60° 10 1.2 N, 24° 57 18 E