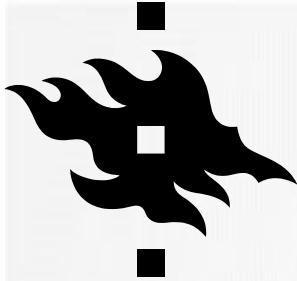


TODAY'S LECTURE

- Two-frame structure from motion (SFM) => estimatin simultaneously the 3D scene structure and camera pose from matched image features
 - Essential Matrix
 - Fundamental matrix
 - Solving ambiguous scale
 - Stereo vision
- Szeliski 7,
- Hartley, Zissermann

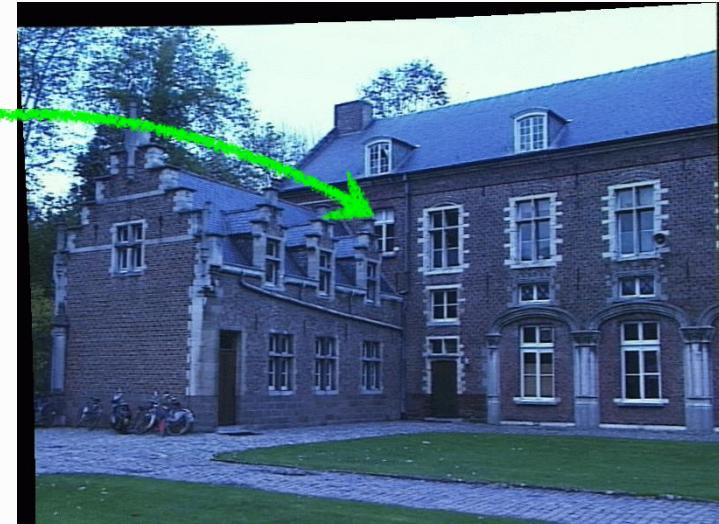


EPIPOLAR CONSTRAINT

Task: Match point in left image to point in right image



Left image

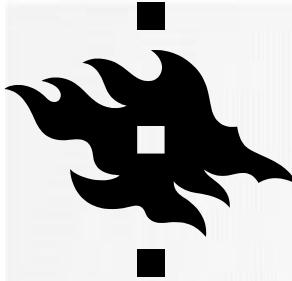


Right image

Epipolar constraint reduces search to a single line

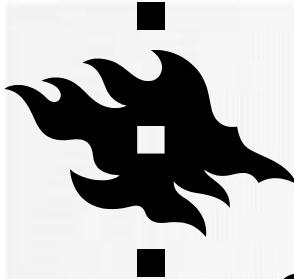
The epipolar constraint is an important concept for stereo vision

How do you compute the epipolar line?



FUNDAMENTAL MATRIX SONG

[https://www.youtube.com/watch?v=DgGV3I82
NTk](https://www.youtube.com/watch?v=DgGV3I82NTk)

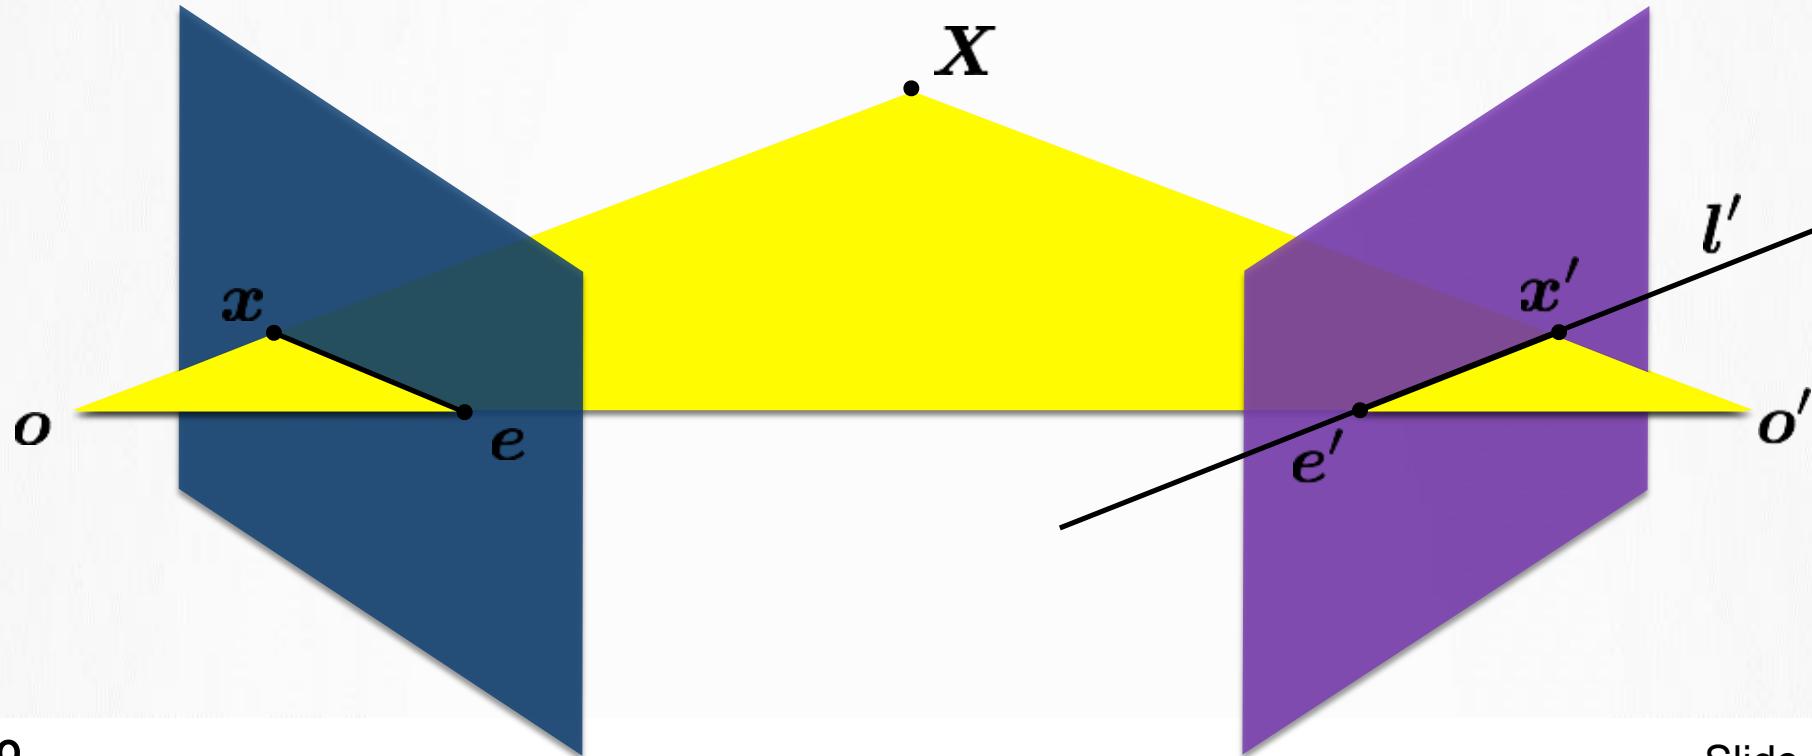


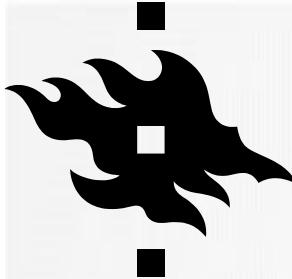
THE ESSENTIAL MATRIX



Given a point in one image, multiplying by the **essential matrix** will tell us the **epipolar line** in the second view.

$$\mathbf{E}x = l'$$





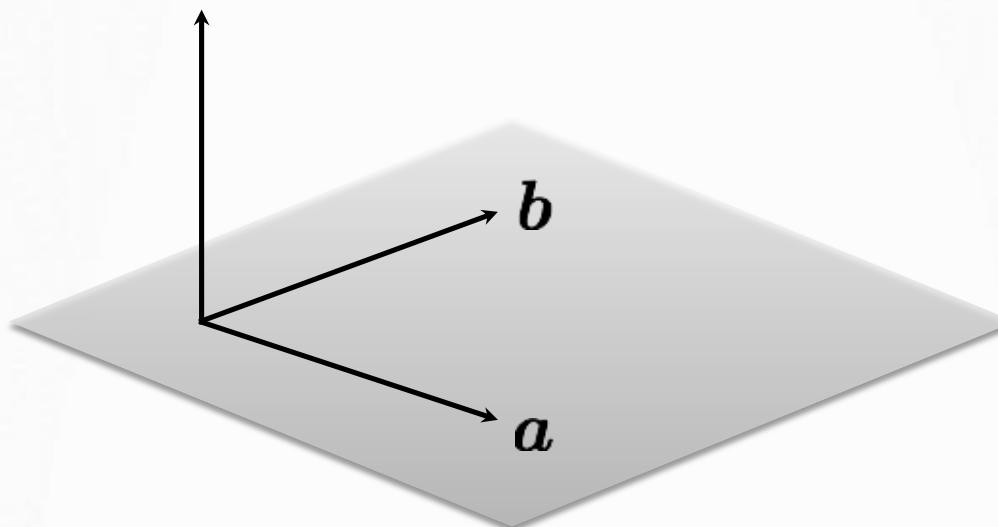
RECALL: CROSS PRODUCT



Vector (cross) product

takes two vectors and returns a vector perpendicular to both

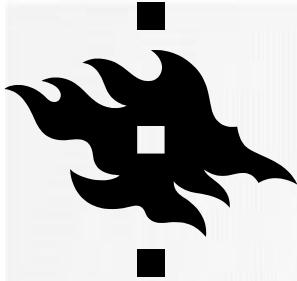
$$\mathbf{c} = \mathbf{a} \times \mathbf{b}$$



$$\mathbf{c} \cdot \mathbf{a} = 0$$

$$\mathbf{c} \cdot \mathbf{b} = 0$$

**dot product of
two orthogonal
vectors is zero**



Cross product

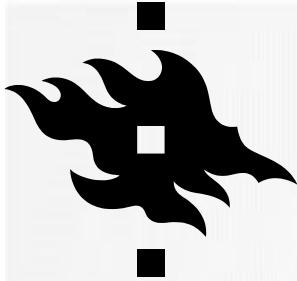
$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} a_2 b_3 - a_3 b_2 \\ a_3 b_1 - a_1 b_3 \\ a_1 b_2 - a_2 b_1 \end{bmatrix}$$



Can also be written as a matrix multiplication

$$\mathbf{a} \times \mathbf{b} = [\mathbf{a}]_{\times} \mathbf{b} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix} \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

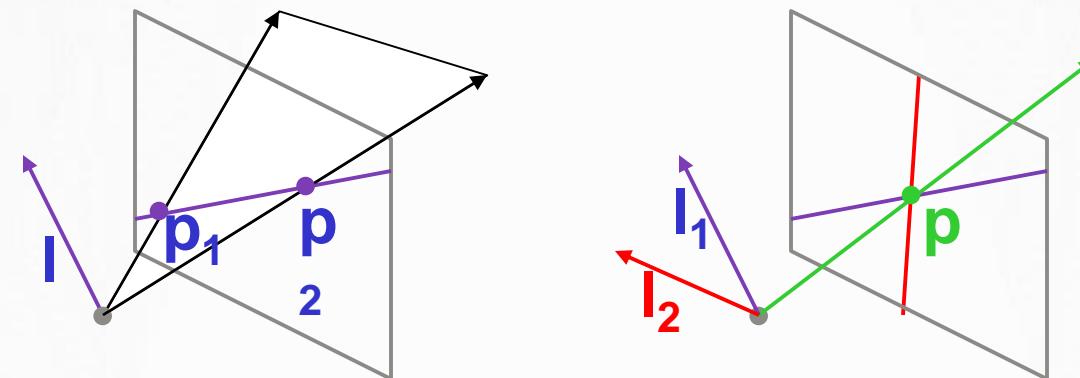
Skew symmetric



POINT AND LINE DUALITY (LECTURE 2)

A line \mathbf{l} is a homogeneous 3-vector

It is \perp to every point (ray) \mathbf{p} on the line: $\mathbf{l} \cdot \mathbf{p} = 0$

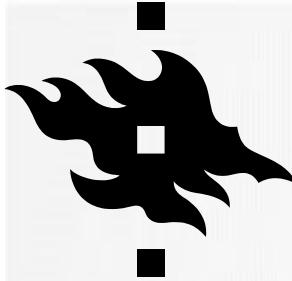


What is the line \mathbf{l} spanned by rays \mathbf{p}_1 and \mathbf{p}_2 ?

- \mathbf{l} is \perp to \mathbf{p}_1 and $\mathbf{p}_2 \Rightarrow \mathbf{l} = \mathbf{p}_1 \times \mathbf{p}_2$
- \mathbf{l} can be interpreted as a *plane normal*

What is the intersection of two lines \mathbf{l}_1 and \mathbf{l}_2 ?

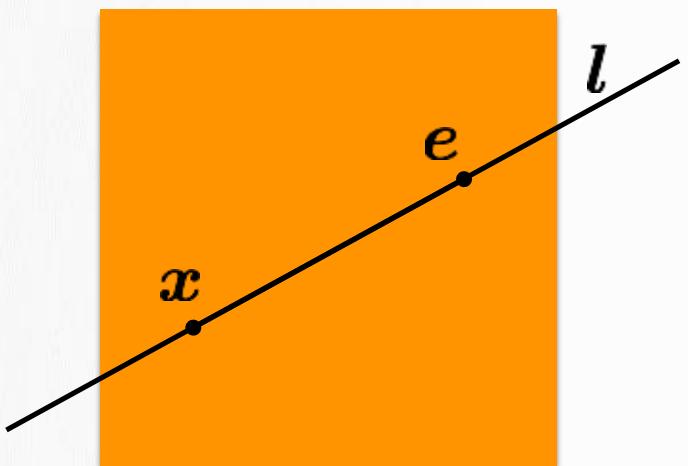
- \mathbf{p} is \perp to \mathbf{l}_1 and $\mathbf{l}_2 \Rightarrow \mathbf{p} = \mathbf{l}_1 \times \mathbf{l}_2$



EPIPOLAR LINE

$$ax + by + c = 0 \quad \text{in vector form}$$

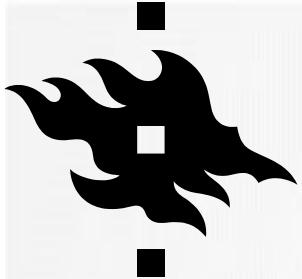
$$\mathbf{l} = \begin{bmatrix} a \\ b \\ c \end{bmatrix}$$



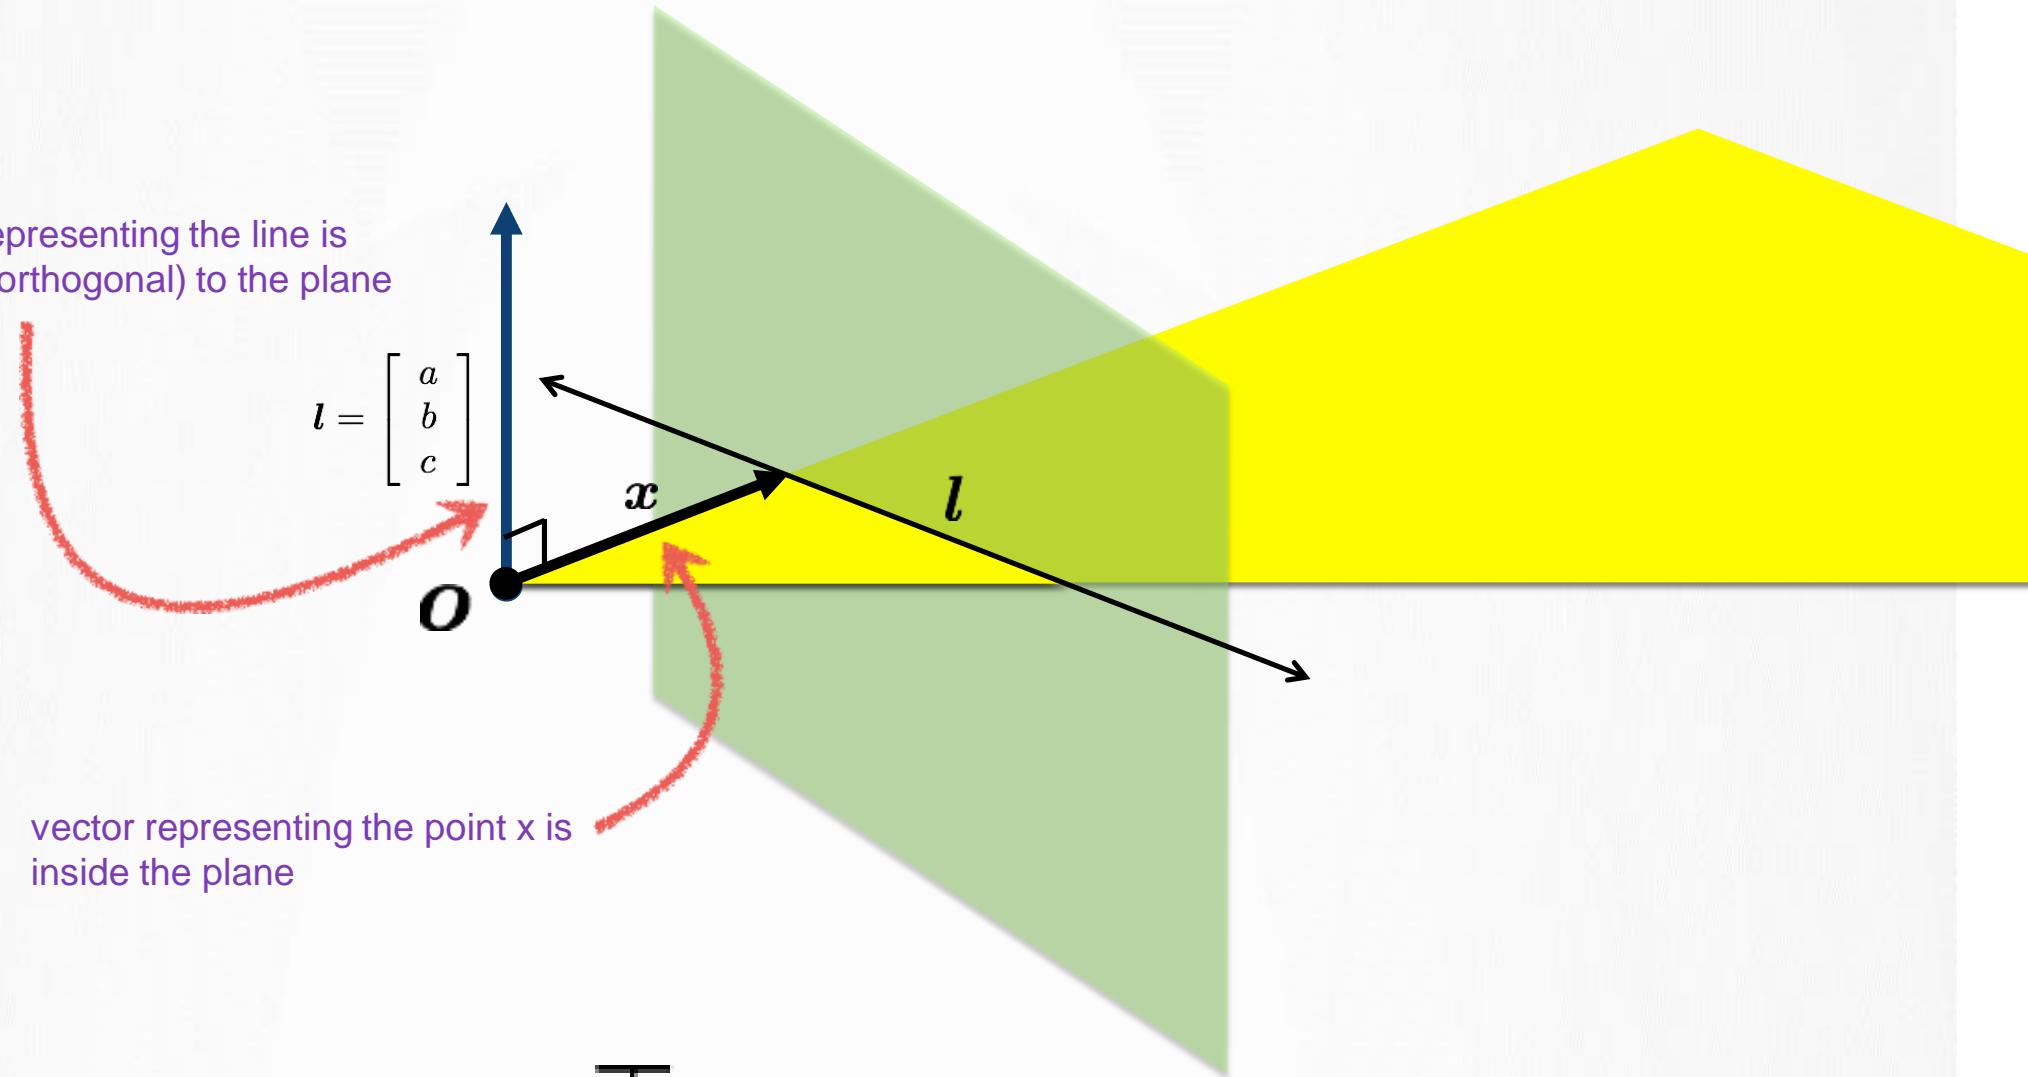
If the point \mathbf{x} is on the epipolar line \mathbf{l} then

$$\mathbf{x}^\top \mathbf{l} = 0$$

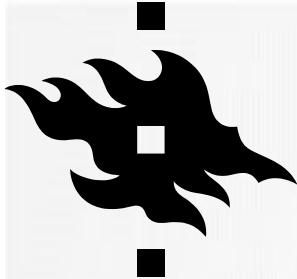




vector representing the line is
normal (orthogonal) to the plane



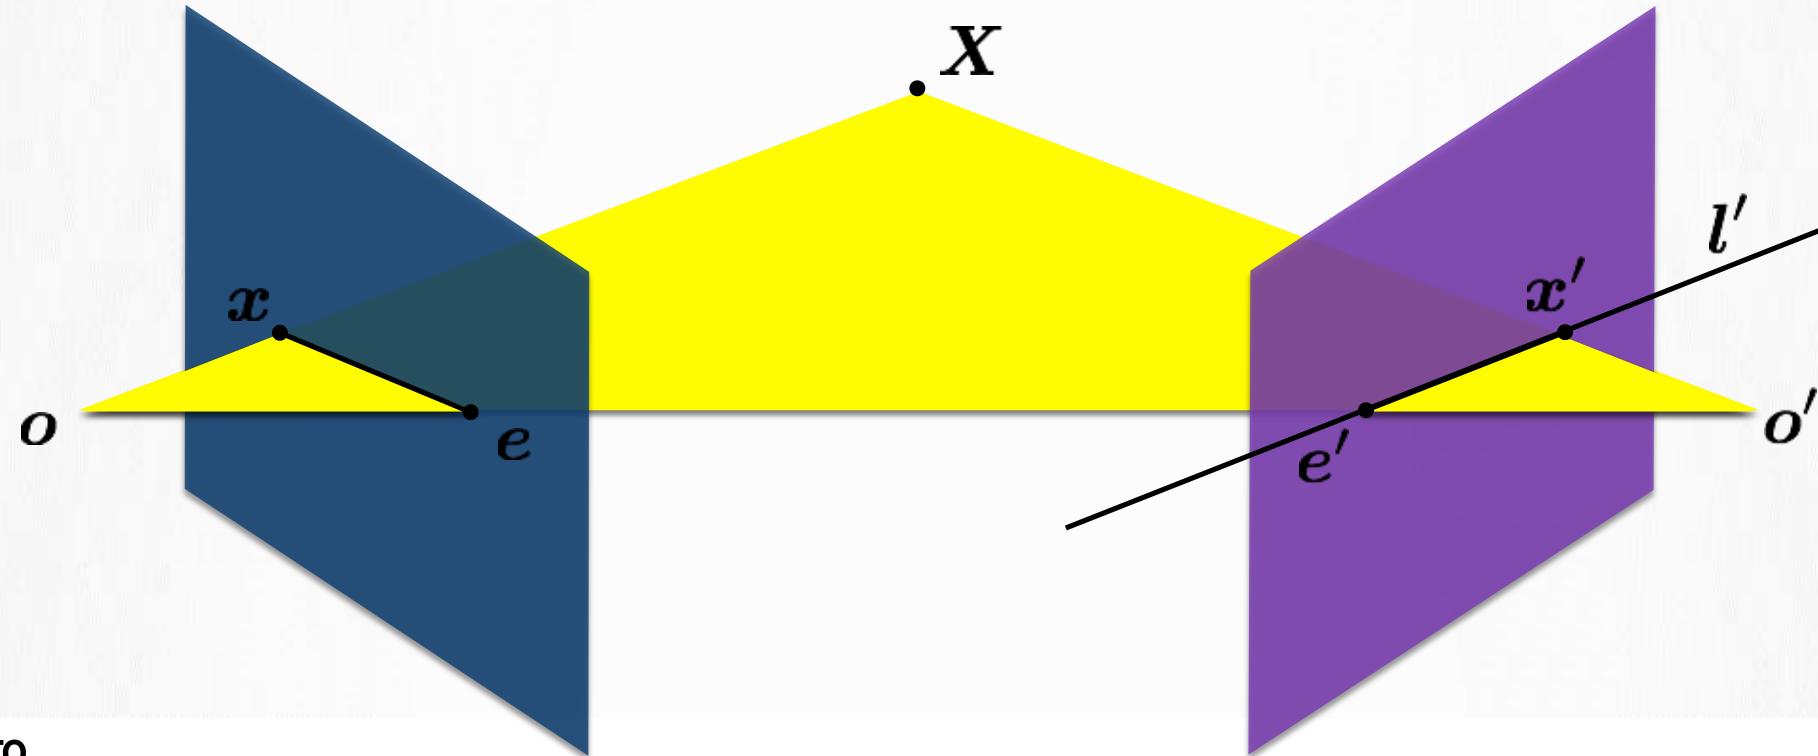
$$x^\top l = 0$$

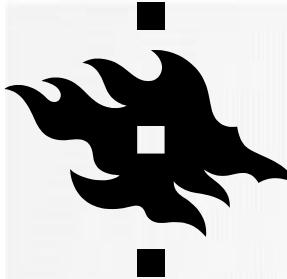


So if $\mathbf{x}^\top \mathbf{l} = 0$ and $\mathbf{E}\mathbf{x} = \mathbf{l}'$ then



$$\mathbf{x}'^\top \mathbf{E}\mathbf{x} = 0$$





PUTTING IT TOGETHER



rigid motion

$$\mathbf{x}' = \mathbf{R}(\mathbf{x} - \mathbf{t}) \quad (\mathbf{x} - \mathbf{t})^\top (\mathbf{t} \times \mathbf{x}) = 0$$

coplanarity

$$(\mathbf{x}'^\top \mathbf{R})(\mathbf{t} \times \mathbf{x}) = 0$$

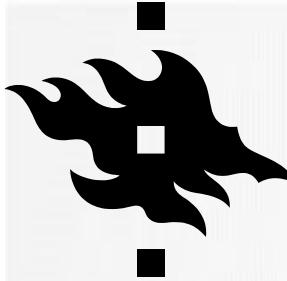
$$(\mathbf{x}'^\top \mathbf{R})([\mathbf{t}_\times] \mathbf{x}) = 0$$

$$\mathbf{x}'^\top (\mathbf{R}[\mathbf{t}_\times]) \mathbf{x} = 0$$

$$\boxed{\mathbf{x}'^\top \mathbf{E} \mathbf{x} = 0}$$

Essential Matrix
[Longuet-Higgins 1981]

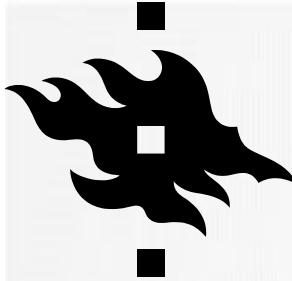
Slide credit:
Kris Kitani



PROPERTIES OF THE E MATRIX

Longuet-Higgins equation

$$\mathbf{x}'^\top \mathbf{E} \mathbf{x} = 0$$



PROPERTIES OF THE E MATRIX

Longuet-Higgins equation

$$\mathbf{x}'^\top \mathbf{E} \mathbf{x} = 0$$

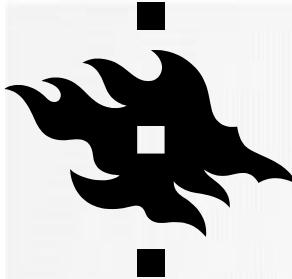
Epipolar lines

$$\mathbf{x}^\top \mathbf{l} = 0$$

$$\mathbf{l}' = \mathbf{E} \mathbf{x}$$

$$\mathbf{x}'^\top \mathbf{l}' = 0$$

$$\mathbf{l} = \mathbf{E}^T \mathbf{x}'$$



PROPERTIES OF THE E MATRIX



Longuet-Higgins equation

$$\mathbf{x}'^\top \mathbf{E} \mathbf{x} = 0$$

Epipolar lines

$$\mathbf{x}^\top \mathbf{l} = 0$$

$$\mathbf{l}' = \mathbf{E} \mathbf{x}$$

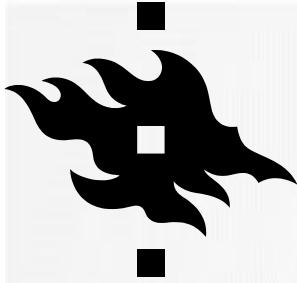
$$\mathbf{x}'^\top \mathbf{l}' = 0$$

$$\mathbf{l} = \mathbf{E}^T \mathbf{x}'$$

Epipoles

$$\mathbf{e}'^\top \mathbf{E} = 0$$

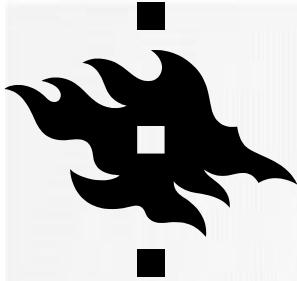
$$\mathbf{E} \mathbf{e} = 0$$



THE FUNDAMENTAL MATRIX

- Essential matrix requires that the points are aligned to camera coordinate axis (calibrated camera)
- When the task needs to be generalized, a Fundamental Matrix is used
- We can set camera pose for image 1 as $t=(0,0,0)'$, R identity





$$\hat{\mathbf{x}}'^\top \mathbf{E} \hat{\mathbf{x}} = 0$$

The Essential matrix operates on image points expressed in
normalized coordinates

(points have been aligned (normalized) to camera coordinates)

$$\hat{\mathbf{x}}' = \mathbf{K}'^{-1} \mathbf{x}'$$

$$\hat{\mathbf{x}} = \mathbf{K}^{-1} \mathbf{x}$$

camera
point

image
point

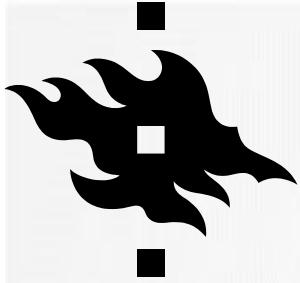
Writing out the epipolar constraint in terms of image coordinates

$$\mathbf{x}'^\top \mathbf{K}'^{-\top} \mathbf{E} \mathbf{K}^{-1} \mathbf{x} = 0$$

$$\mathbf{x}'^\top (\mathbf{K}'^{-\top} \mathbf{E} \mathbf{K}^{-1}) \mathbf{x} = 0$$

$$\mathbf{x}'^\top \mathbf{F} \mathbf{x} = 0$$

maps pixels to epipolar lines

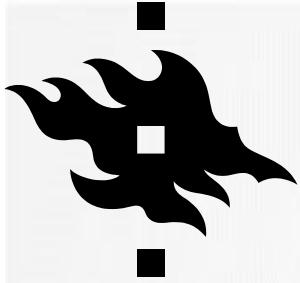


Breaking down the fundamental matrix

$$\mathbf{F} = \mathbf{K}'^{-\top} \mathbf{E} \mathbf{K}^{-1}$$

$$\mathbf{F} = \mathbf{K}'^{-\top} [\mathbf{t}_x] \mathbf{R} \mathbf{K}^{-1}$$

Depends on both intrinsic and extrinsic parameters



Breaking down the fundamental matrix

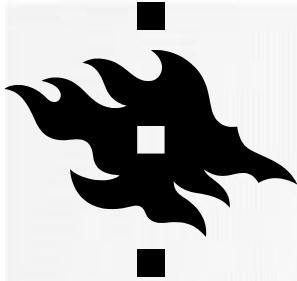
$$\mathbf{F} = \mathbf{K}'^{-\top} \mathbf{E} \mathbf{K}^{-1}$$

$$\mathbf{F} = \mathbf{K}'^{-\top} [\mathbf{t}_x] \mathbf{R} \mathbf{K}^{-1}$$

Depends on both intrinsic and extrinsic parameters

How would you solve for F?

$$\mathbf{x}_m'^\top \mathbf{F} \mathbf{x}_m = 0$$

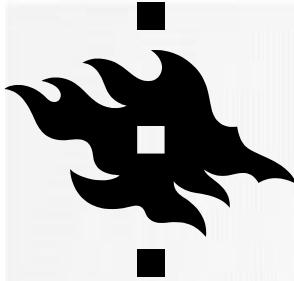


ESTIMATING F



If we don't know K_1 , K_2 , R , or t , can we estimate F for two images?

Yes, given enough correspondences



ESTIMATING F – 8-POINT ALGORITHM

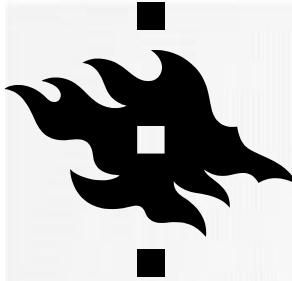
The fundamental matrix F is defined by

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$

for any pair of matches \mathbf{x} and \mathbf{x}' in two images.

- Let $\mathbf{x}=(u,v,1)^T$ and $\mathbf{x}'=(u',v',1)^T$, $\mathbf{F} = \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix}$
each match gives a linear equation

$$uu'f_{11} + vu'f_{12} + u'f_{13} + uv'f_{21} + vv'f_{22} + v'f_{23} + uf_{31} + vf_{32} + f_{33} = 0$$

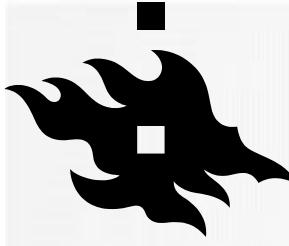


8-POINT ALGORITHM



$$\begin{bmatrix} u_1 u_1' & v_1 u_1' & u_1' & u_1 v_1' & v_1 v_1' & v_1' & u_1 & v_1 & 1 \\ u_2 u_2' & v_2 u_2' & u_2' & u_2 v_2' & v_2 v_2' & v_2' & u_2 & v_2 & 1 \\ \vdots & \vdots \\ u_n u_n' & v_n u_n' & u_n' & u_n v_n' & v_n v_n' & v_n' & u_n & v_n & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{12} \\ f_{13} \\ f_{21} \\ f_{22} \\ f_{23} \\ f_{31} \\ f_{32} \\ f_{33} \end{bmatrix} = 0$$

Like with homographies, instead of solving $\mathbf{Af} = 0$, we seek \mathbf{f} to minimize $\|\mathbf{Af}\|$, least eigenvector of $\mathbf{A}^T \mathbf{A}$.



8-POINT ALGORITHM – PROBLEM?

- \mathbf{F} should have rank 2

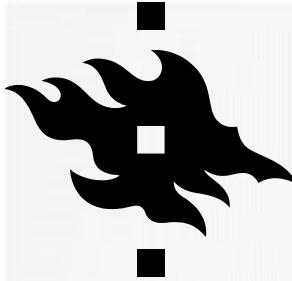
To enforce that \mathbf{F} is of rank 2, \mathbf{F} is replaced by \mathbf{F}' that minimizes subject to the rank constraint.

$$\|\mathbf{F} - \mathbf{F}'\|$$

- This is achieved by SVD. Let $\mathbf{F} = \mathbf{U}\Sigma\mathbf{V}^T$ where

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{bmatrix}, \text{ let } \Sigma' = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

then $\mathbf{F}' = \mathbf{U}\Sigma'\mathbf{V}^T$ is the solution.



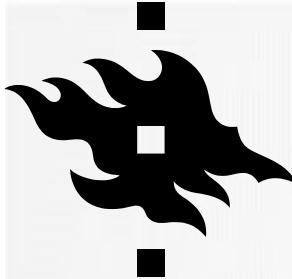
8-POINT ALGORITHM

```
% Build the constraint matrix
A = [x2(1,:)'.*x1(1,:) ' x2(1,:)'.*x1(2,:) ' x2(1,:)' ...
      x2(2,:)'.*x1(1,:) ' x2(2,:)'.*x1(2,:) ' x2(2,:)' ...
      x1(1,:) ' x1(2,:) ' ones(npts,1) ];

[U,D,V] = svd(A);

% Extract fundamental matrix from the column of V
% corresponding to the smallest singular value.
F = reshape(V(:,9),3,3)';

% Enforce rank2 constraint
[U,D,V] = svd(F);
F = U*diag([D(1,1) D(2,2) 0])*V';
```



MOTION FROM E

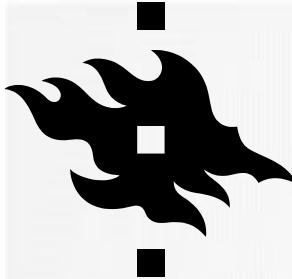
$$E = [\hat{t}]_x R = U \Sigma V^T = \begin{bmatrix} u_0 & u_1 & \hat{t} \end{bmatrix} \begin{bmatrix} 1 & & \\ & 1 & \\ & & 0 \end{bmatrix} \begin{bmatrix} v_0^T \\ v_1^T \\ v_2^T \end{bmatrix}$$



t = singular value associated with smallest singular vector

R = four possible matrices $R = \pm U R_{\pm 90^\circ}^T V^T$

1. Keep the ones with $\det |R|=1$
2. Pair the remaining ones with both possible signs of translation direction
3. Select the combination for which the largest number of points is seen in front of both cameras



PROBLEM WITH 8-POINT ALGORITHM

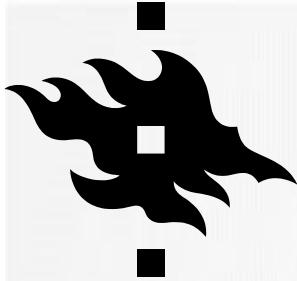


$$\begin{bmatrix} u_1 u_1' & v_1 u_1' & u_1' & u_1 v_1' & v_1 v_1' & v_1' & u_1 & v_1 & 1 \\ u_2 u_2' & v_2 u_2' & u_2' & u_2 v_2' & v_2 v_2' & v_2' & u_2 & v_2 & 1 \\ \vdots & \vdots \\ u_n u_n' & v_n u_n' & u_n' & u_n v_n' & v_n v_n' & v_n' & u_n & v_n & 1 \end{bmatrix} = 0$$

$\sim 10000 \quad \sim 10000 \quad \sim 100 \quad \sim 10000 \quad \sim 10000 \quad \sim 100 \quad \sim 100 \quad \sim 100 \quad 1$



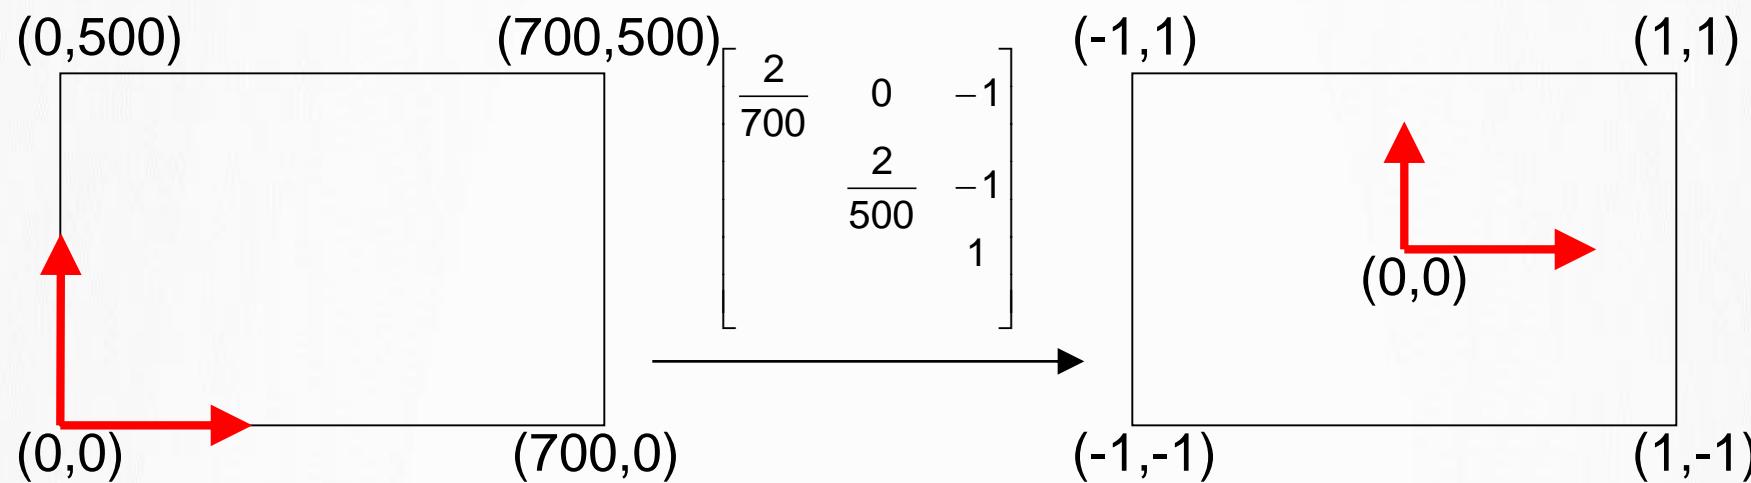
Orders of magnitude difference
between column of data matrix
→ least-squares yields poor results

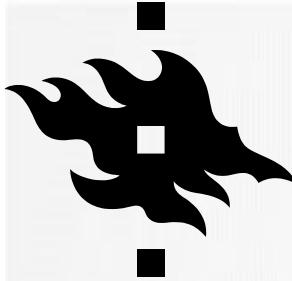


NORMALIZED 8-POINT ALGORITHM

normalized least squares yields good results

Transform image to $\sim[-1,1] \times [-1,1]$

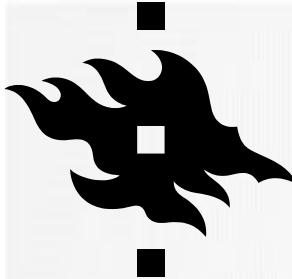




NORMALIZED 8-POINT ALGORITHM

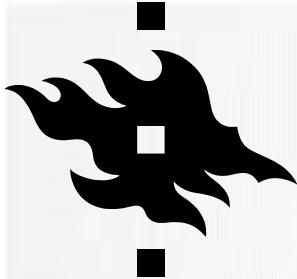
1. Transform input by $\hat{\mathbf{x}}_i = \mathbf{T}\mathbf{x}_i$, $\hat{\mathbf{x}}'_i = \mathbf{T}\mathbf{x}'_i$
2. Call 8-point on $\hat{\mathbf{x}}_i, \hat{\mathbf{x}}'_i$ to obtain $\hat{\mathbf{F}}$
3. $\mathbf{F} = \mathbf{T}'^T \hat{\mathbf{F}} \mathbf{T}$

$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$
$$\hat{\mathbf{x}}'^T \mathbf{T}'^{-T} \hat{\mathbf{F}} \mathbf{T}^{-1} \hat{\mathbf{x}} = 0$$
$$\underbrace{\hat{\mathbf{x}}'^T \mathbf{T}'^{-T} \hat{\mathbf{F}} \mathbf{T}^{-1} \hat{\mathbf{x}}}_{\hat{\mathbf{F}}} = 0$$

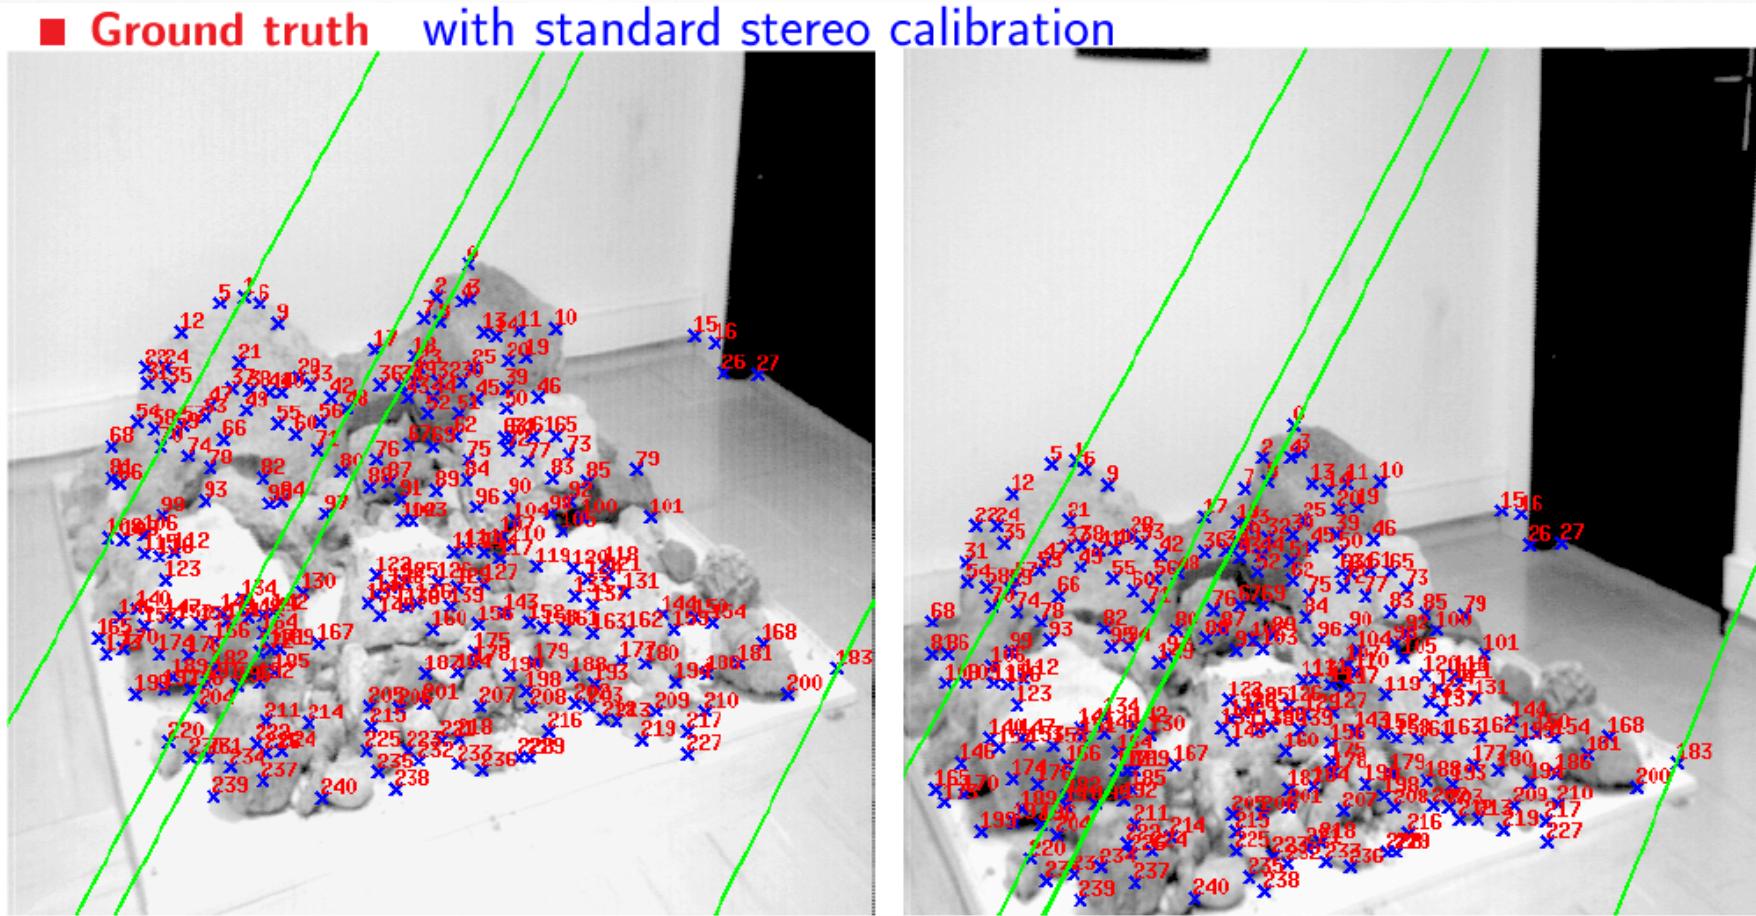


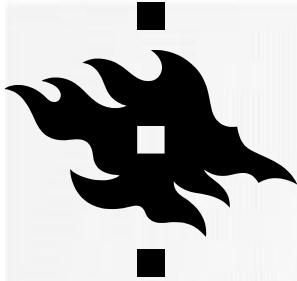
NORMALIZED 8-POINT ALGORITHM

```
[x1, T1] = normalise2dpts(x1);  
[x2, T2] = normalise2dpts(x2);  
  
A = [x2(1,:)'.*x1(1,:)'    x2(1,:}'.*x1(2,:)'    x2(1,:)' ...  
      x2(2,:)'.*x1(1,:)'    x2(2,:}'.*x1(2,:)'    x2(2,:)' ...  
      x1(1,:)'                x1(2,:)'                ones(npts,1)];  
  
[U,D,V] = svd(A);  
F = reshape(V(:,9),3,3)';  
[U,D,V] = svd(F);  
F = U*diag([D(1,1) D(2,2) 0])*V';  
  
% Denormalise  
F = T2'*F*T1;
```

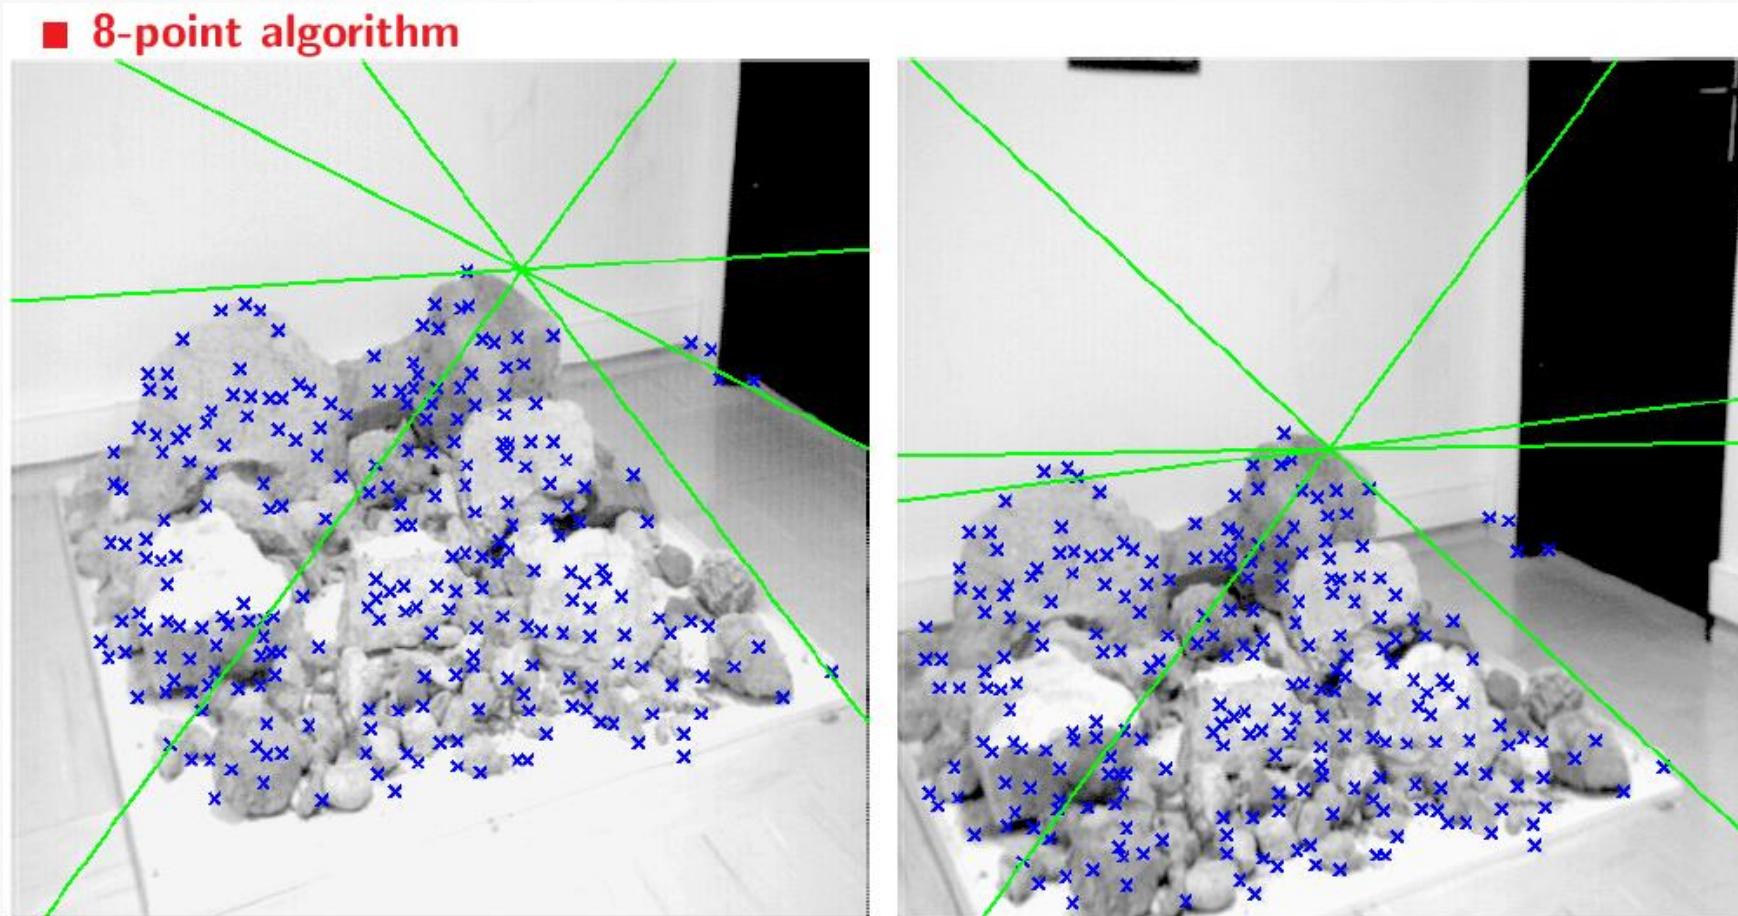


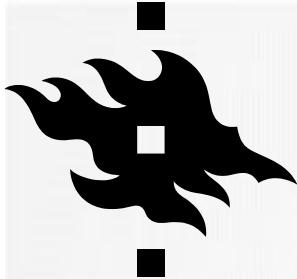
RESULTS (GROUND TRUTH)



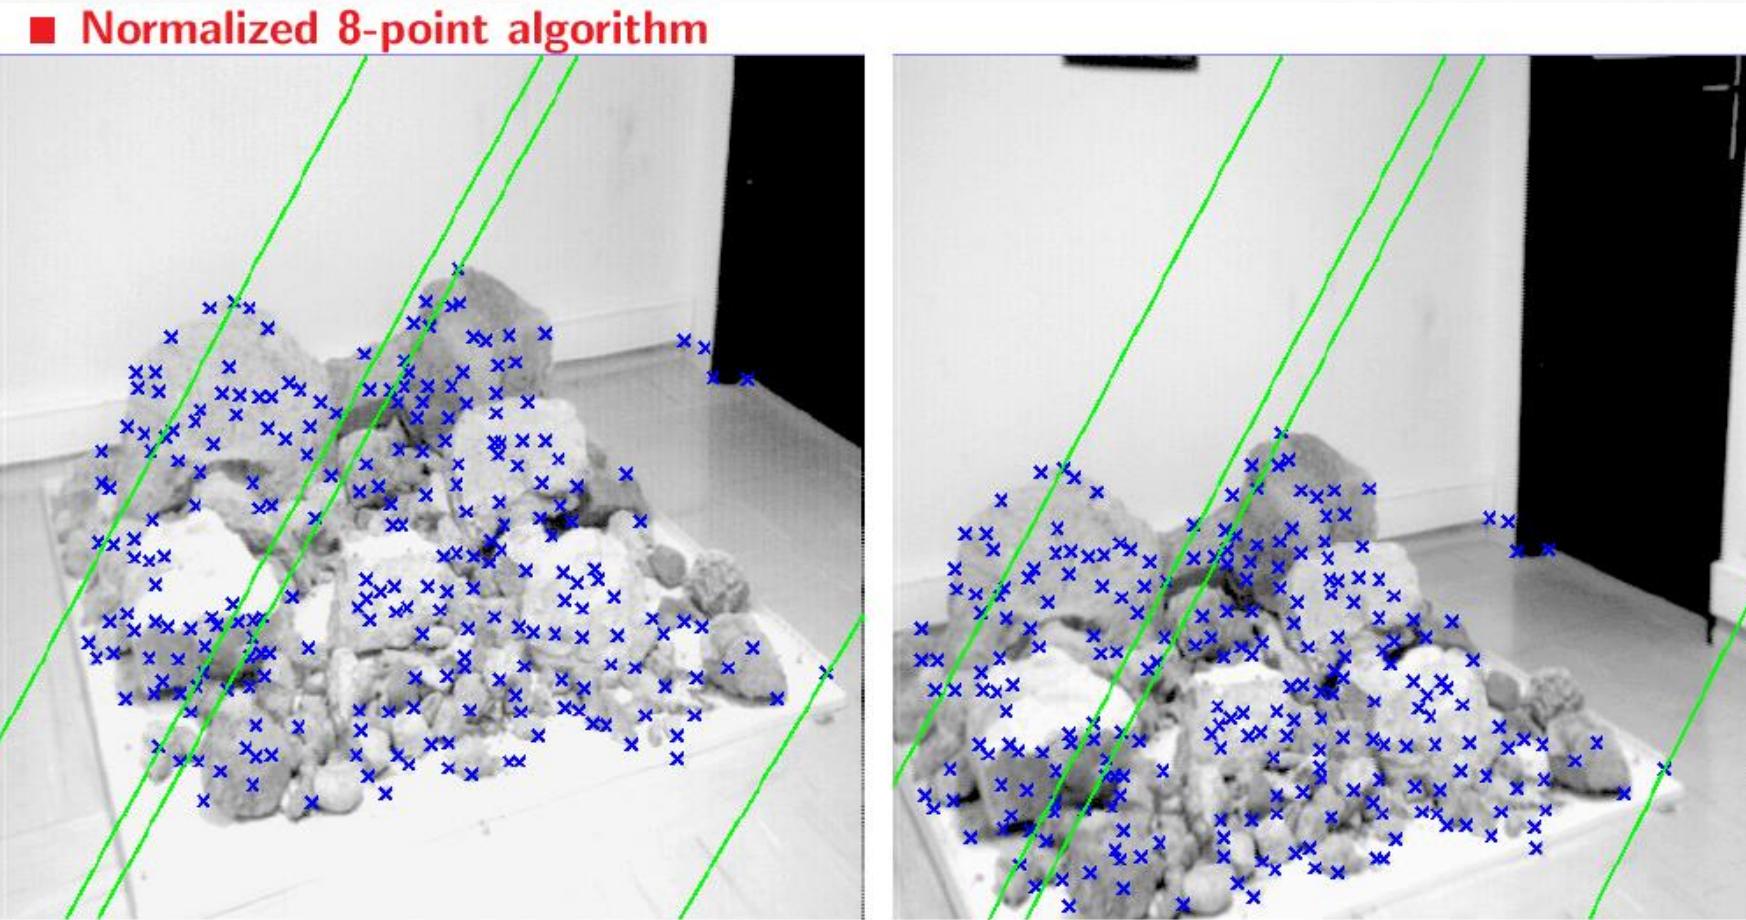


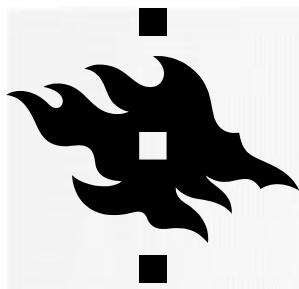
RESULTS (8-POINT ALGORITHM)





RESULTS (NORMALIZED 8-POINT ALGORITHM)



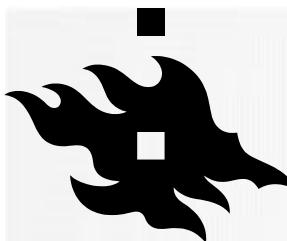


WHAT ABOUT MORE THAN TWO VIEWS?

The geometry of three views is described by a $3 \times 3 \times 3$ tensor called the *trifocal tensor*

The geometry of four views is described by a $3 \times 3 \times 3 \times 3$ tensor called the *quadrifocal tensor*

After this it starts to get complicated...



SOLVING THE SCALE PROBLEM

■ Solutions for finding scale (Z) :

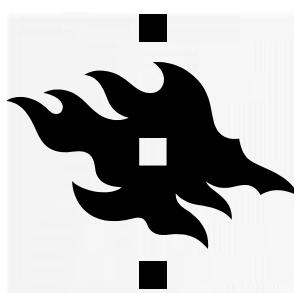
LiDars, RGB-D cameras => include Z

Stereo camera => Known pose of the cameras give Z by triangulation
: **length of baseline important**

Monocular camera

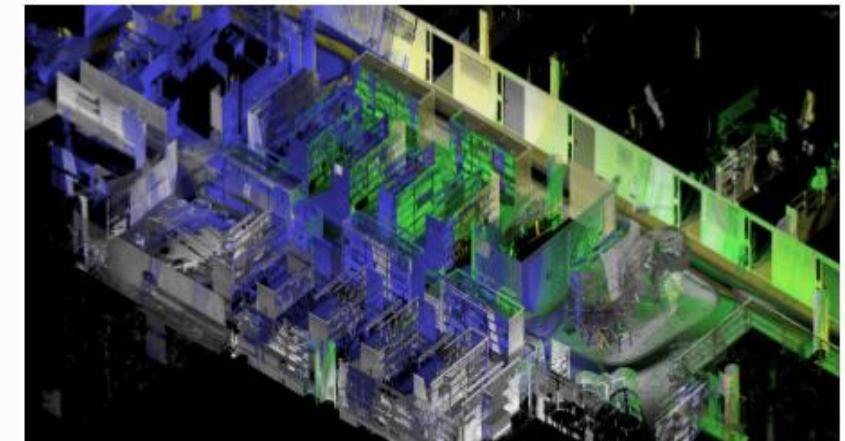
- Objects with known size: **needs a priori preparation and is restricted to one area**
- Camera facing down with known height : **problematic indoors due to uniform texture of the floor**
- Estimation, this we'll look at when talking about SLAM

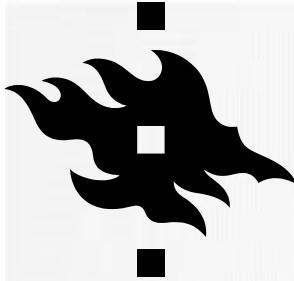




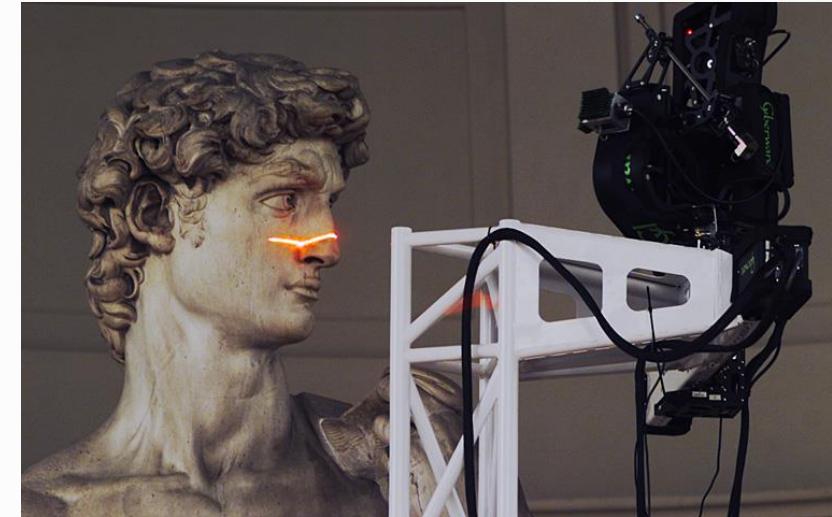
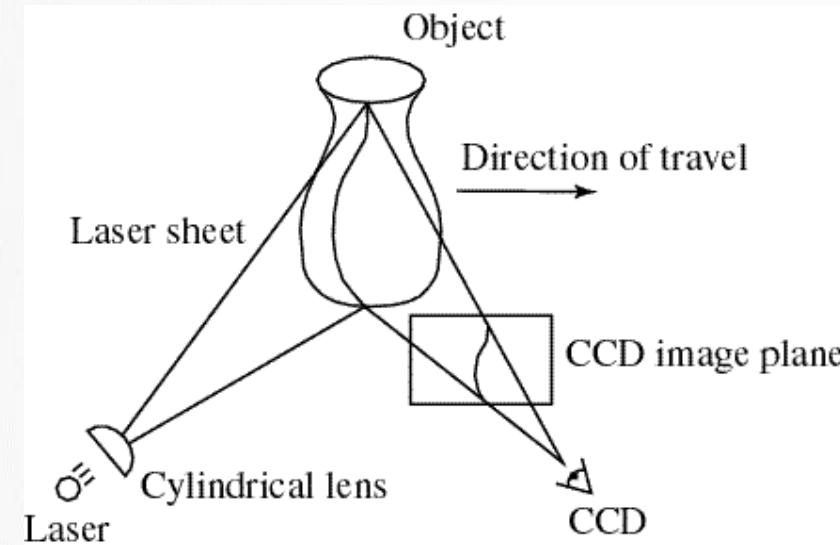
LIGHT DETECTION AND RANGING (LIDAR)

- Controlled steering of laser beams followed by a distance measurement (time it took for the signal to bounce and get back to the system at every pointing direction => 3D model of objects)
- Doesn't require light
- Traditionally expensive and large => technology evolving fast
- Materials of surroundings matter, e.g. glass lets the rays through
- Quite slow to process
- Motion estimation very accurate, cm level accuracy





LASER SCANNING



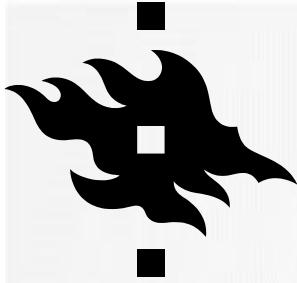
Digital Michelangelo Project
<http://graphics.stanford.edu/projects/mich/>

Optical triangulation

Project a single stripe of laser light

Scan it across the surface of the object

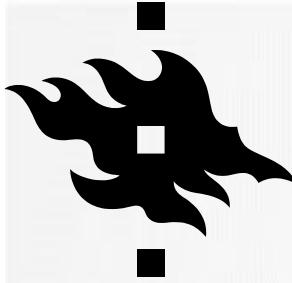
This is a very precise version of structured light scanning



LASER SCANNED MODELS



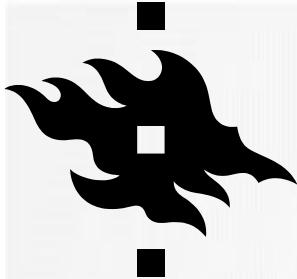
The Digital Michelangelo Project, Levoy et al.



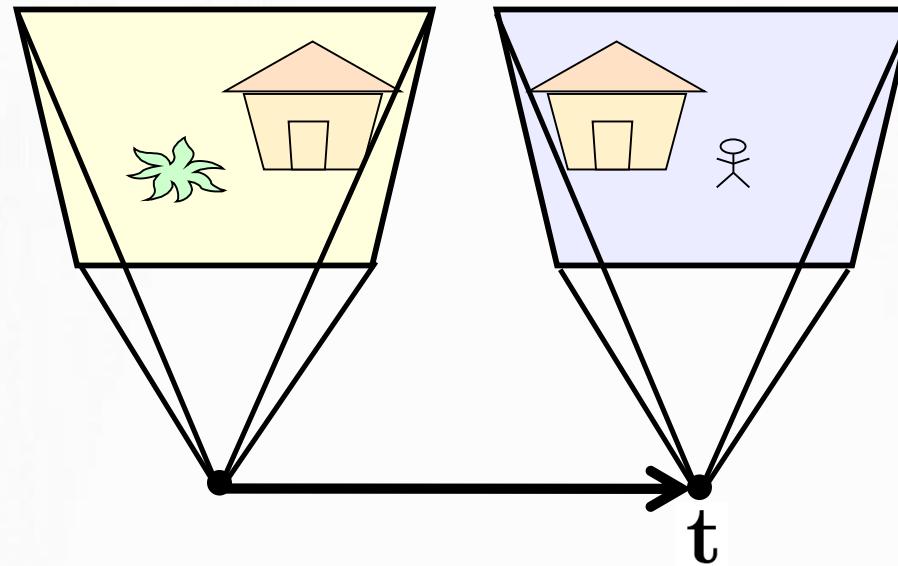
RGB-D CAMERAS

- Combine color information with per pixel depth information using infrared sensors
- Limitations due to field of view, max object distance from sensor, huge amount of data
- Inexpensive, small and light cameras have entered the market





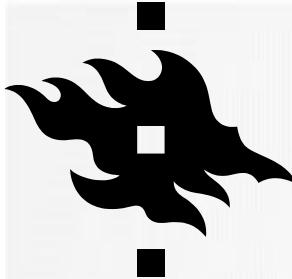
RECTIFIED CASE



$$\mathbf{R} = \mathbf{I}_{3 \times 3}$$

$$\mathbf{t} = [\begin{array}{ccc} 1 & 0 & 0 \end{array}]^T$$

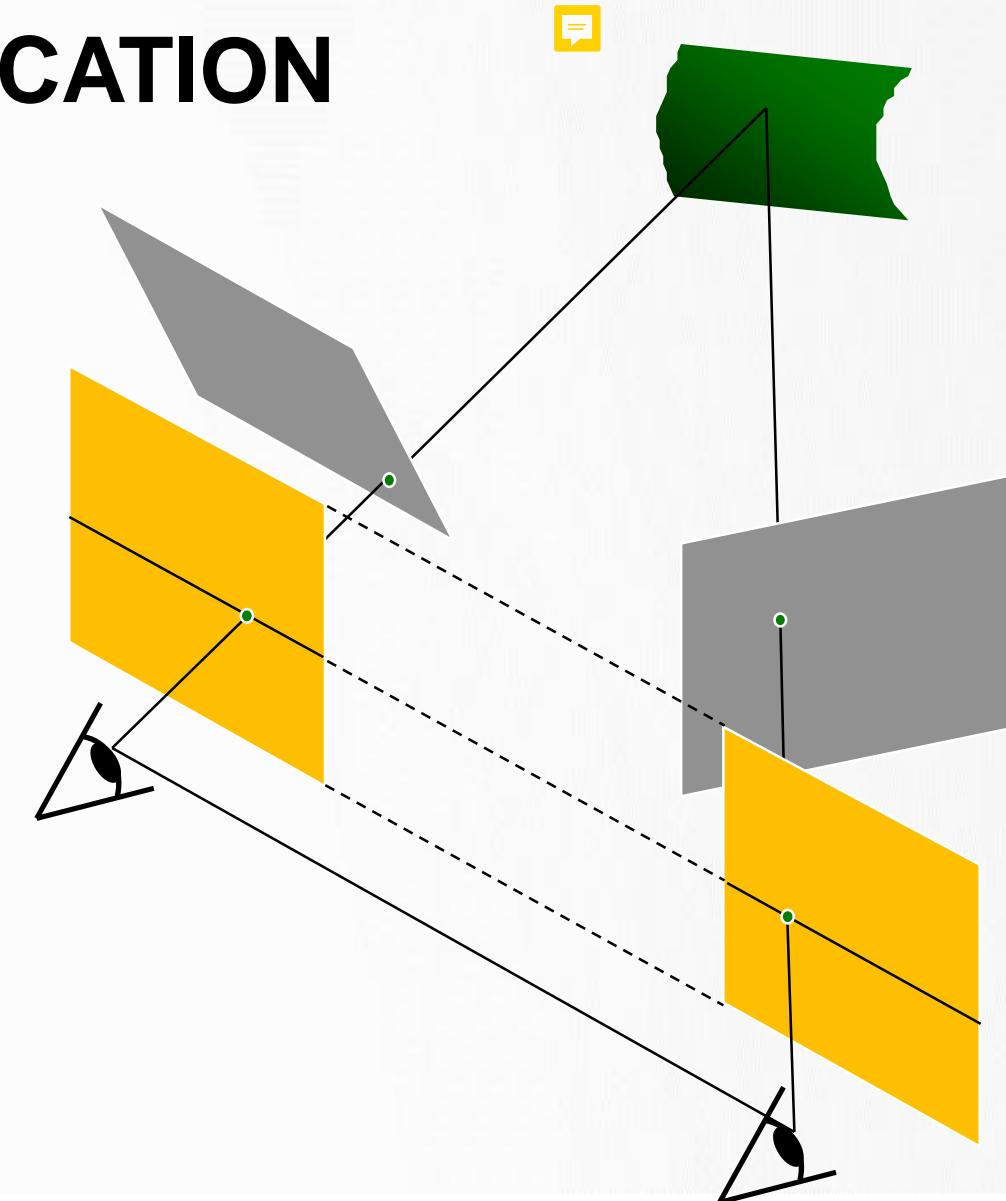
$$\mathbf{E} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{bmatrix}$$

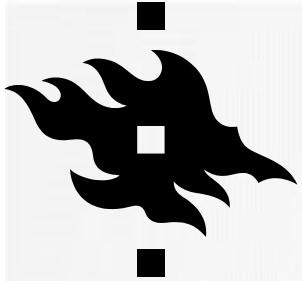


STEREO IMAGE RECTIFICATION

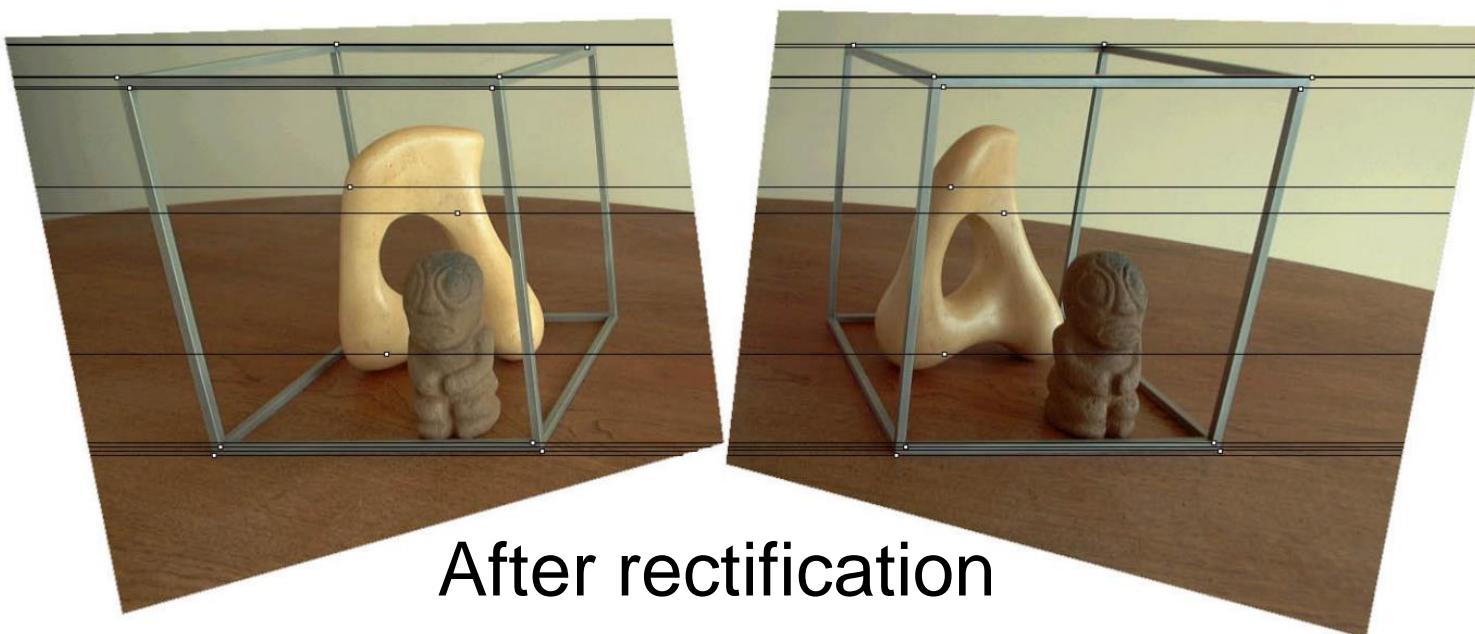


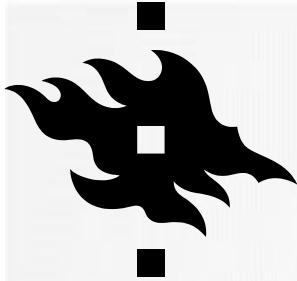
- reproject image planes onto a common plane parallel to the line between optical centers
 - pixel motion is horizontal after this transformation
 - two homographies (3×3 transform), one for each input image reprojection
- C. Loop and Z. Zhang. [Computing Rectifying Homographies for Stereo Vision](#). IEEE Conf. Computer Vision and Pattern Recognition, 1999.



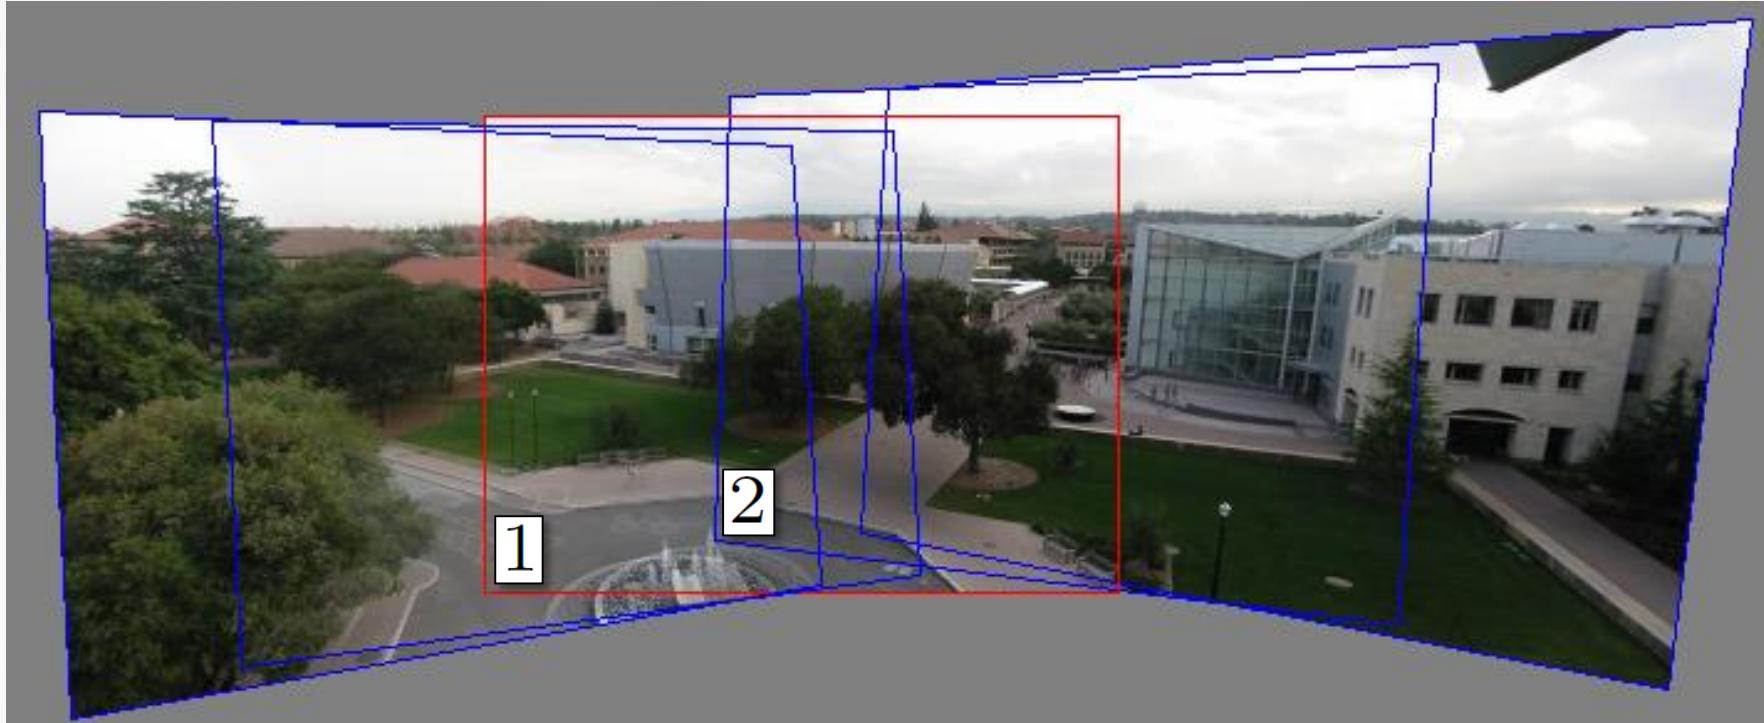


Original stereo pair

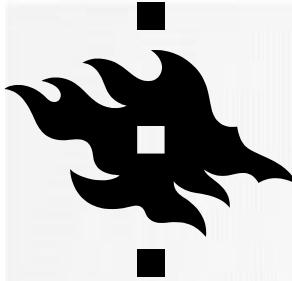




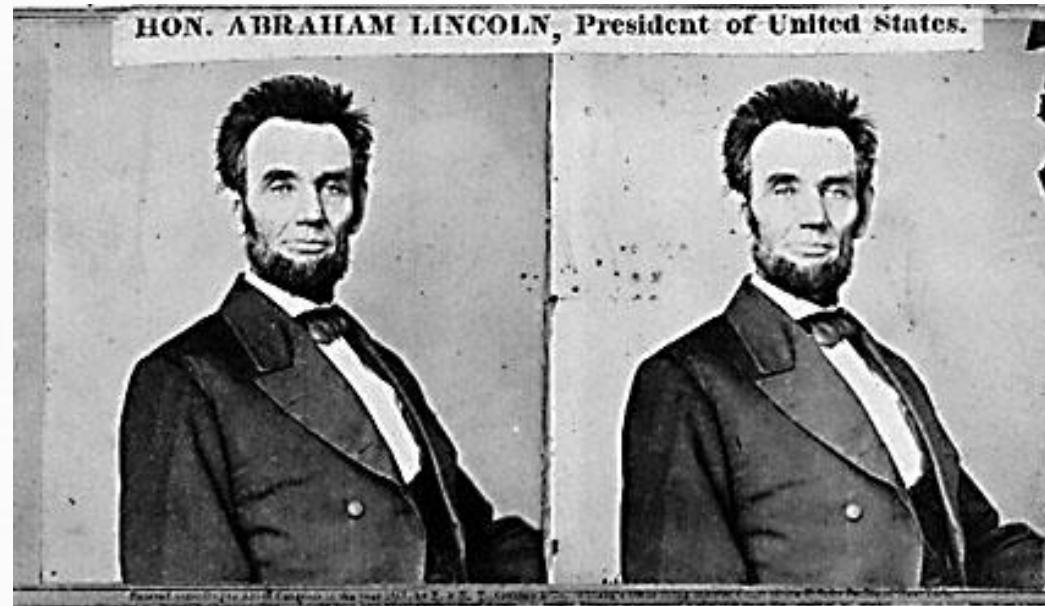
RELATIONSHIP WITH HOMOGRAPHY?



Images taken from the same center of projection? Use a homography!



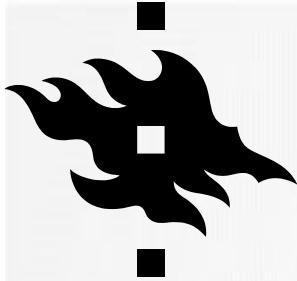
STEREO



Given two images from different viewpoints

How can we compute the depth of each point in the image?

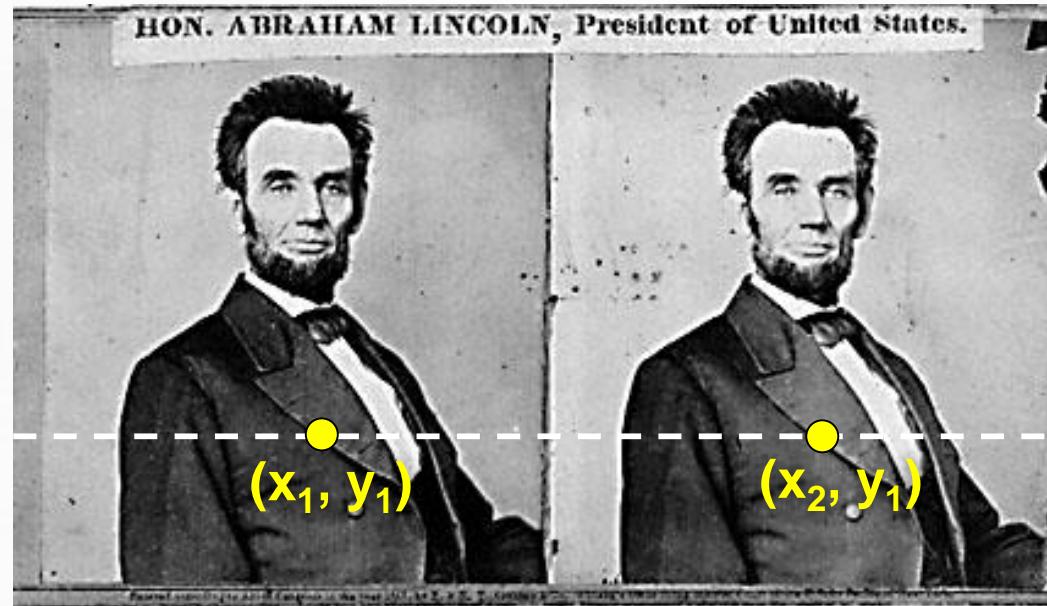
Based on *how much each pixel moves* between the two images



EPIPOLAR GEOMETRY

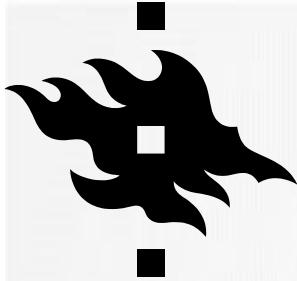


*epipolar
lines*

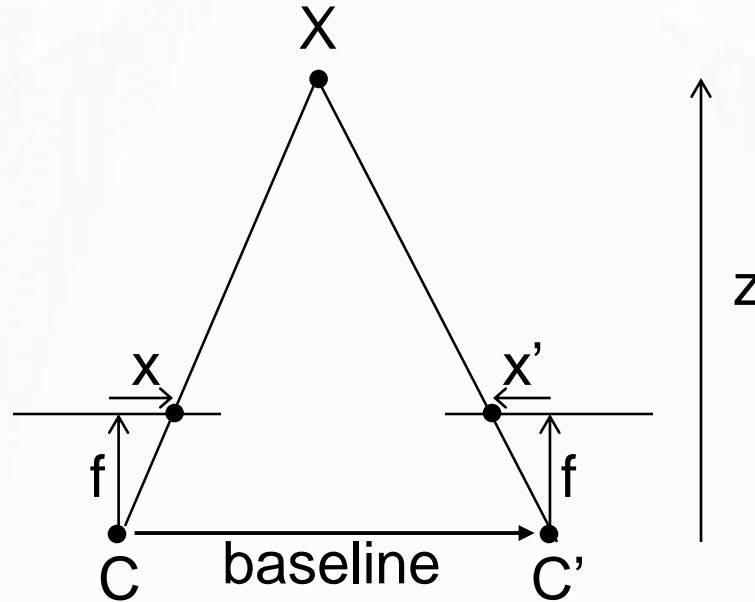


Two images captured by a purely horizontal translating camera
(*rectified* stereo pair)

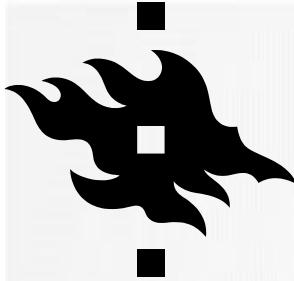
$x_2 - x_1 = \text{the } \textbf{\textit{disparity}} \text{ of pixel } (x_1, y_1)$



DEPTH FROM DISPARITY



$$disparity = x - x' = \frac{baseline * f}{z}$$



BASIC STEREO MATCHING ALGORITHM

Match Pixels in Conjugate Epipolar Lines

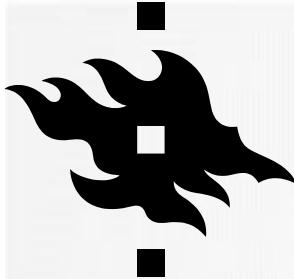


Assume brightness constancy

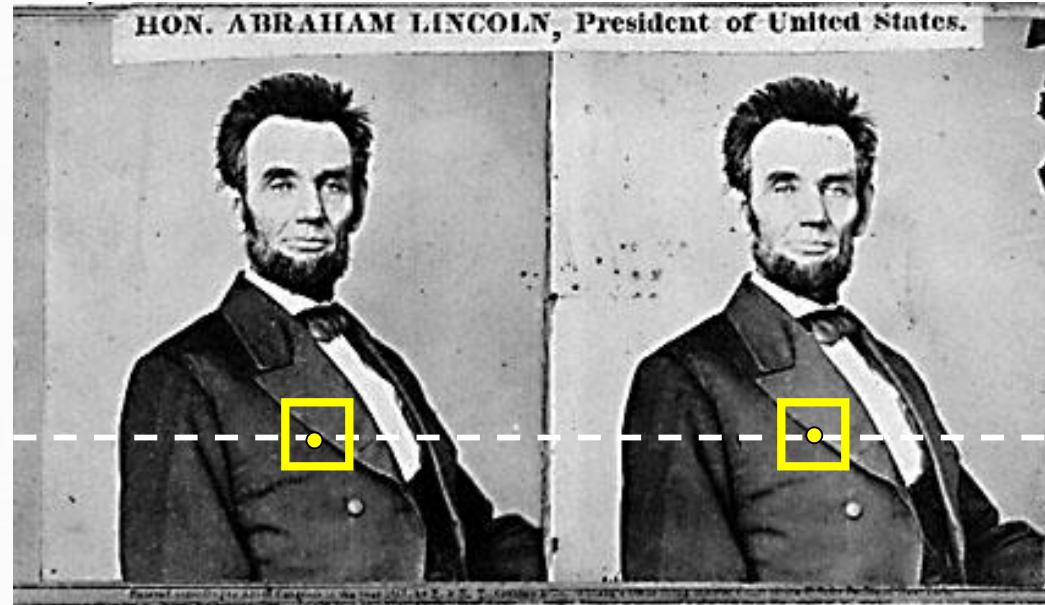
This is a challenging problem

Hundreds of approaches

– A good survey and evaluation: <http://www.middlebury.edu/stereo/>



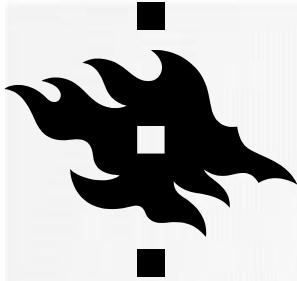
BASIC STEREO ALGORITHM



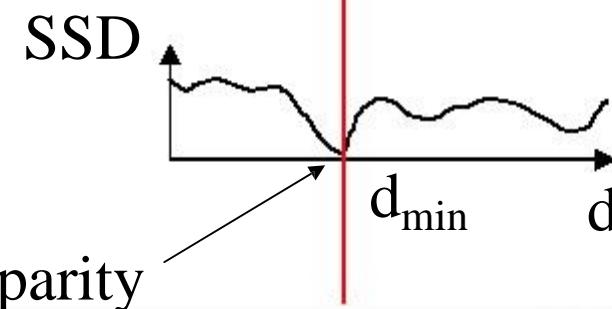
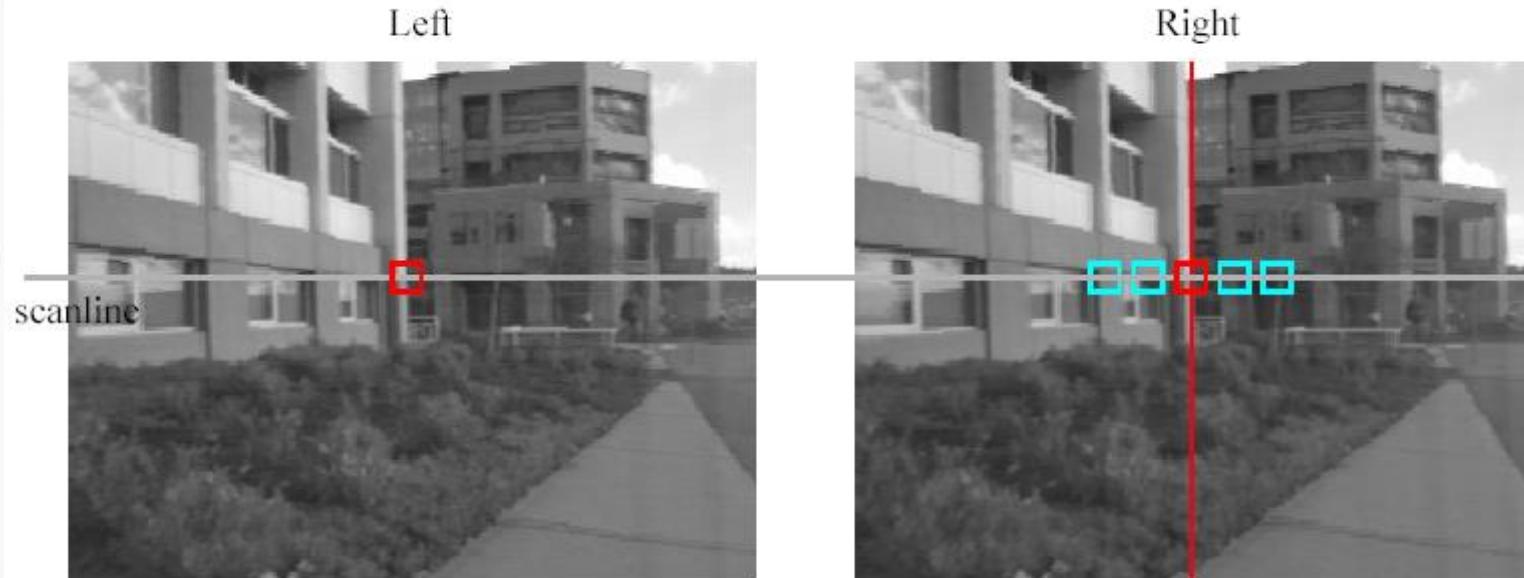
For each epipolar line

For each pixel in the left image

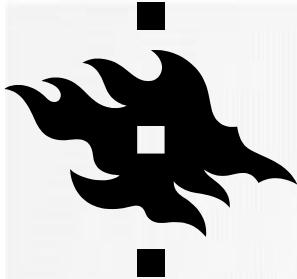
- compare with every pixel on same epipolar line in right image
- pick pixel with minimum match cost



STEREO MATCHING BASED ON SSD



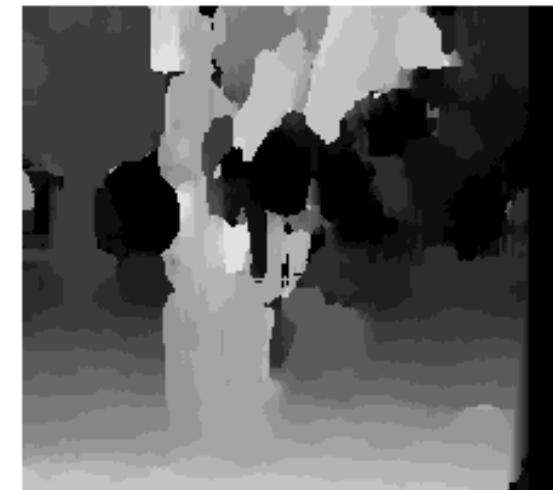
Best matching disparity



WINDOW SIZE



$W = 3$



$W = 20$

Effect of window size
Smaller window

+

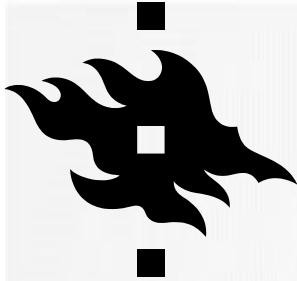
-

Larger window

+

Better results with *adaptive window*

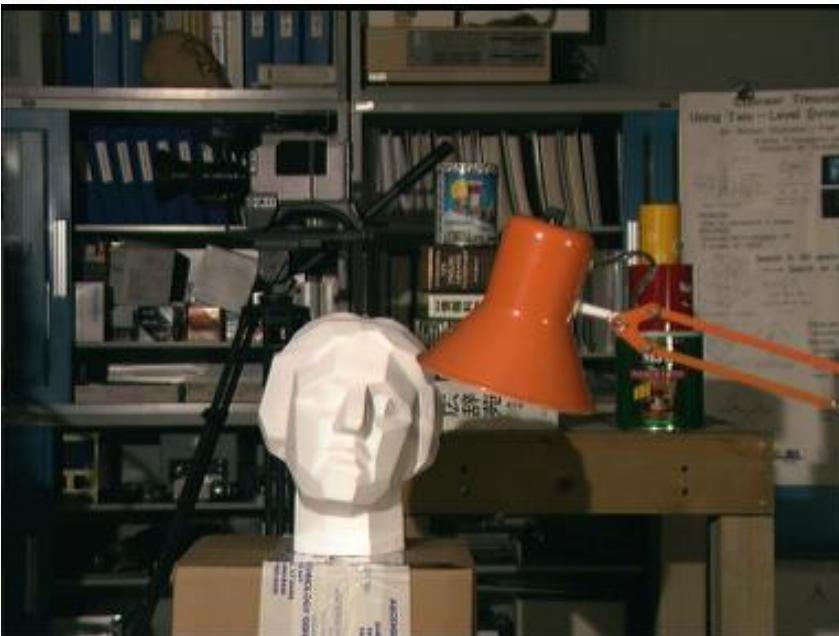
- T. Kanade and M. Okutomi, [A Stereo Matching Algorithm with an Adaptive Window: Theory and Experiment](#), Proc. International Conference on Robotics and Automation, 1991.
- D. Scharstein and R. Szeliski. [Stereo matching with nonlinear diffusion](#). International Journal of Computer Vision, 28(2):155-174, July 1998



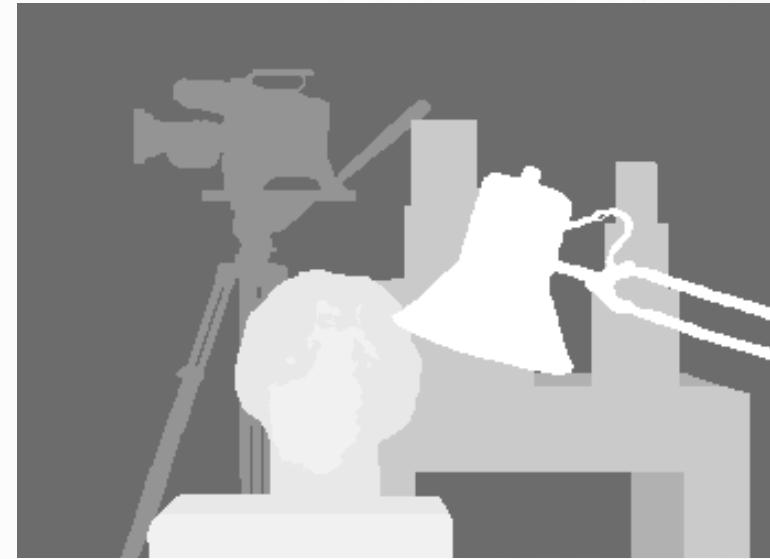
STEREO RESULTS

Data from University of Tsukuba

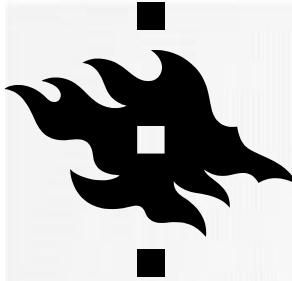
Similar results on other images without ground truth



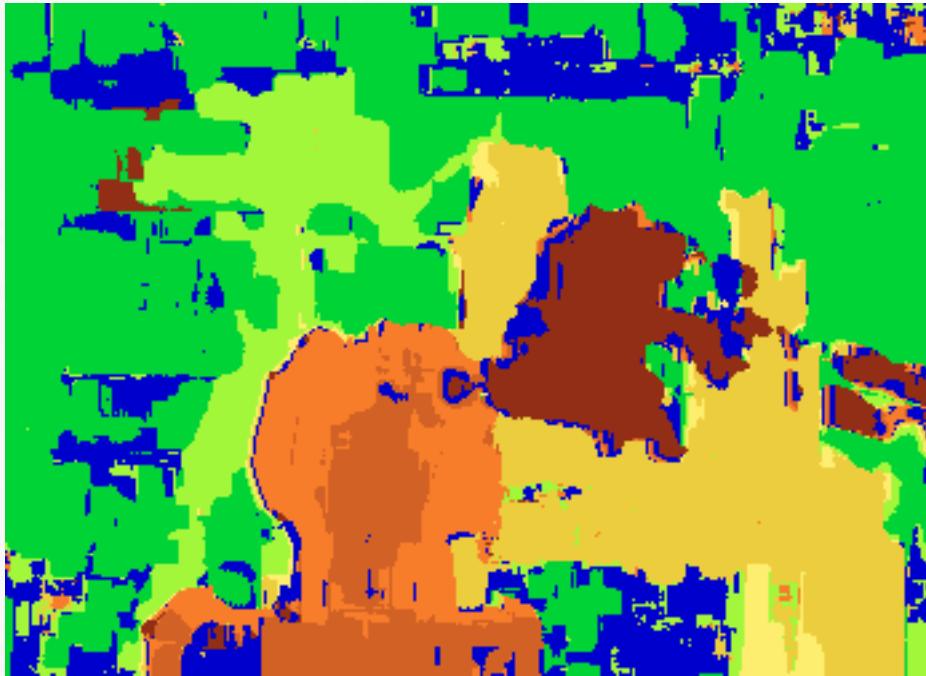
Scene



Ground truth



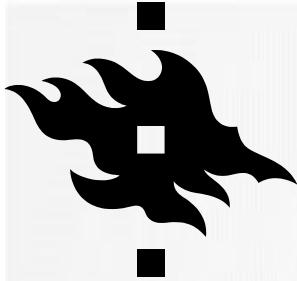
RESULTS WITH WINDOW SEARCH



Window-based matching
(best window size)



Ground truth



BETTER METHODS EXIST...



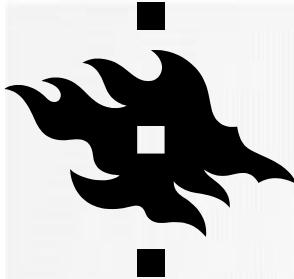
State of the art method

Boykov et al., [Fast Approximate Energy Minimization via Graph Cuts](#),

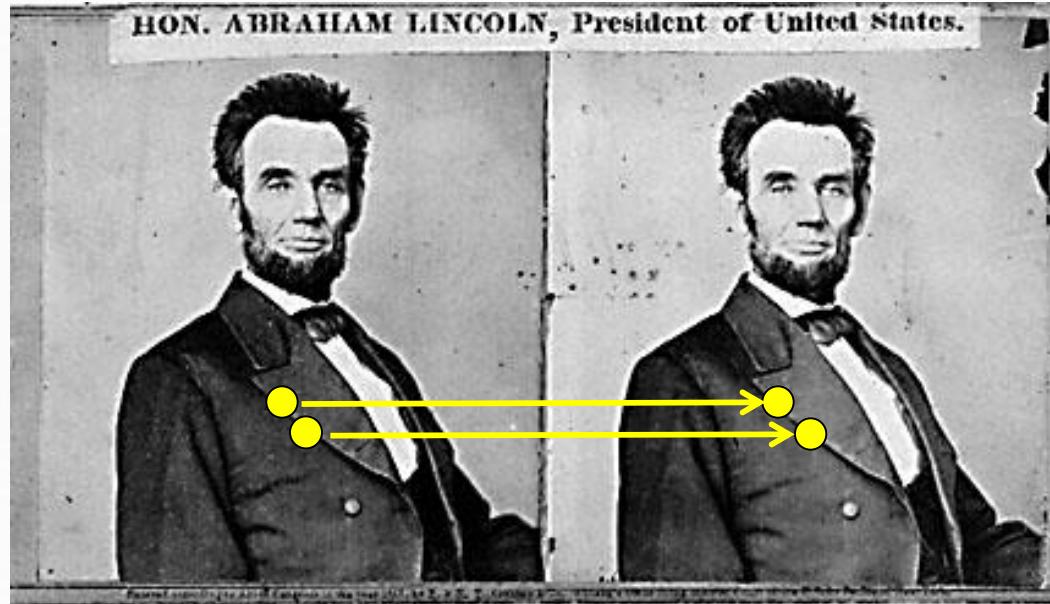
International Conference on Computer Vision, September 1999.



Ground truth

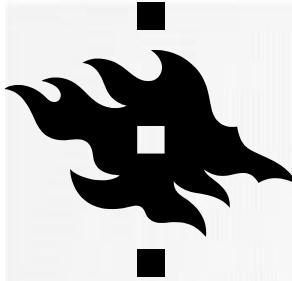


STEREO AS ENERGY MINIMIZATION



What defines a good stereo correspondence?

1. Match quality
 - Want each pixel to find a good match in the other image
2. Smoothness
 - If two pixels are adjacent, they should (usually) move about the same amount



STEREO AS ENERGY MINIMIZATION

- Find disparity map d that minimizes an energy function

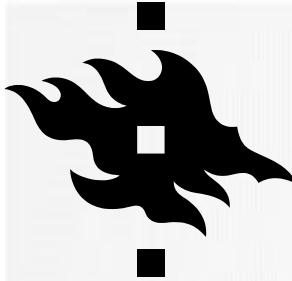
$$E(d)$$

- Simple pixel / window matching

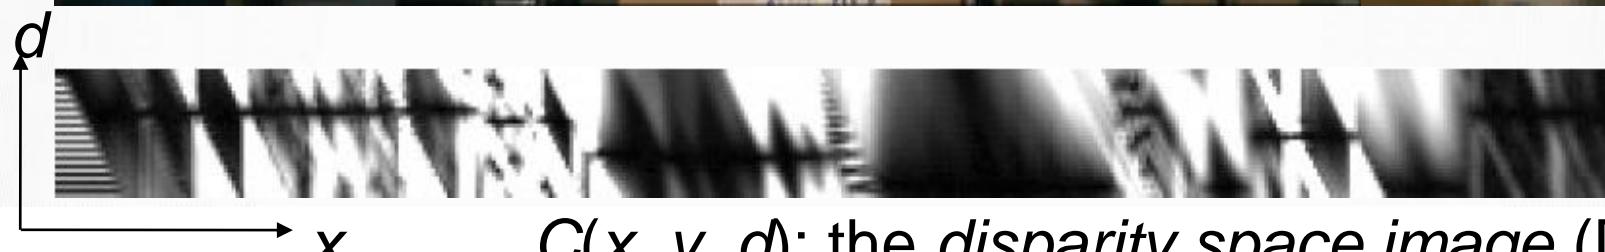
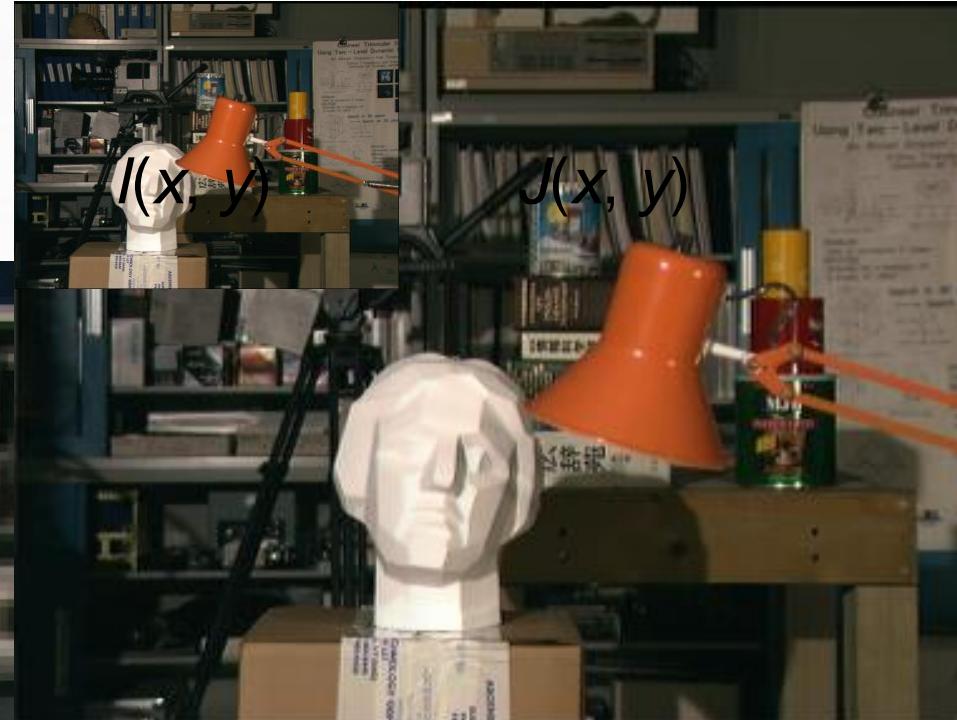
$$\begin{aligned}x' &= x + d(x, y) \\y' &= y\end{aligned}$$

$$E(d) = \sum_{(x,y) \in I} C(x, y, d(x, y))$$

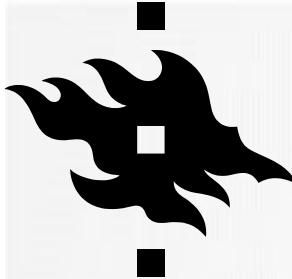
$C(x, y, d(x, y)) =$ SSD distance between windows
 $I(x, y)$ and $J(x + d(x, y), y)$



STEREO AS ENERGY MINIMIZATION



$C(x, y, d)$; the *disparity space image* (DSI)

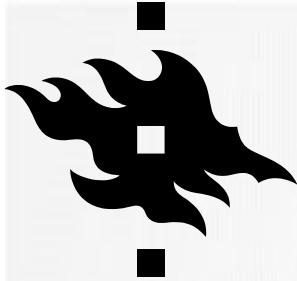


STEREO AS ENERGY MINIMIZATION

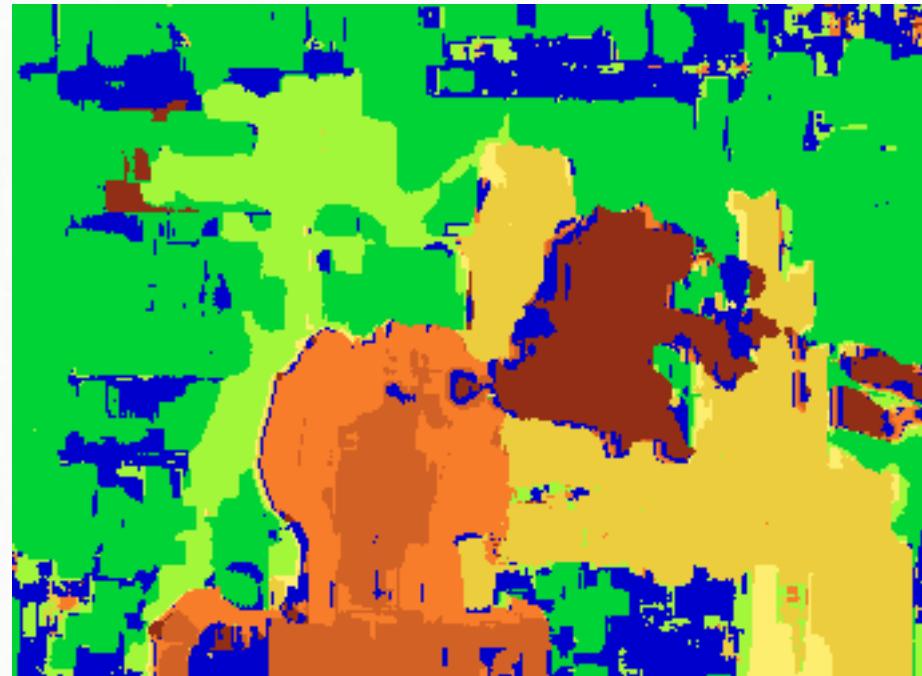


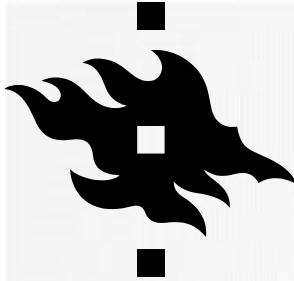
Simple pixel / window matching: choose the minimum of each column in the DSI independently:

$$d(x, y) = \arg \min_{d'} C(x, y, d')$$



GREEDY SELECTION OF BEST MATCH





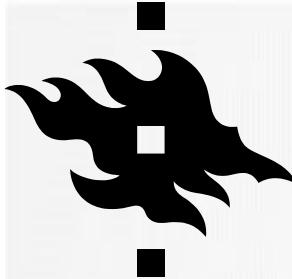
STEREO AS ENERGY MINIMIZATION

Better objective function

$$E(d) = \underbrace{E_d(d)}_{\text{match cost}} + \lambda \underbrace{E_s(d)}_{\text{smoothness cost}}$$

Want each pixel to find a good
match in the other image

Adjacent pixels should
(usually) move about the same
amount



STEREO AS ENERGY MINIMIZATION



$$E(d) = E_d(d) + \lambda E_s(d)$$

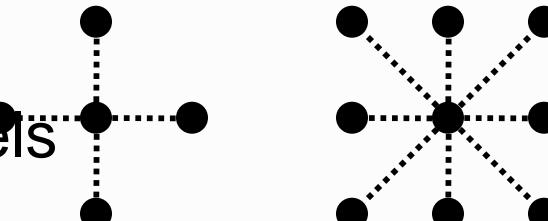
match cost:

$$E_d(d) = \sum_{(x,y) \in I} C(x, y, d(x, y))$$

smoothness
cost:

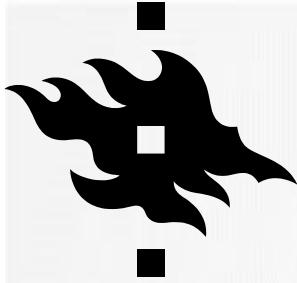
$$E_s(d) = \sum_{(p,q) \in \mathcal{E}} V(d_p, d_q)$$

\mathcal{E} : set of neighboring pixels



4-connected
neighborhood

8-connected
neighborhood



SMOOTHNESS COST

$$E_s(d) = \sum_{(p,q) \in \mathcal{E}} V(d_p, d_q)$$

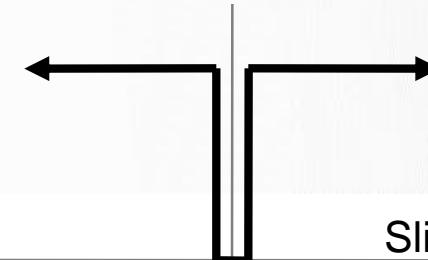
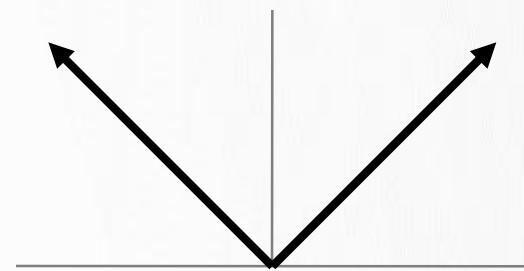
How do we choose V ?

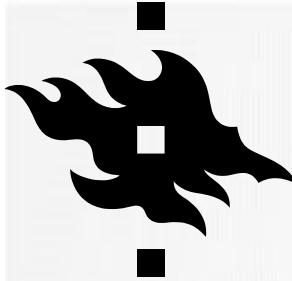
$$V(d_p, d_q) = |d_p - d_q|$$

L_1 distance

$$V(d_p, d_q) = \begin{cases} 0 & \text{if } d_p = d_q \\ 1 & \text{if } d_p \neq d_q \end{cases}$$

“Potts model”





STEREO AS A MINIMIZATION PROBLEM

$$E(d) = E_d(d) + \lambda E_s(d)$$

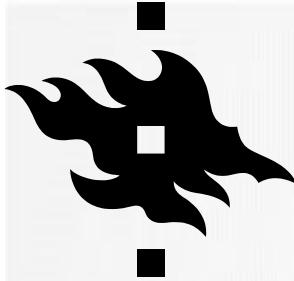
The 2D problem has many local minima

Gradient descent doesn't work well

And a large search space

$n \times m$ image w/ k disparities has k^{nm} possible solutions

Finding the global minimum is NP-hard in general



DYNAMIC PROGRAMMING

$$E(d) = E_d(d) + \lambda E_s(d)$$

Can minimize this independently per scanline using dynamic programming (DP)

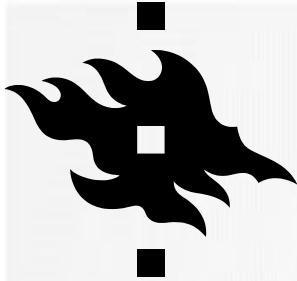
Basic idea: incrementally build a table of costs D one column at a time



$D(x, y, i)$: minimum cost of solution such that $d(x, y) = i$

Base case: $D(0, y, i) = C(0, y, i), i = 0, \dots, L$ (L = max disparity)

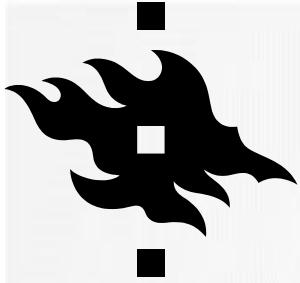
Recurrence: $D(x, y, i) = C(x, y, i) + \min_{j \in \{0, 1, \dots, L\}} D(x - 1, y, j) + \lambda|i - j|$



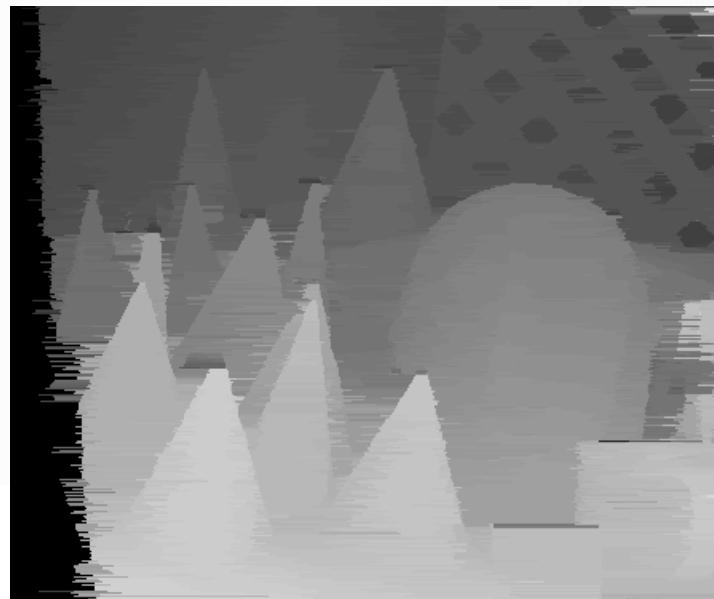
DYNAMIC PROGRAMMING

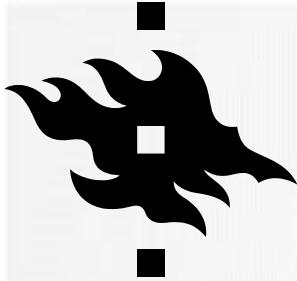


Finds “smooth”, low-cost path through DPI from left to right



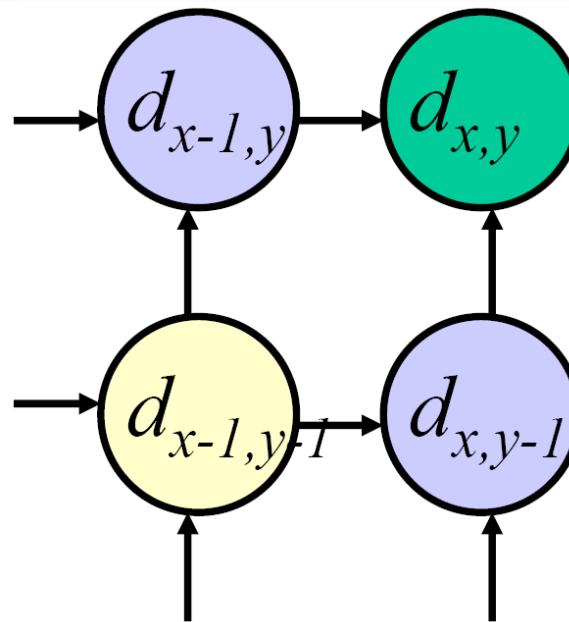
DYNAMIC PROGRAMMING

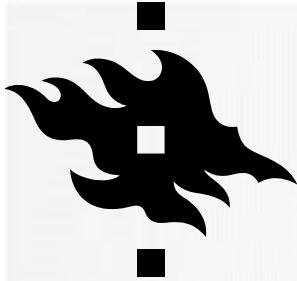




DYNAMIC PROGRAMMING

Can we apply this trick in 2D as well?





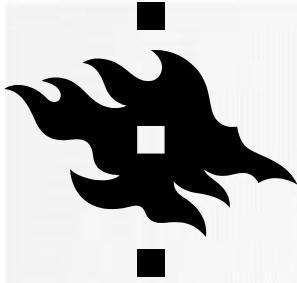
REAL-TIME STEREO



[Nomad robot](#) searches for meteorites in Antarctica
<http://www.frc.ri.cmu.edu/projects/meteorobot/index.html>

Used for robot navigation (and other tasks)

Several real-time stereo techniques have been developed (most based on simple discrete search)



STEREO RECONSTRUCTION PIPELINE

Steps



Calibrate cameras

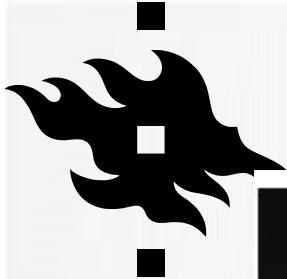
Rectify images

Compute disparity

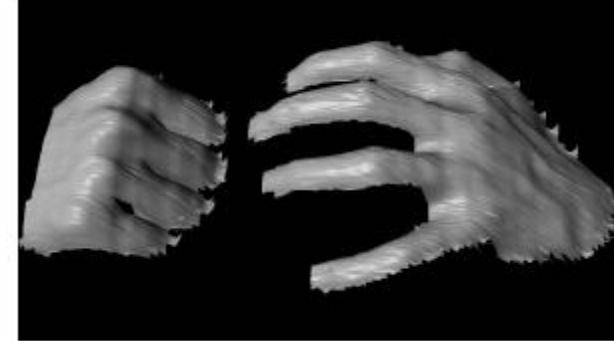
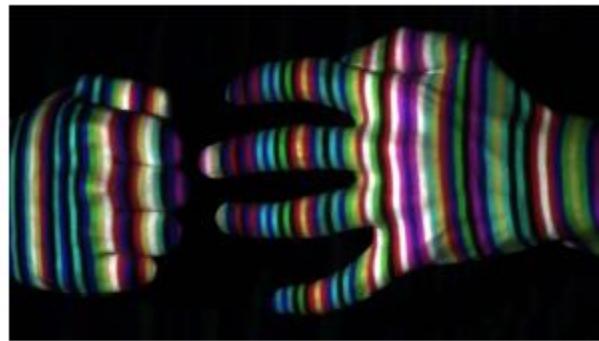
Estimate depth

What will cause errors?

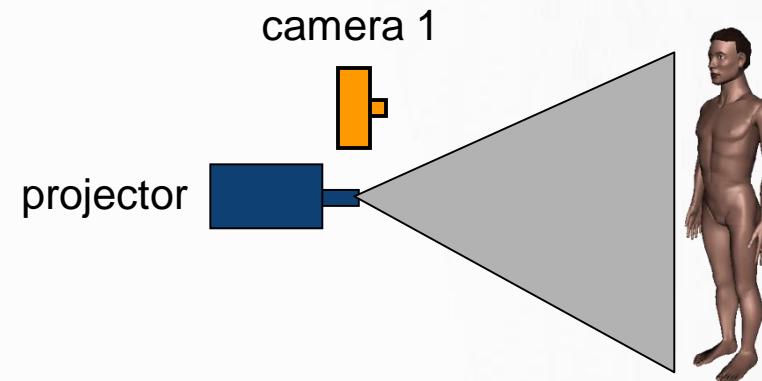
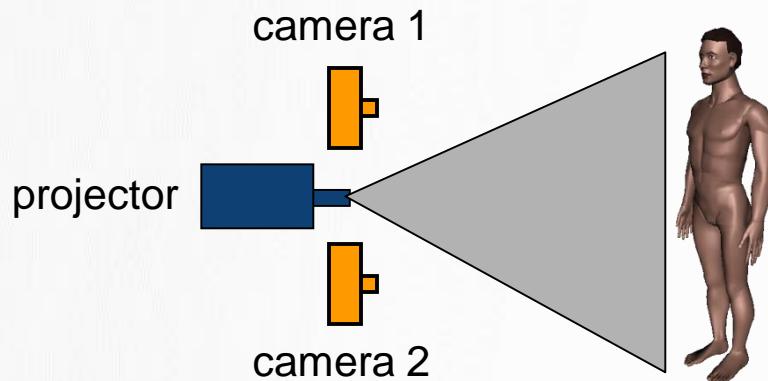
- Camera calibration errors
- Poor image resolution
- Occlusions
- Violations of brightness constancy (specular reflections)
- Large motions
- **Low-contrast image regions**



ACTIVE STEREO WITH STRUCTURED LIGHT

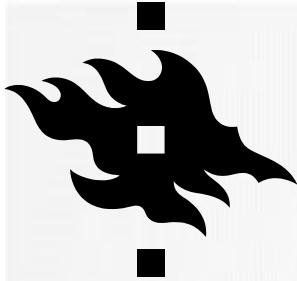


Li Zhang's one-shot stereo



Project “structured” light patterns onto the object
simplifies the correspondence problem

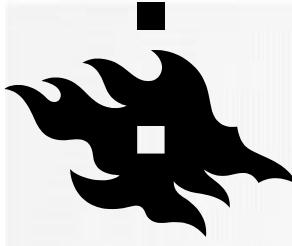
basis for active depth sensors, such as Kinect and iPhone X (using IR)



ACTIVE STEREO WITH STRUCTURED LIGHT

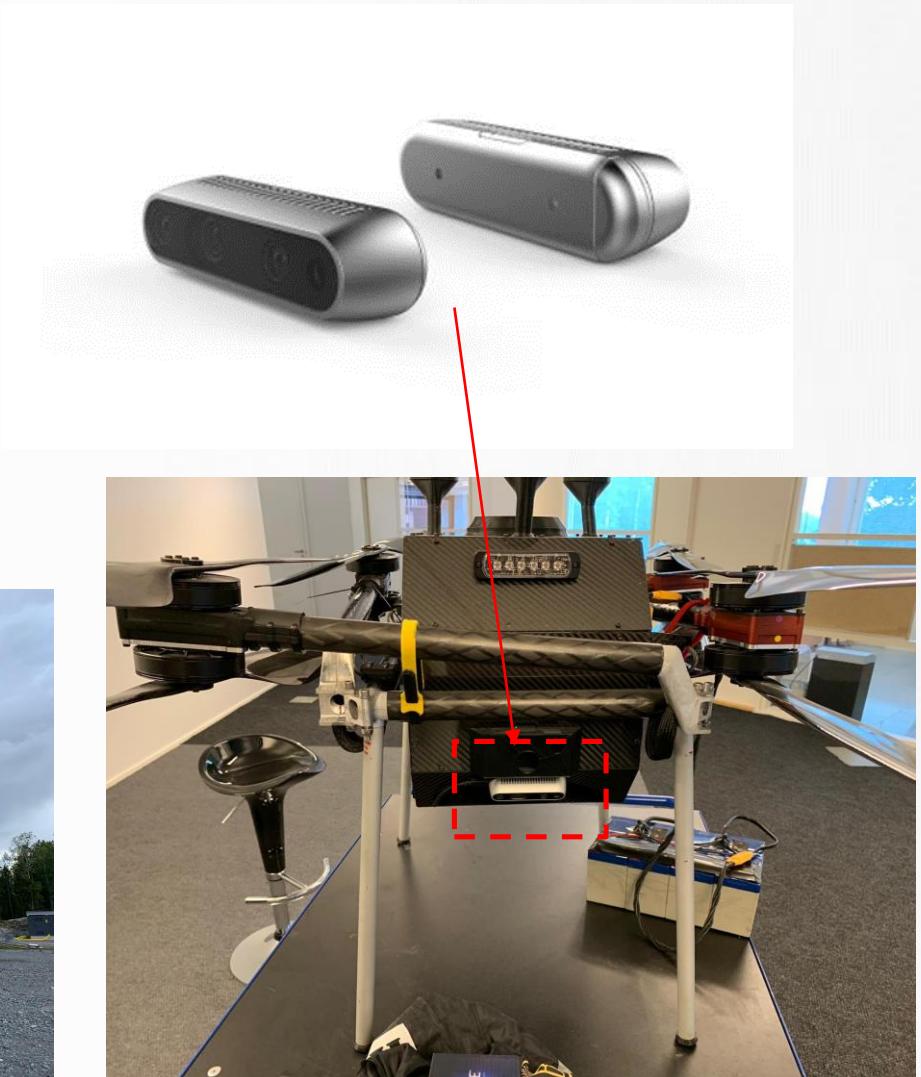


<https://ios.gadgethacks.com/news/watch-iphone-xs-30k-ir-dots-scan-your-face-0180944/>



INTEL REALSENSE

- RealSense cameras
 - Stereoscopic, subpixel disparity accuracy
 - Function well **outdoors**
 - Low power consumption, 100x10x4 mm in size
 - Includes 3 cameras
 - Parameters can't be modified, depth solution not optimal





60° 10 1.2 N, 24° 57 18 E