


# COMPUTER VISION SLAM

LECTURE 11 9.10.2019

Laura Ruotsalainen, Associate Professor  
Department of Computer Science



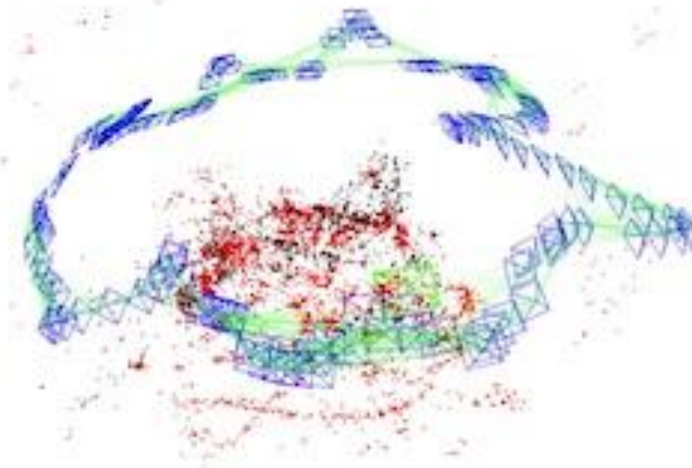
# VISUAL SLAM

- Feature based Visual SLAM
  - Tracking features, e.g. SIFT, ORB
- Direct Visual SLAM
  - Uses image pixels directly
- Pipeline
  - Tracking (computing the motion and updating positions based on that) between consecutive images
  - Mapping
  - Global optimization (loop-closure and refining map based on that)
  - Relocalization





## Feature-Based SLAM




Minimize **Feature Reprojection Error**

**Sparse** Reconstruction

e.g. ORB-SLAM

## Direct SLAM



Minimize **Photometric Error**

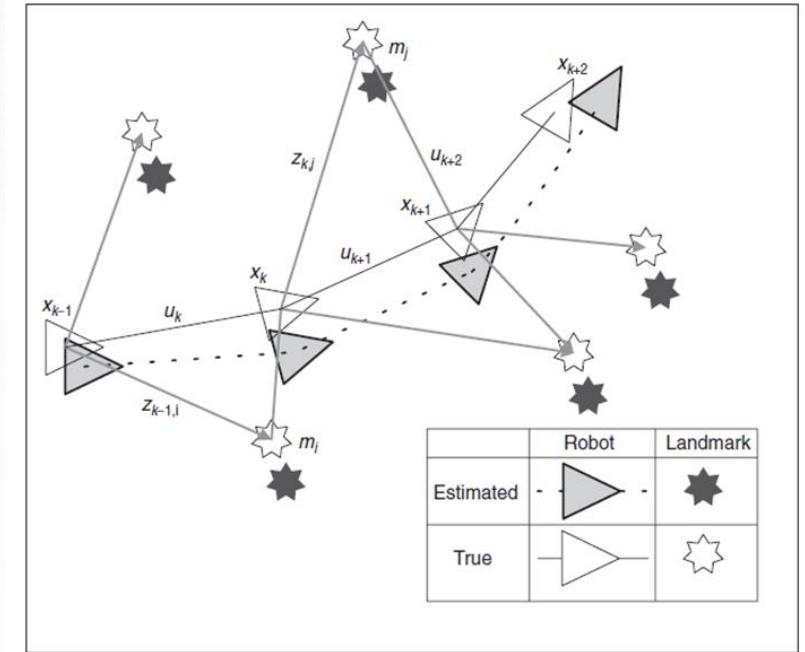
**Semi Dense / Dense** Reconstruction

e.g. LSD-SLAM



# SLAM - PRINCIPLES

- The user does not know its location state vector  $x_k$
- There is no map  $M$  of the area
- $m$  is the location of  $i$ :th landmark,  
 $Z$  all landmark observations
- $U$  control inputs
- In earlier developments the probabilistic form probability function  $p(x_k, m | Z_{0:k}, U_{0:k}, x_0)$  was solved using EKF or Particle filtering





# SLAM'S VARIANTS



- Feature based vs direct slam
  - Direct SLAM provides dense, more representative map
  - Feature-based methods are more invariant to viewpoint and illumination changes
  - Feature-based methods are less affected by dynamic objects
- Filtering vs keyframe-based
  - Monocular SLAMs are either filter based (e.g. Kalman filtering) or
  - Keyframe-based where they optimize the motion and map simultaneously



# TEMPORAL STATE MODELS





Represent the 'world' as a set of random variables  **$X$**

$X = \{x, y\}$       location on the ground plane

$X = \{x, y, z\}$       position in the 3D world

$X = \{x, \dot{x}\}$       position and velocity


$X = \{x, \dot{x}, f_1, \dots, f_n\}$       position, velocity and  
location of landmarks



## Object tracking (localization)

$$\mathbf{X} = \{\mathbf{x}, \mathbf{y}\}$$


e.g., location on the ground plane



## Object location and world landmarks (localization and mapping)

$$\mathbf{X} = \{\mathbf{x}, \dot{\mathbf{x}}, \mathbf{f}_1, \dots, \mathbf{f}_n\}$$

e.g., position and velocity of robot  
and location of landmarks







$X_t$



The state of the world changes over time



$$X_t$$



The state of the world changes over time

So we use a sequence of random variables:

$$X_0, X_1, \dots, X_t$$


$$\mathbf{X}_t$$


The state of the world changes over time

So we use a sequence of random variables:

$$\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t$$

The state of the world is usually **uncertain** so we think in terms of a distribution

$$P(\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t)$$

*How big is the space of this distribution?*





When we use a sensor (camera),  
we don't have direct access to the state but noisy  
observations of the state

$$\mathbf{E}_t$$

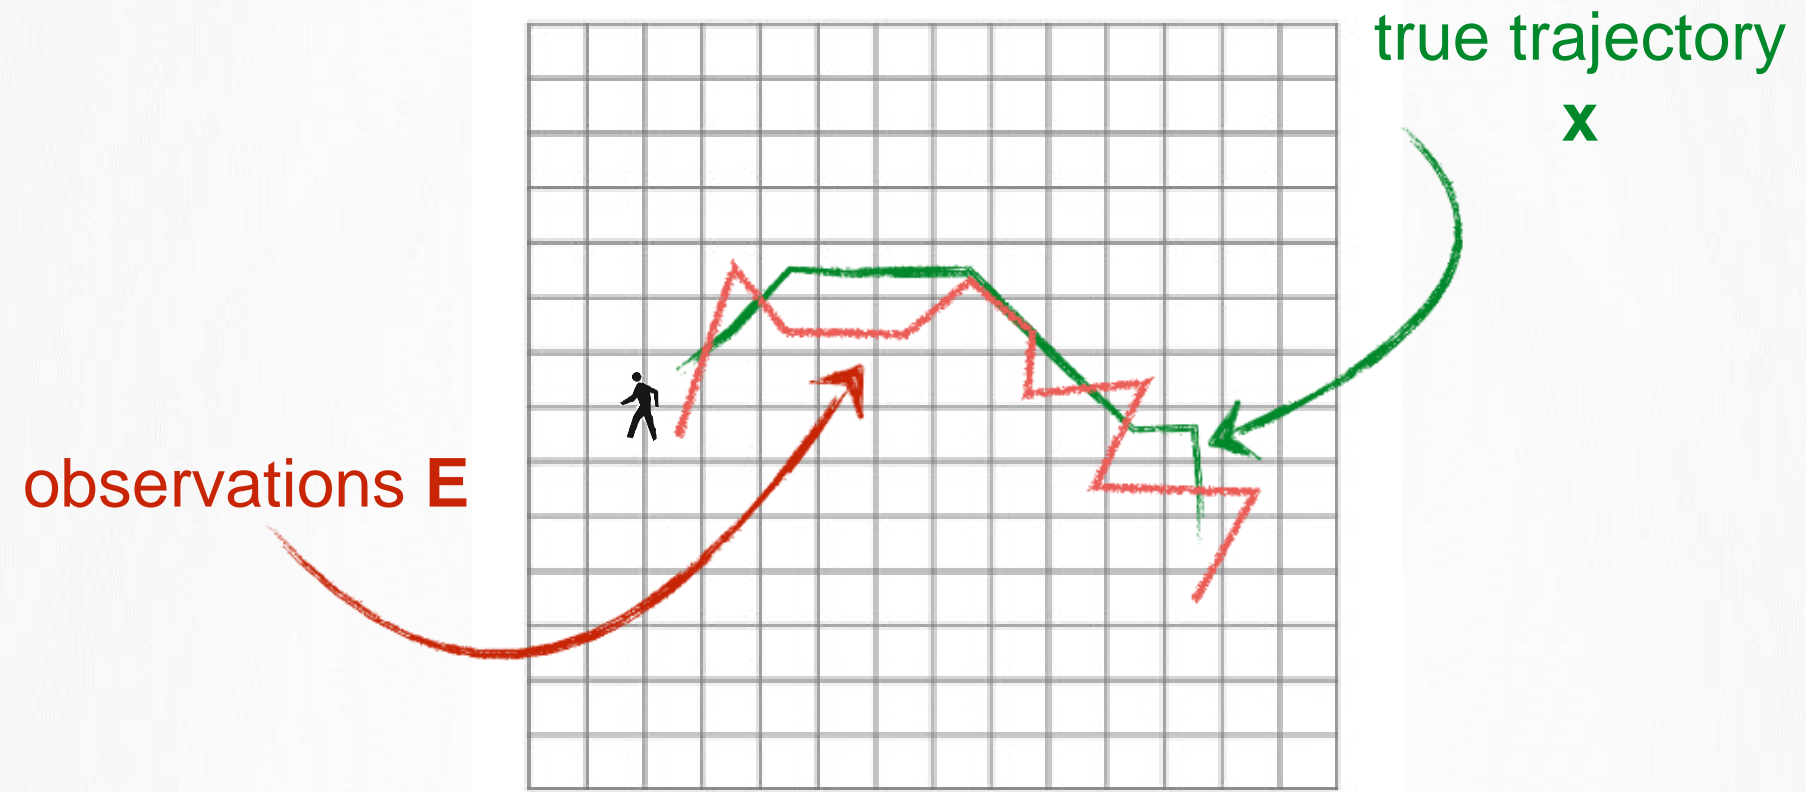
$$\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t, \mathbf{E}_1, \mathbf{E}_2, \dots, \mathbf{E}_t$$

(**all** possible ways of observing **all** possible trajectories)

*How big is the space of this distribution?*



all possible ways of observing all possible trajectories of length  $t$







So we think of the world in terms of the distribution

$$P(\underbrace{X_0, X_1, \dots, X_t}_{\text{unobserved variables (hidden state)}}, \underbrace{E_1, E_2, \dots, E_t}_{\text{observed variables (evidence)}})$$





## Reduction 1. Stationary process assumption:

*‘a process of change that is governed by laws that do not themselves change over time.’*

$$P(\mathbf{E}_t | \mathbf{X}_t) = P_t(\mathbf{E}_t | \mathbf{X}_t)$$

the model doesn't change over time



## Reduction 1. Stationary process assumption:


*‘a process of change that is governed by laws that do not themselves change over time.’*

$$P(\mathbf{E}_t | \mathbf{X}_t) = P_t(\mathbf{E}_t | \mathbf{X}_t)$$



the model doesn't change over time


Only have to store **one** model.




## Reduction 2. Markov Assumption:


*'the current state only depends on a finite history of previous states.'*

First-order Markov Model:  $P(\mathbf{X}_t | \mathbf{X}_{t-1})$ .



Second-order Markov Model:  $P(\mathbf{X}_t | \mathbf{X}_{t-1}, \mathbf{X}_{t-2})$





## Reduction 2. Markov Assumption:


*‘the current observation only depends on current state.’*

The current observation is usually  
most influenced by the current state


$$P(\mathbf{E}_t | \mathbf{X}_t)$$

(this relationship is called the **observation** model)

*Can you think of an observation of a state?*



For example, GPS is a noisy observation of location.




But GPS tells you that you are here  
with probability  $P(\mathbf{E}_t | \mathbf{X}_t)$





## Reduction 3. Prior State Assumption:

*‘we know where the process (probably) starts’*





# Joint Probability of a Temporal Sequence

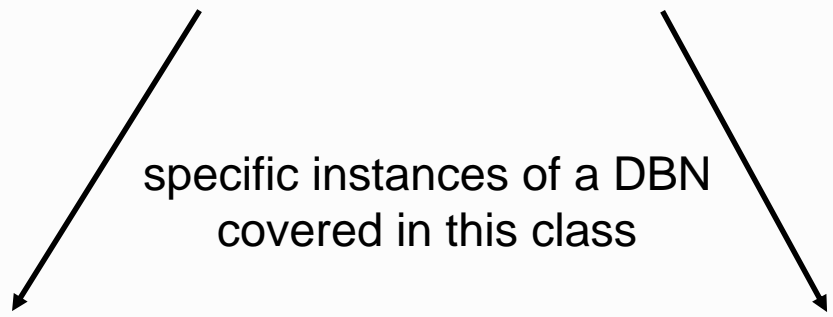
$$P(\mathbf{X}_0) \prod_{t=1}^T P(\mathbf{X}_t | \mathbf{X}_{t-1}) P(\mathbf{E}_t | \mathbf{X}_t)$$

state prior  
prior

motion model  
transition model

sensor model  
observation model

## Joint Distribution for a Dynamic Bayesian Network



### Hidden Markov Model

(typically taught as discrete but not necessarily)

### Kalman Filter

(Gaussian motion model, prior and observation model)



# Filtering

$$P(\mathbf{X}_t | e_{1:t})$$

Posterior probability over the **current** state, given all evidence up to present

## Where am I now?



# KALMAN FILTER 1/2

the Kalman Filter is “the best known filter, a simple and elegant algorithm, as an optimal recursive Bayesian estimator for a somewhat restricted class of linear Gaussian problems” (B. Ristic et al., Beyond the Kalman filter, particle filters for tracking applications, 2004)

Estimate the **state**  $x \in \mathcal{R}^n$ ,

State transition model

$$x_k = Ax_{k-1} + Bu_{k-1} + \varepsilon_{k-1}$$

process noise  
 $p(\varepsilon) \sim N(0, Q)$


using **measurements**  $z \in \mathcal{R}^m$

Optional control input

$$z_k = Hx_k + \delta_k$$

measurement noise  
 $p(\delta) \sim N(0, R)$

Observation model



# KALMAN FILTER 2/2

- Feedback control:

time update  
"predict"

measurement update  
"correct"

$$\hat{x}_k^- = A\hat{x}_{k-1}$$

$$P_k^- = AP_{k-1}A^T + Q$$

*a priori estimate*

P = estimate covariance  
Q = process covariance

Time update:

Measurement update:

R = measurement covariance


K = Kalman gain

$$K_k = P_k^- H^T (HP_k^- H^T + R)^{-1}$$

$$\hat{x}_k = \hat{x}_k^- + K_k (z_k - H\hat{x}_k^-)$$


$$P_k = (I - K_k H)P_k^-$$

*a posteriori estimate*



# 2D EXAMPLE





state


$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

### Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$

Constant position


$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

system noise

$$\epsilon_t \sim \mathcal{N}(\mathbf{0}, R)$$

$$B\mathbf{u} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$R = \begin{bmatrix} \sigma_r^2 & 0 \\ 0 & \sigma_r^2 \end{bmatrix}$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$


measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$



Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$

zero-mean measurement noise

$$\mathbf{C} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\delta_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q})$$

$$\mathbf{Q} = \begin{bmatrix} \sigma_q^2 & 0 \\ 0 & \sigma_q^2 \end{bmatrix}$$



# Algorithm for the 2D object tracking example



$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

motion model

$$C = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

observation model

$$[x \ P] = \text{KF\_constPos}(x, P, z)$$

$$P = P + Q;$$

$$K = P / (P + R);$$

$$x = x + K * (z - x);$$



$$P = (\text{eye}(\text{size}(K, 1)) - K) * P;$$



Motion model of the Kalman filter is linear

$$x_t = Ax_{t-1} + Bu_t + \epsilon_t$$

but motion is not always linear









# How do you deal with non-linear models?


Output:  
Gaussian



approximate with a  
***linear function***

*When does this trick work?*

Input:  
Gaussian



# Extended Kalman Filter



- Does not assume linear Gaussian models
- Assumes Gaussian noise
- Uses local linear approximations of model to keep the efficiency of the KF framework

## Kalman Filter

linear motion model

$$x_t = Ax_{t-1} + Bu_t + \epsilon_t$$

linear sensor model

$$z_t = C_t x_t + \delta_t$$

## Extended Kalman Filter

non-linear motion model

$$x_t = g(x_{t-1}, u_t) + \epsilon_t$$

non-linear sensor model

$$z_t = H(x_t) + \delta_t$$



# Motion model linearization

$$g(x_{t-1}, u_t) \approx g(\mu_{t-1}, u_t) + \frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}} (x_{t-1} - \mu_{t-1})$$

Taylor series expansion



# Motion model linearization

$$\begin{aligned}g(x_{t-1}, u_t) &\approx g(\mu_{t-1}, u_t) + \frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}} (x_{t-1} - \mu_{t-1}) \\&\approx g(\mu_{t-1}, u_t) + G_t (x_{t-1} - \mu_{t-1})\end{aligned}$$



*What's this called?*



# Motion model linearization

$$\begin{aligned} g(x_{t-1}, u_t) &\approx g(\mu_{t-1}, u_t) + \frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}} (x_{t-1} - \mu_{t-1}) \\ &\approx g(\mu_{t-1}, u_t) + G_t (x_{t-1} - \mu_{t-1}) \end{aligned}$$



*What's this called?*

**Jacobian Matrix**

'the rate of change in x'  
'slope of the function'



## Motion model linearization


$$\begin{aligned}g(x_{t-1}, u_t) &\approx g(\mu_{t-1}, u_t) + \frac{\partial g(\mu_{t-1}, u_t)}{\partial x_{t-1}} (x_{t-1} - \mu_{t-1}) \\ &\approx g(\mu_{t-1}, u_t) + G_t (x_{t-1} - \mu_{t-1})\end{aligned}$$

### Jacobian Matrix

'the rate of change in x'  
'slope of the function'

## Sensor model linearization

$$\begin{aligned}h(x_t) &\approx h(\bar{\mu}_t) + \frac{\partial h(\bar{\mu}_t)}{\partial x_t} (x_{t-1} - \bar{\mu}_t) \\ &\approx h(\bar{\mu}_t) + H_t (x_t - \bar{\mu}_t)\end{aligned}$$



# New EKF Algorithm

(pretty much the same)

## Kalman Filter

$$\bar{\mu}_t = A_t \mu_{t-1} + B u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^\top + R$$

$$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

## Extended KF

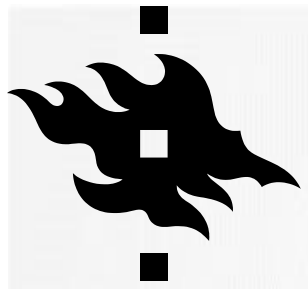
$$\bar{\mu}_t = g(\mu_{t-1}, u_t)$$

$$\bar{\Sigma}_t = G_t \bar{\Sigma}_{t-1} G_t^\top + R$$

$$K_t = \bar{\Sigma}_t H_t^\top (H_t \bar{\Sigma}_t H_t^\top + Q)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t))$$


$$\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$$



# 2D EXAMPLE








**state:** position-velocity

$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \\ y \\ \dot{y} \end{bmatrix} \begin{array}{l} \text{position} \\ \text{velocity} \\ \text{position} \\ \text{velocity} \end{array}$$

constant velocity motion model

$$A = \begin{bmatrix} 1 & \Delta t & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & \Delta t \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

with additive Gaussian noise



**measurement:** range-bearing


$$\begin{aligned} z &= \begin{bmatrix} r \\ \theta \end{bmatrix} \\ &= \begin{bmatrix} \sqrt{x^2 + y^2} \\ \tan^{-1}(y/x) \end{bmatrix} \end{aligned}$$

measurement model

*Is the measurement model linear?*

$$z = h(r, \theta)$$

with additive Gaussian noise



**measurement:** range-bearing

$$z = \begin{bmatrix} r \\ \theta \end{bmatrix} = \begin{bmatrix} \sqrt{x^2 + y^2} \\ \tan^{-1}(y/x) \end{bmatrix}$$

measurement model

*Is the measurement model linear?*

$$z = h(r, \theta)$$

with additive Gaussian noise

non-linear!

*What should we do?*



**linearize** the observation/measurement model!

$$\begin{aligned} z &= \begin{bmatrix} r \\ \theta \end{bmatrix} \\ &= \begin{bmatrix} \sqrt{x^2 + y^2} \\ \tan^{-1}(y/x) \end{bmatrix} \end{aligned}$$

$$H = \frac{\partial z}{\partial x} = ?$$

*What is the Jacobian?*

$$H = \begin{bmatrix} \frac{\partial r}{\partial x} & \frac{\partial r}{\partial \dot{x}} & \frac{\partial r}{\partial y} & \frac{\partial r}{\partial \dot{y}} \\ \frac{\partial \theta}{\partial x} & \frac{\partial \theta}{\partial \dot{x}} & \frac{\partial \theta}{\partial y} & \frac{\partial \theta}{\partial \dot{y}} \end{bmatrix} =$$



$$\begin{aligned} z &= \begin{bmatrix} r \\ \theta \end{bmatrix} \\ &= \begin{bmatrix} \sqrt{x^2 + y^2} \\ \tan^{-1}(y/x) \end{bmatrix} \end{aligned}$$

$$H = \frac{\partial z}{\partial x} = ?$$

*What is the Jacobian?*

Jacobian used in the Taylor series expansion looks like ...

$$H = \begin{bmatrix} \frac{\partial r}{\partial x} & \frac{\partial r}{\partial \dot{x}} & \frac{\partial r}{\partial y} & \frac{\partial r}{\partial \dot{y}} \\ \frac{\partial \theta}{\partial x} & \frac{\partial \theta}{\partial \dot{x}} & \frac{\partial \theta}{\partial y} & \frac{\partial \theta}{\partial \dot{y}} \end{bmatrix} = \begin{bmatrix} \cos(\theta) & 0 & \sin(\theta) & 0 \\ -\sin(\theta)/r & 0 & \cos(\theta)/r & 0 \end{bmatrix}$$



```
[x P] = EKF(x, P, z, dt)
```

```
r = sqrt(x(1)^2 + x(3)^2);  
b = atan2(x(3), x(1));  
y = [r; b];
```

```
H = [ cos(b)    0    sin(b)    0;  
      -sin(b)/r 0    cos(b)/r  0];
```

```
x = F*x;  
P = F*P*F' + Q;
```

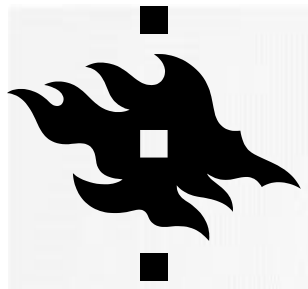
```
K = P*H' / (H*P*H' + R);
```

```
x = x + K*(z - y);  
P = (eye(size(K, 1)) - K*H) * P;
```

Parameters:

```
Q = diag([0 .1 0 .1]);  
R = diag([50^2 0.005^2]);  
F = [ 1 dt 0 0;  
      0 1 0 0;  
      0 0 1 dt;  
      0 0 0 1];
```

extra computation for  
the EKF measurement  
model Jacobian




# Problems with EKF

Taylor series expansion = poor approximation of non-linear functions  
success of linearization depends on limited uncertainty and amount of  
local non-linearity


Computing partial derivatives is a pain

Drifts when linearization is a bad approximation

Cannot handle multi-modal (multi-hypothesis) distributions



# Simultaneous Localization and Mapping



Given a **single camera** feed,  
estimate the 3D **position of the camera** and  
the 3D **positions of all landmark** points in the world





## *What is the camera (robot) state?*

*What are the dimensions?*

$$\mathbf{x}_c = \begin{bmatrix} \mathbf{r} \\ \mathbf{q} \\ \mathbf{v} \\ \boldsymbol{\omega} \end{bmatrix}$$


position

rotation (quaternion)

velocity

angular velocity

**13 total**






## What is the camera (robot) state?

What are the dimensions?

$$\mathbf{x}_c = \begin{bmatrix} \mathbf{r} \\ \mathbf{q} \\ \mathbf{v} \\ \boldsymbol{\omega} \end{bmatrix}$$

position	3
rotation (quaternion)	4
velocity	3
angular velocity	3

**13 total**





## *What is the world (robot+environment) state?*



*What are the dimensions?*

$$\mathbf{x} = \begin{bmatrix} \mathbf{x}_c \\ \mathbf{y}_1 \\ \mathbf{y}_2 \\ \vdots \\ \mathbf{y}_N \end{bmatrix}$$

state of the camera

location of feature 1

location of feature 2

location of feature N



## *What is the world (robot+environment) state?*


*What are the dimensions?*

$$\mathbf{x} = \begin{bmatrix} \mathbf{x}_c \\ \mathbf{y}_1 \\ \mathbf{y}_2 \\ \vdots \\ \mathbf{y}_N \end{bmatrix}$$

state of the camera	13
location of feature 1	3
location of feature 2	3
location of feature N	3
<b>13+3N total</b>	



Observations are...



detected visual features of landmark points.  
(e.g., Harris corners)

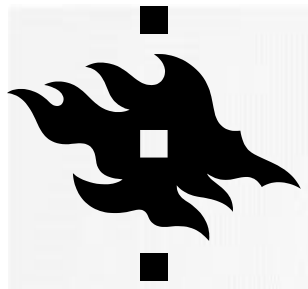


*What is the motion model?*  $P(\mathbf{x}_t | \mathbf{x}_{t-1})$

**Landmarks:**  
constant position  
(identity matrix)

**Camera:**  
constant velocity  
(not identity matrix and non-linear) **EKF!**

*What is the form of the belief?*  $P(\mathbf{x}_t | \mathbf{z}_{1:t-1})$



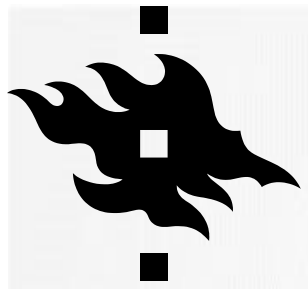
*What is the motion model?*  $P(\mathbf{x}_t | \mathbf{x}_{t-1})$

**Landmarks:**  
constant position  
(identity matrix)

**Camera:**  
constant velocity  
(not identity matrix and non-linear)

*What is the form of the belief?*  $P(\mathbf{x}_t | \mathbf{z}_{1:t-1})$

**Gaussian!**  
(everything will be parametrized by a mean and variance)



## Constant Velocity Motion Model

$$\mathbf{r}_t = \mathbf{r}_{t-1} + \mathbf{v}_{t-1} \Delta t$$

position

$$\mathbf{q}_t = \mathbf{q}_{t-1} \times [\mathbf{q}(\omega) \Delta t]$$

rotation (quaternion)

$$\mathbf{v}_t = \mathbf{v}_{t-1}$$

velocity

$$\omega_t = \omega_{t-1}$$

angular velocity





## Gaussian noise uncertainty (only on velocity)

$$\mathbf{v}_t = \mathbf{v}_{t-1} + \mathbf{V}$$

$$\omega_t = \omega_{t-1} + \Omega$$

$$\mathbf{V} \sim \mathcal{N}(\mathbf{0}, \begin{bmatrix} \sigma_v & 0 & 0 \\ 0 & \sigma_v & 0 \\ 0 & 0 & \sigma_v \end{bmatrix})$$

$$\Omega \sim \mathcal{N}(0, \begin{bmatrix} \sigma_w & 0 & 0 \\ 0 & \sigma_w & 0 \\ 0 & 0 & \sigma_w \end{bmatrix})$$



## Prediction (**mean** of camera state):

$$P(\mathbf{x}_t | \mathbf{z}_{1:t-1}) = \int_{\mathbf{x}_{t-1}} P(\mathbf{x}_t | \mathbf{x}_{t-1}) P(\mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}) d\mathbf{x}_{t-1}$$

$$\mathbf{f}_t = \begin{bmatrix} \mathbf{r}_t \\ \mathbf{q}_t \\ \mathbf{v}_t \\ \omega_t \end{bmatrix} = \begin{bmatrix} \mathbf{r}_{t-1} + \mathbf{v}_{t-1} \Delta t \\ \mathbf{q}_{t-1} + \mathbf{q}(\omega)_{t-1} \Delta t \\ \mathbf{v}_{t-1} \\ \omega_{t-1} \end{bmatrix}$$



## Prediction (**covariance** of camera state):



$$P(\mathbf{x}_t | \mathbf{z}_{1:t-1}) = \int_{\mathbf{x}_{t-1}} P(\mathbf{x}_t | \mathbf{x}_{t-1}) P(\mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}) d\mathbf{x}_{t-1}$$

$$\bar{\Sigma}_{\mathbf{x}\mathbf{x}} = \boxed{\frac{\partial \mathbf{f}_t}{\partial \mathbf{x}}} \Sigma_{\mathbf{x}\mathbf{x}} \frac{\partial \mathbf{f}_t}{\partial \mathbf{x}}^\top + \mathbf{Q}_t$$

new covariance      change around new state      old covariance      change around new state      system noise (process noise)



Bit of a pain to compute this term...



We just covered the **prediction** step for the camera state

$$P(\mathbf{x}_t | \mathbf{z}_{1:t-1}) = \int_{\mathbf{x}_{t-1}} P(\mathbf{x}_t | \mathbf{x}_{t-1}) P(\mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}) d\mathbf{x}_{t-1}$$

$$\mathbf{f}_t = \begin{bmatrix} \mathbf{r}_t \\ \mathbf{q}_t \\ \mathbf{v}_t \\ \omega_t \end{bmatrix} = \begin{bmatrix} \mathbf{r}_{t-1} + \mathbf{v}_{t-1} \\ \mathbf{q}_{t-1} + \mathbf{q}(\omega)_{t-1} \\ \mathbf{v}_{t-1} \\ \omega_{t-1} \end{bmatrix}$$

$$\bar{\Sigma}_{\mathbf{xx}} = \frac{\partial \mathbf{f}_t}{\partial \mathbf{x}} \Sigma_{\mathbf{xx}} \frac{\partial \mathbf{f}_t}{\partial \mathbf{x}}^\top + \mathbf{Q}_t$$

Now we need to do the **update** step!



## General Filtering Equations

$$P(\mathbf{x}_t | \mathbf{z}_{1:t}) \propto P(\mathbf{z}_t | \mathbf{x}_t) \int_{\mathbf{x}_{t-1}} P(\mathbf{x}_t | \mathbf{x}_{t-1}) P(\mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}) d\mathbf{x}_{t-1}$$

**Prediction:**

$$P(\mathbf{x}_t | \mathbf{z}_{1:t-1}) = \int_{\mathbf{x}_{t-1}} P(\mathbf{x}_t | \mathbf{x}_{t-1}) P(\mathbf{x}_{t-1} | \mathbf{z}_{1:t-1}) d\mathbf{x}_{t-1}$$

**Update:**

$$P(\mathbf{x}_t | \mathbf{z}_{1:t}) = P(\mathbf{z}_t | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{z}_{1:t-1})$$



Belief state


State observation

Predicted State

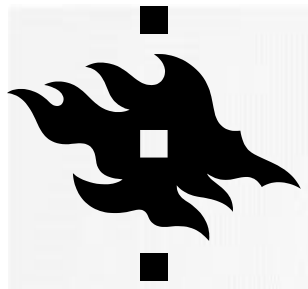
$$P(\mathbf{x}_t | \mathbf{z}_{1:t}) = P(\mathbf{z}_t | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{z}_{1:t-1})$$



*What are the observations?*




2D projections of 3D landmarks



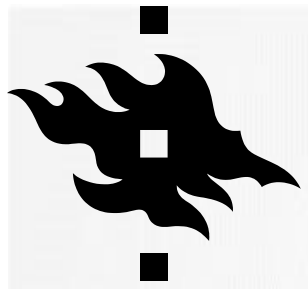
Recall, the state includes the 3D location of landmarks

$$\mathbf{x} = \begin{bmatrix} \mathbf{x}_c \\ \mathbf{y}_1 \\ \mathbf{y}_2 \\ \vdots \\ \mathbf{y}_N \end{bmatrix}$$

*What is the projection from 3D point to 2D image point?*



Slide credit: Kris Kitani



# Observation Model

$$P(z_t | x_t)$$

*If you know the 3D location of a landmark, what is the 2D projection?*

Non-linear observation model

$$\mathbf{h} \sim \mathbf{P} \mathbf{y}$$

2D Image Point      Camera matrix      3D World Point

$$\mathbf{P} = \mathbf{K} [\mathbf{R} | \mathbf{T}]$$

*What do we know about  $\mathbf{P}$ ?*

*How do we make the observation model linear?*





$$H = \frac{\partial h}{\partial x}$$

(2n x 13)

n: number of visible points



$$P(\mathbf{x}_t | \mathbf{z}_{1:t}) = P(\mathbf{z}_t | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{z}_{1:t-1})$$

Update step (mean):

$$\underset{\text{Updated state}}{\mathbf{x}_t} = \underset{\text{Predicted state}}{\mathbf{x}_t} + \overset{\text{Kalman gain}}{\mathbf{K}_t} (\underset{\text{Matched 2D features}}{\mathbf{z}_t} - \underset{\text{2D projection of 3D point}}{\mathbf{h}(\mathbf{y}; \mathbf{x}_t)})$$


Update step (covariance):

$$\underset{\text{Covariance (updated)}}{\Sigma_t} = (\underset{\text{Identity}}{\mathbf{I}} - \overset{\text{Kalman gain}}{\mathbf{K}_t} \underset{\text{Jacobian}}{\mathbf{H}_t}) \underset{\text{Covariance (predicted)}}{\Sigma_t}$$



# KEYFRAME- BASED SLAM

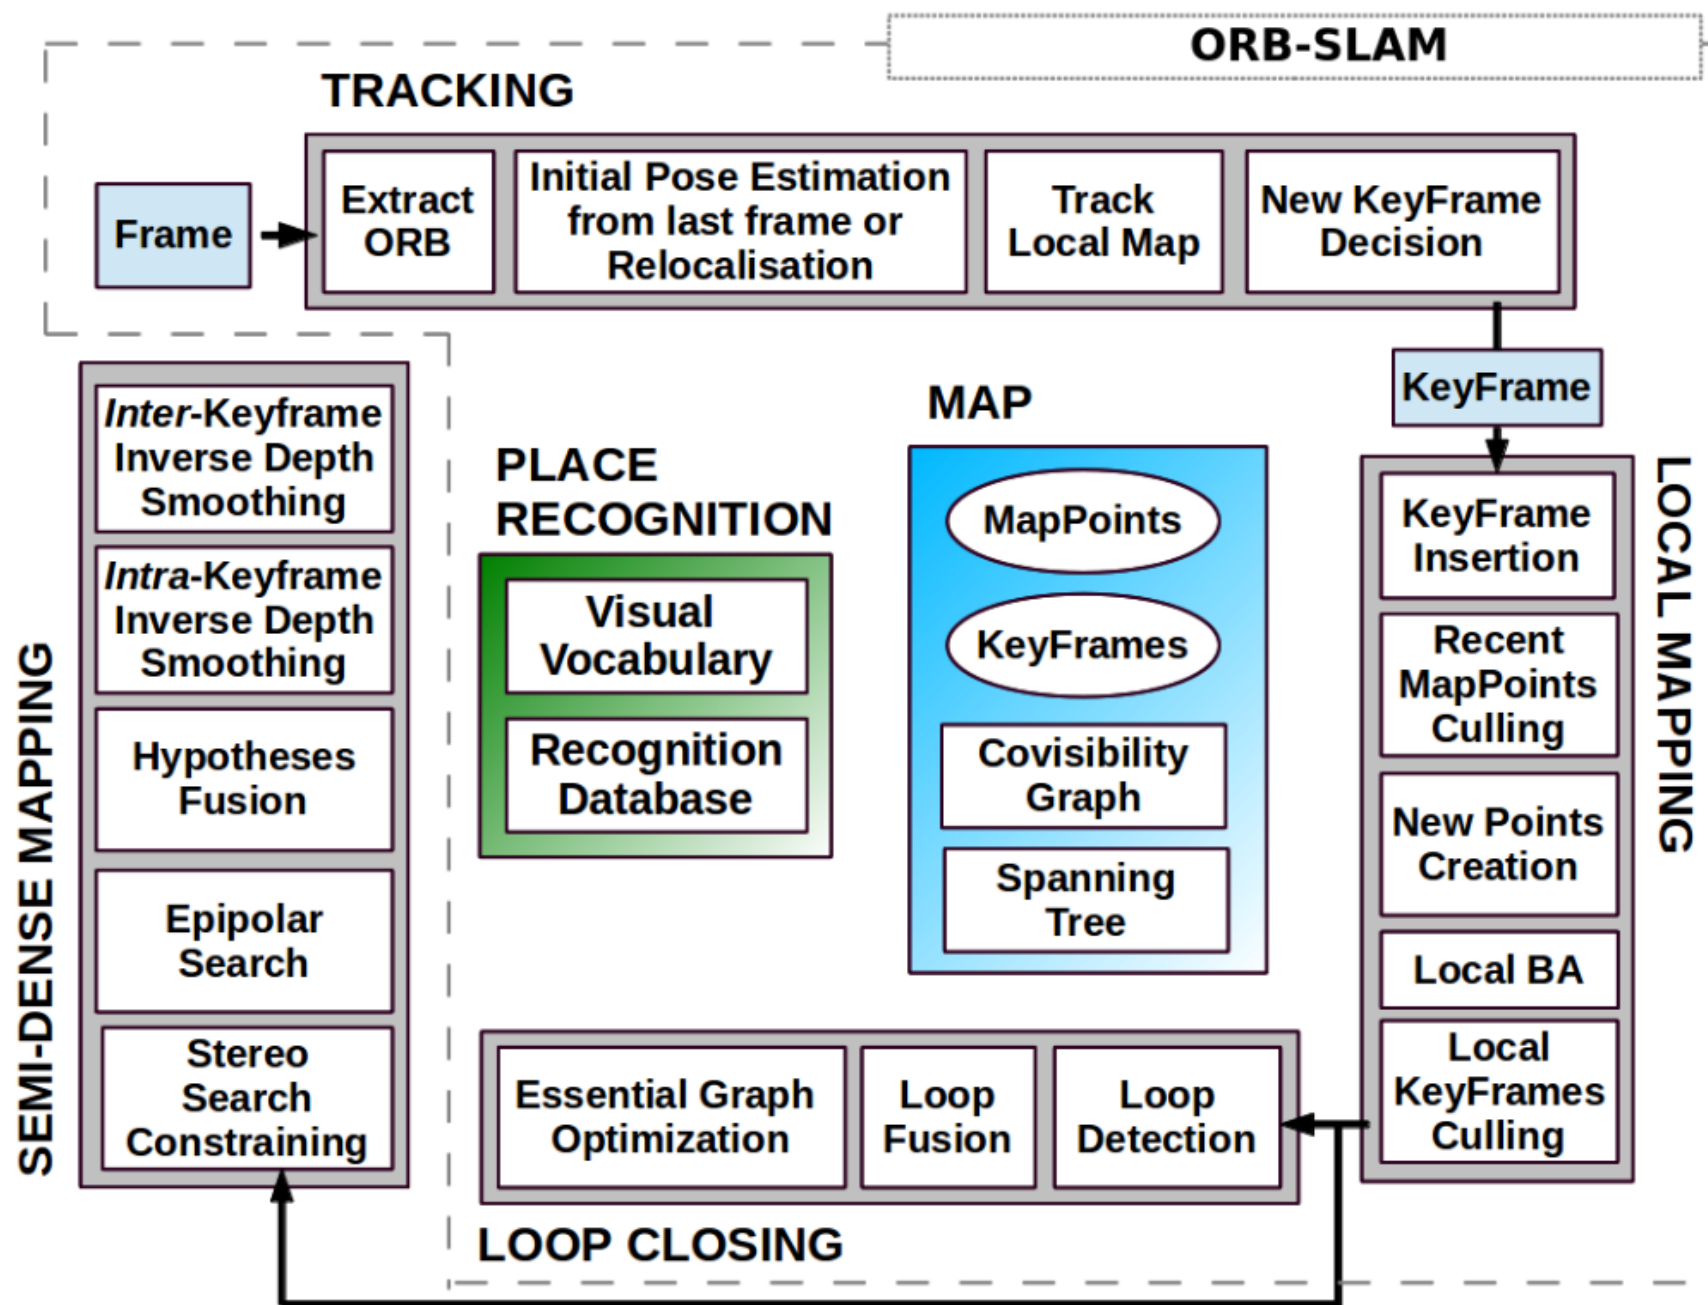

- Filterig is replaced with bundle adjustment
  - Keep every measurement that goes into the map
  - Keeping all previous poses in estimation is too computationally heavy => using only subset of poses = **Key-frames**
  - Split map-making and camera pose tracking into two separate threads




# ORB-SLAM

- State-of-the-art feature and keyframe based SLAM
- ORB name comes from the ORB features that are tracked
  - Similar to SIFT but much faster to be computed
- Semi-dense mapping

<https://www.youtube.com/watch?v=HlBmq70LKrQ&feature=youtu.be>







# PROBABILISTIC SEMI-DENSE MAPPING



KeyFrame

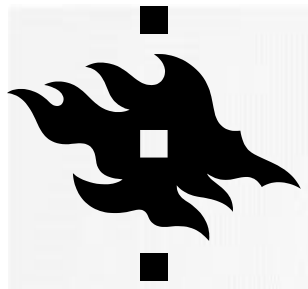


High Gradient Pixels





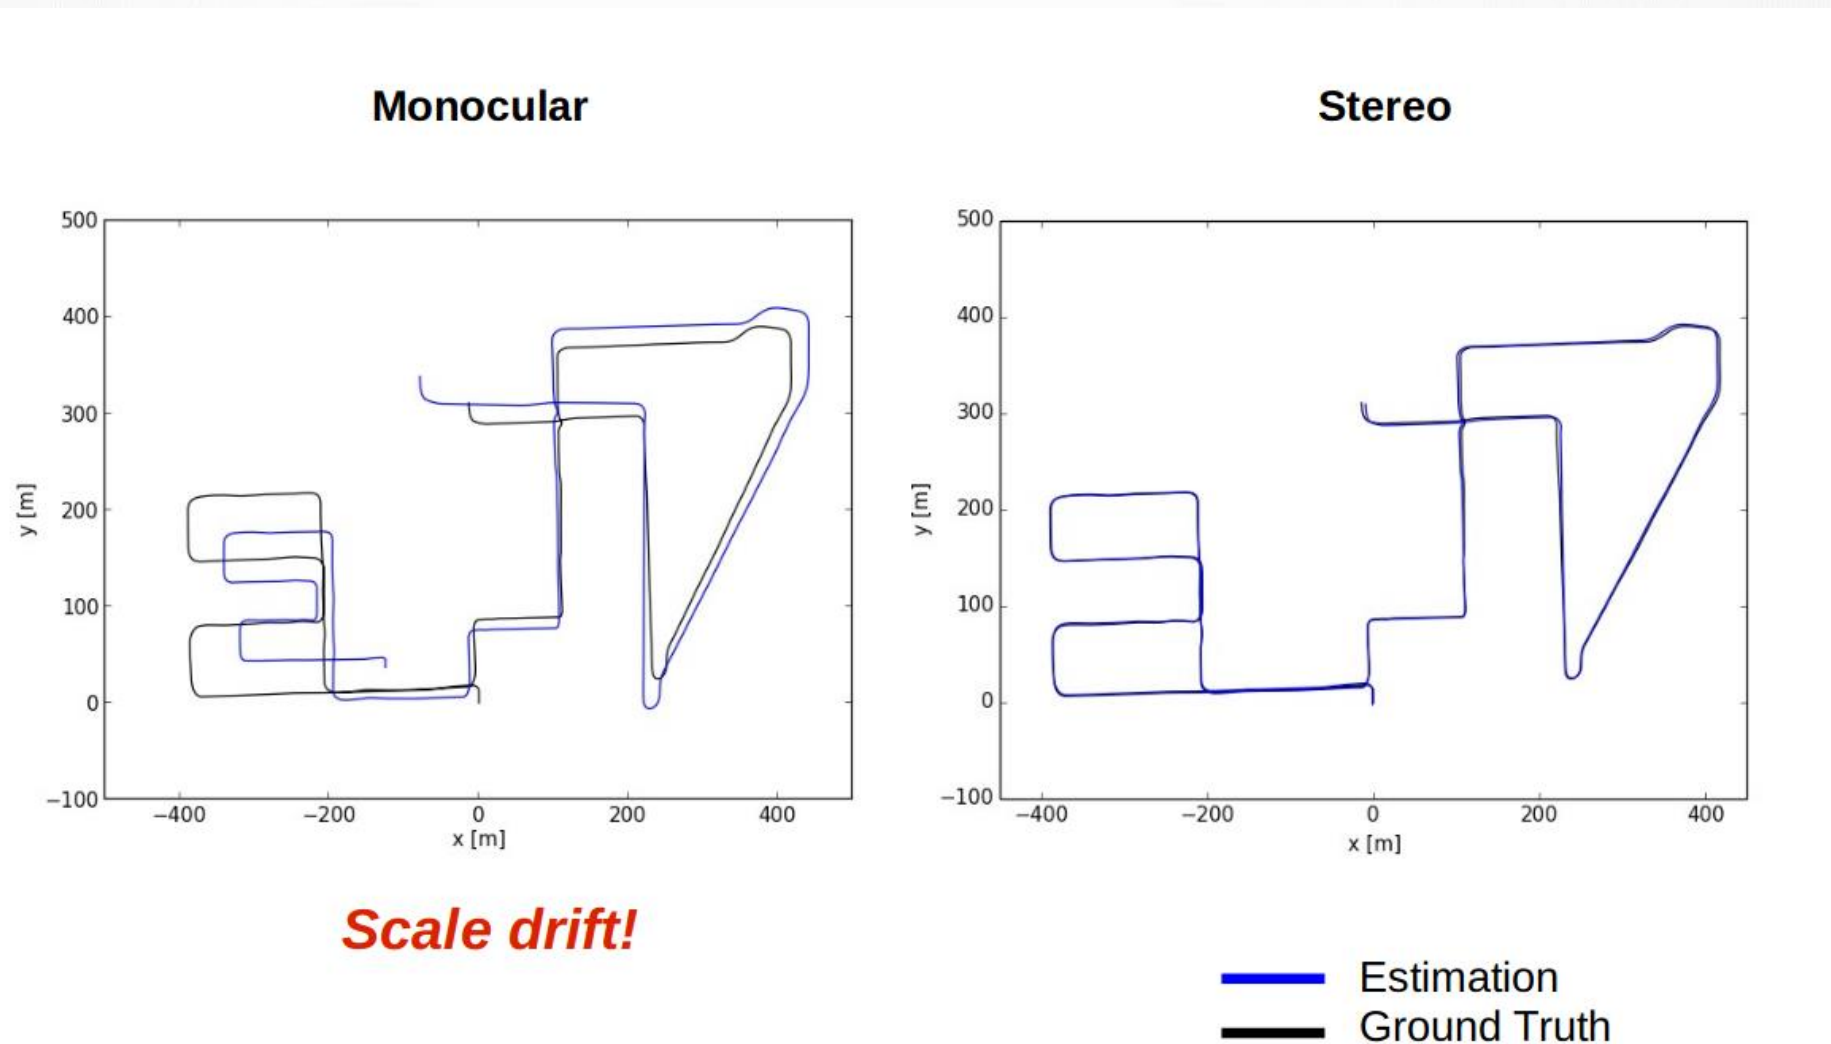
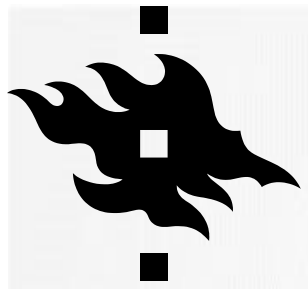
Inverse Depth Map  
& uncertainty

Compute each inverse depth map from scratch using neighbor keyframes



## Per Pixel Operations






**Scale drift!**





# CNN-SLAM: REAL-TIME DENSE MONOCULAR SLAM WITH LEARNED DEPTH PREDICTION

[https://www.youtube.com/watch?v=z\\_NJxbkQnBU](https://www.youtube.com/watch?v=z_NJxbkQnBU)





60° 10 1.2 N, 24° 57 18 E